

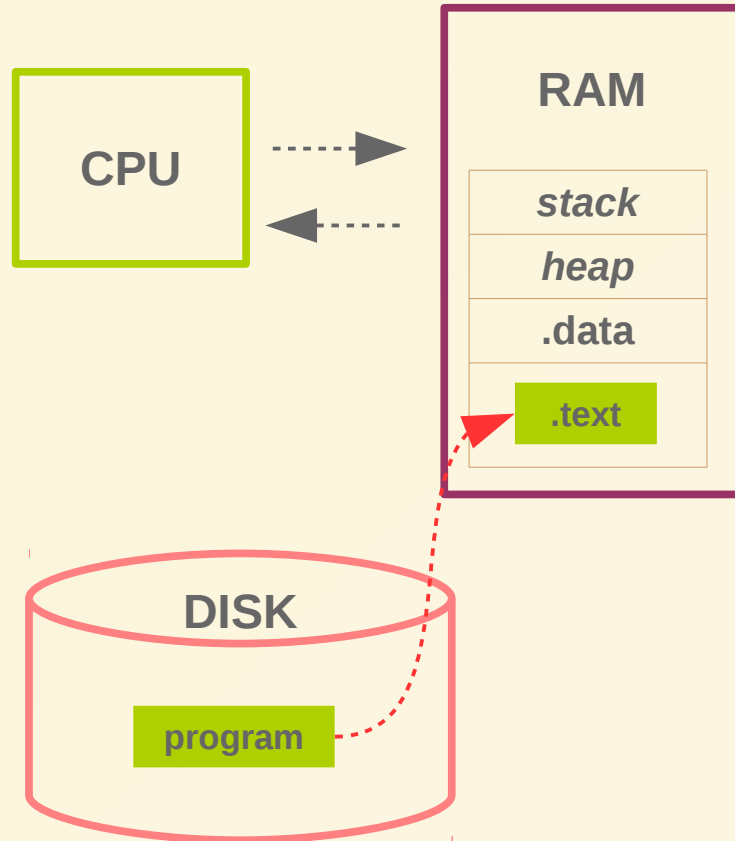
***“Goodbye ! printf()”* hands-on tutorials with uftime: function graph tracer for C/C++**

Taeung Song 송태웅 , KOSSLAB

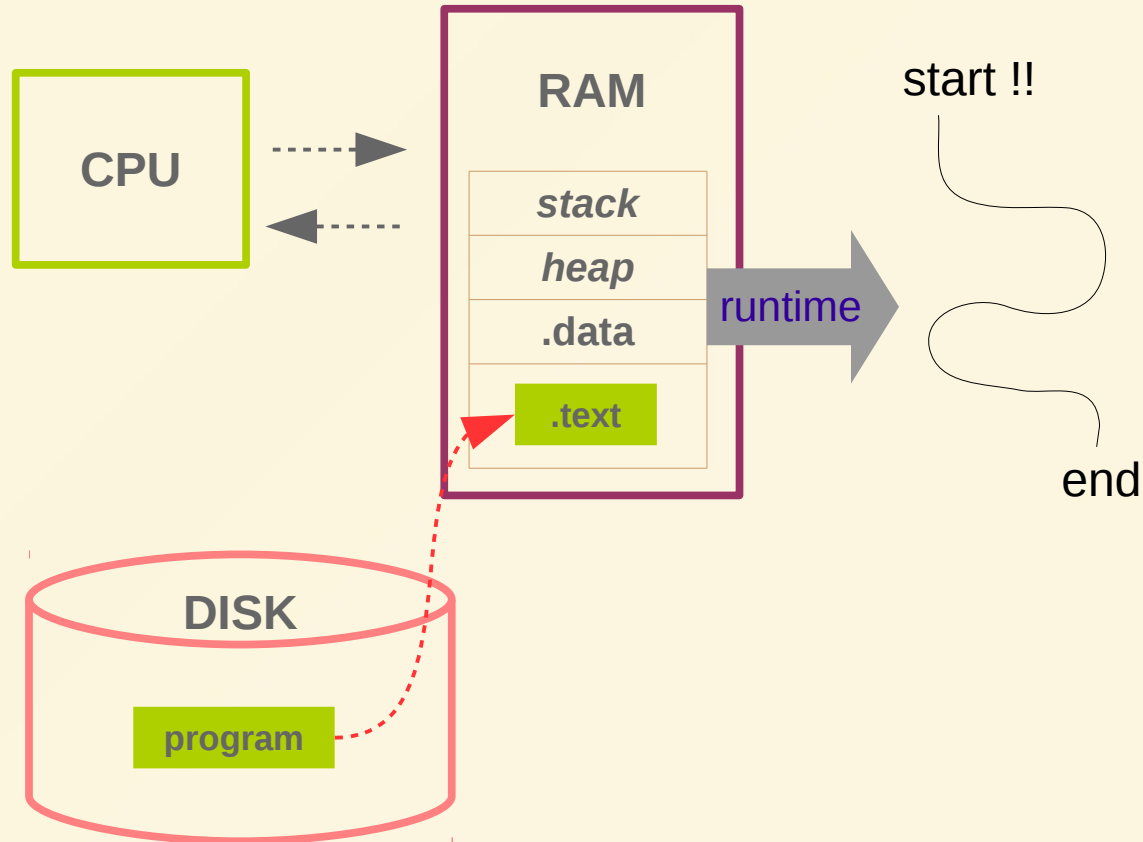
Why use **printf()** ?
during development

As you know..


Program Execution



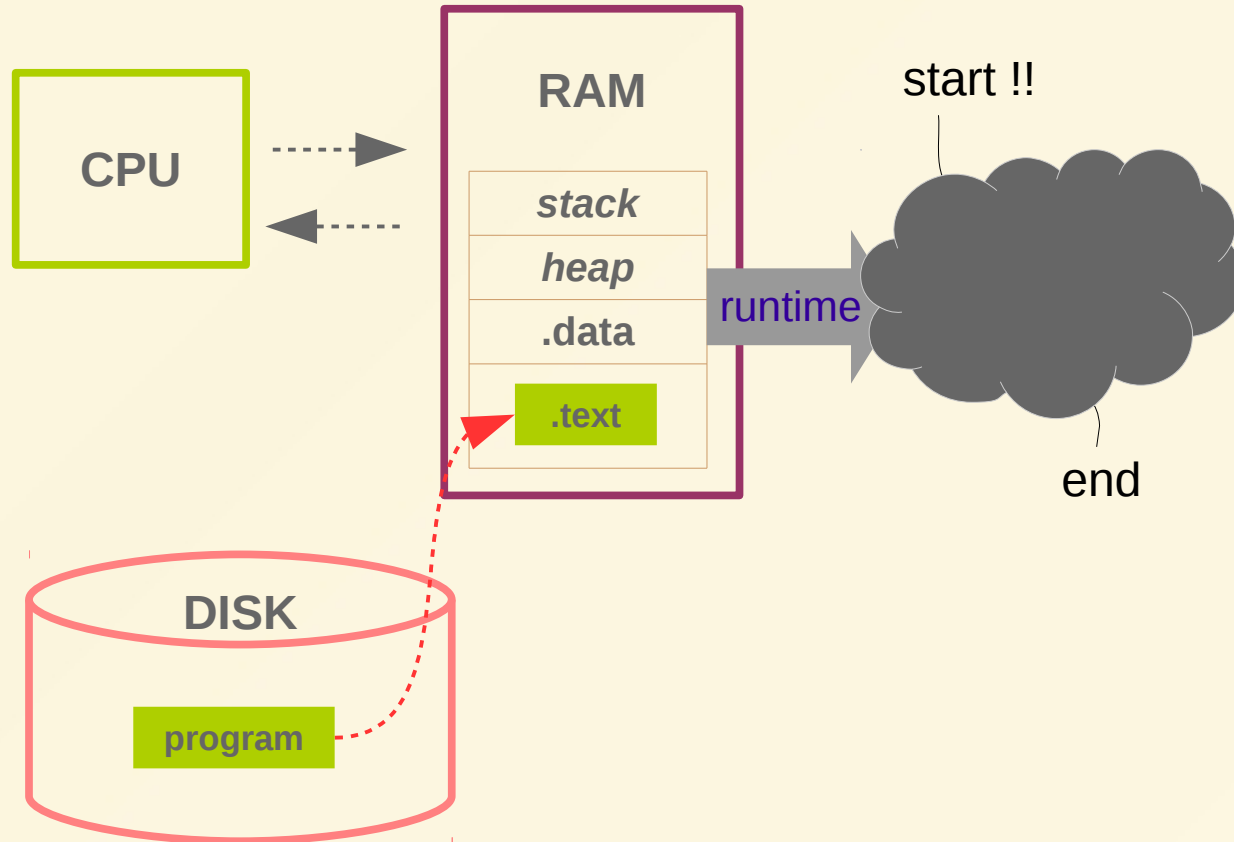
Execution Flow is



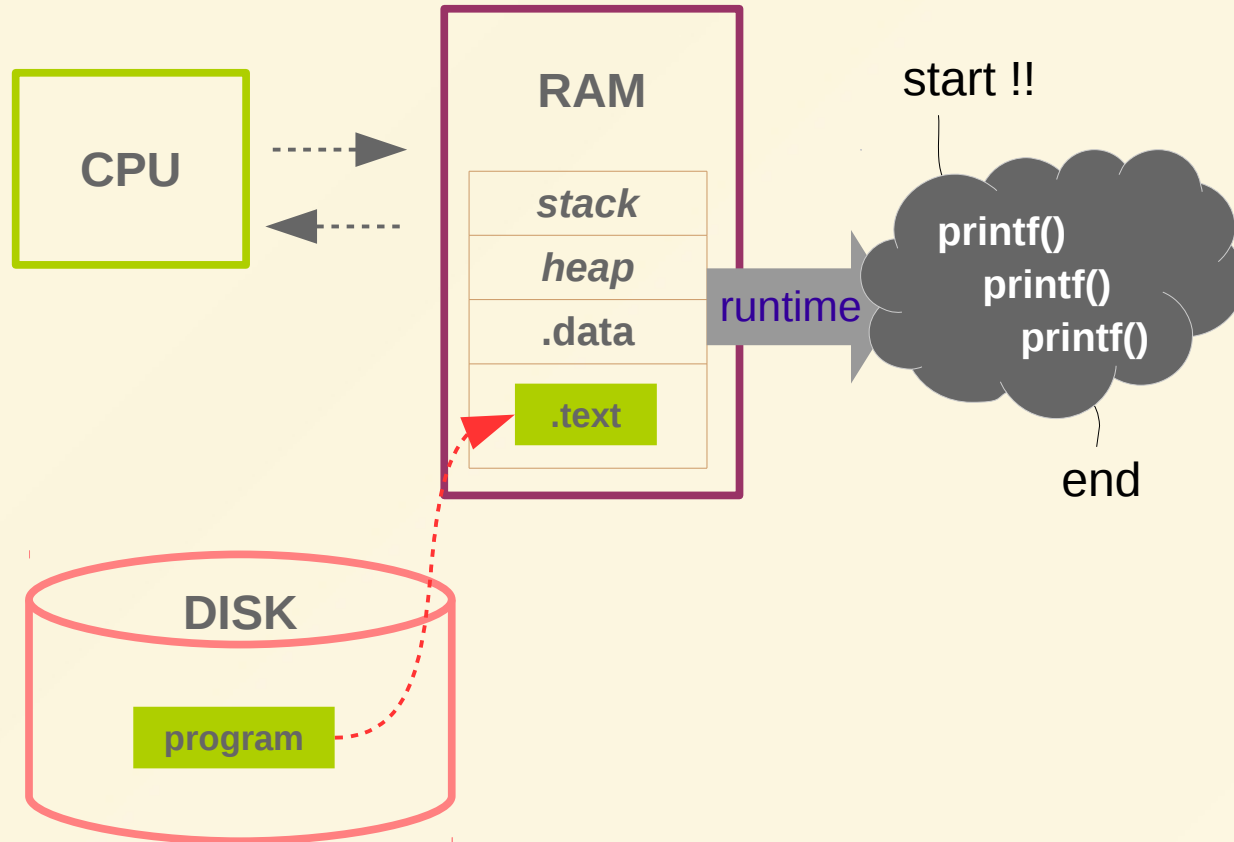
Execution Flow is **invisible**



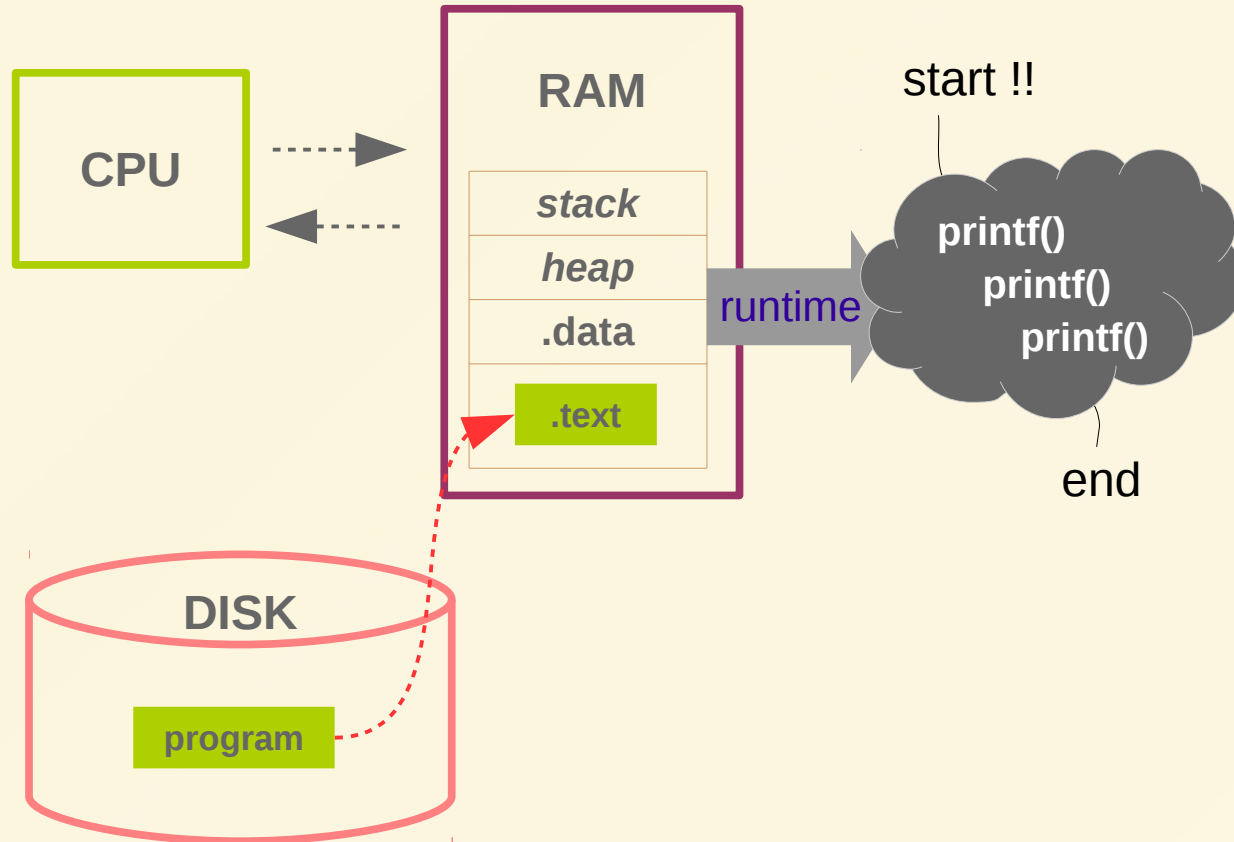
How to **solve** it ?



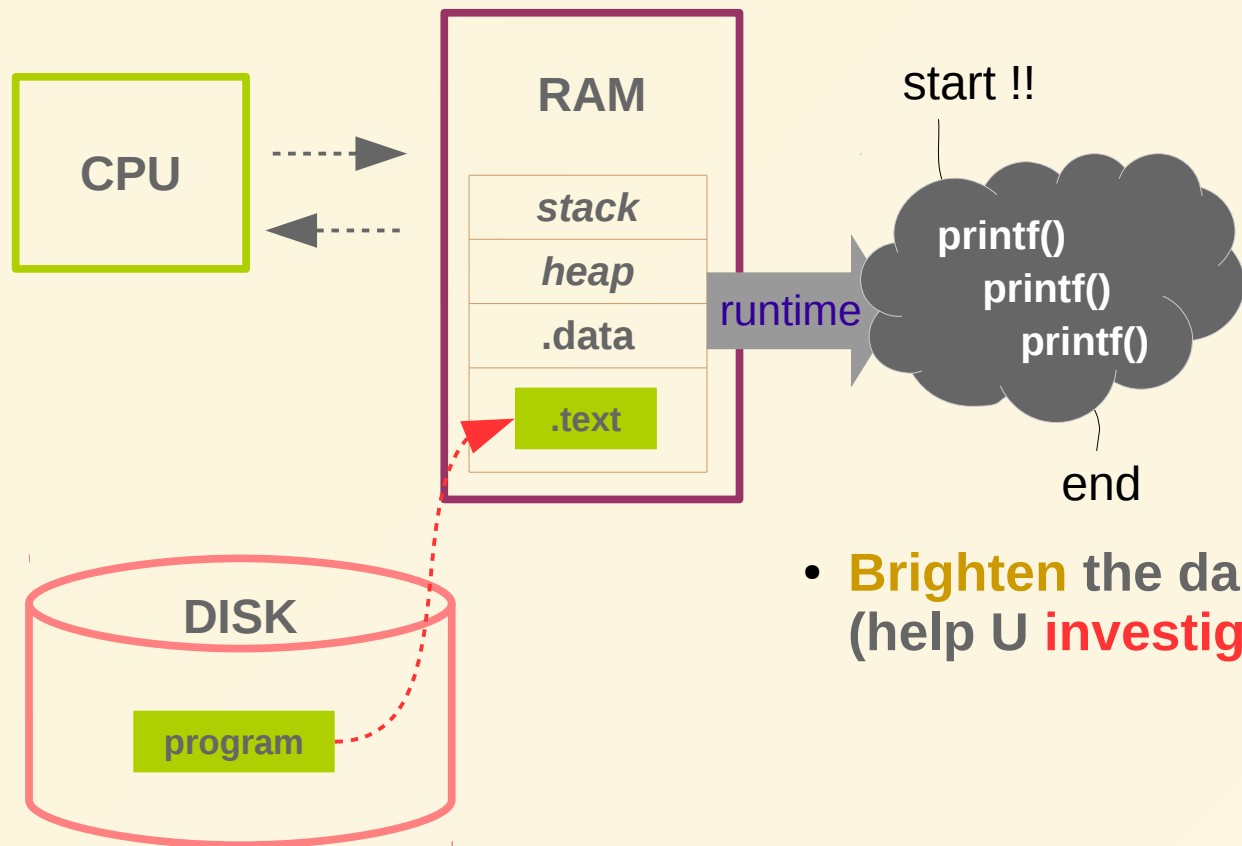
The proper solution is **printf()** !!



Why **printf()** ?

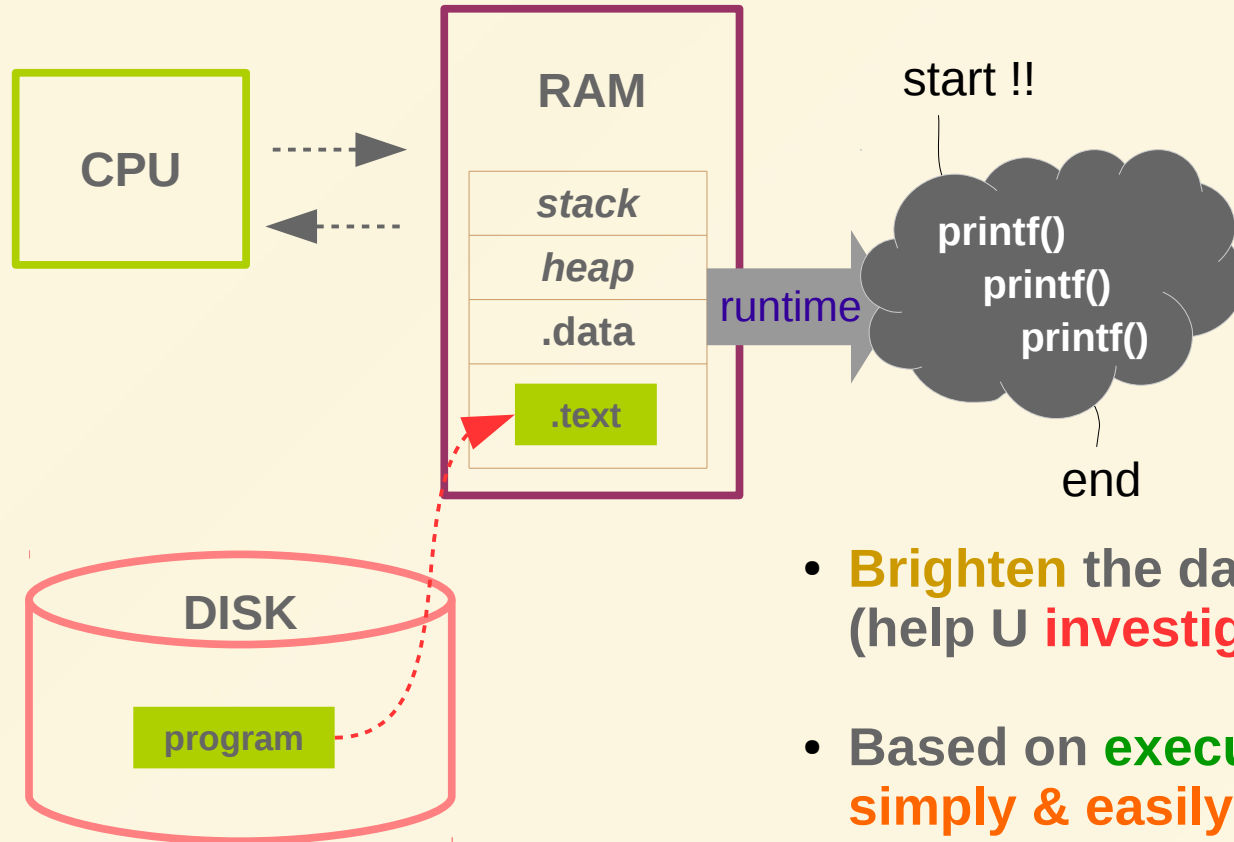


Why **printf()** ?



- **Brighten** the dark
(help U **investigate** it's internals)

Why **printf()** ?



- **Brighten** the dark (help U **investigate** it's internals)
- Based on **execution flow** **simply & easily**

Benefit of printf()

- **Brighten** the dark of process
(help U **investigate** internals)
- Based on **execution flow**
simply & easily

uftrace

+α

Benefit of printf()

- **Brighten** the dark of process
(help U **investigate** internals)
- Based on **execution flow**
simply & easily

uftrace

+α : without source code modification

Benefit of printf()

- Brighten the dark of process (help U investigate internals)
- Based on execution flow simply & easily

uftrace

+ α : without source code modification
+ β , ...

Benefit of printf()

- Brighten the dark of process (help U investigate internals)
- Based on execution flow simply & easily



“Goodbye ! printf()”

UFTRACE

<https://github.com/namhyung/uftrace>

<https://gitter.im/uftrace/uftrace>



This repository

Search

Pull requests

Issues

Marketplace

Explore



namhyung / uftace

Unwatch

46

★ Unstar

497

Fork

82

<> Code

! Issues 27

Pull requests 7

Wiki

Insights

Function (graph) tracer for user-space

trace

function

tracer

2,129 commits

24 branches

9 releases

14 contributors

GPL-2.0

Branch: master

New pull request

Create new file

Upload files

Find file

Clone or download



namhyung

Merge branch 'filter-fix'



Latest commit d116982 2 days ago

arch	mcount: Remove a residue 'cmpq' after mcount_entry()	4 days ago
check-deps	build: Remove all executable files for check-clean	2 months ago
doc	trigger: Add "filter" and "notrace" actions	5 days ago
libmcount	filter: Allow setting filter on kernel functions	2 days ago
libtraceevent	libtraceevent: Adjust switch statements for fall-through warning on g...	6 months ago

build passing coverity passed

uftrace

The uftrace tool is to trace and analyze execution of a program written in C/C++. It was heavily inspired by the ftrace framework of the Linux kernel (especially function graph tracer) and supports userspace programs. It supports various kind of commands and filters to help analysis of the program execution and performance.

- Homepage: <https://github.com/namhyung/uftrace>
- Tutorial: <https://github.com/namhyung/uftrace/wiki/Tutorial>
- Chat: <https://gitter.im/uftrace/uftrace>
- Mailing list: uftrace@googlegroups.com

build passing coverity passed

uftrace

The uftrace tool is to trace and analyze execution of a program written in C/C++. It was heavily inspired by the ftrace framework of the Linux kernel (especially function graph tracer) and supports userspace programs. It supports various kind of commands and filters to help analysis of the program execution and performance.

- Homepage: <https://github.com/namhyung/uftrace>
- Tutorial: <https://github.com/namhyung/uftrace/wiki/Tutorial>
- Chat: <https://gitter.im/uftrace/uftrace>
- Mailing list: uftrace@googlegroups.com



Come in and ask questions !!

Introduction to **uftrace**



Function **tracer** for C/C++ programs


- created by Namhyung Kim
 - LG Electronics open-source contribution team
 - Linux kernel developer (since 2010)
 - perf, ftrace, ...
- inspired by ftrace framework in the kernel
- **record** and **replay** model

```
$ gcc -pg -o fibonacci tests/s-fibonacci.c
```

```
$ ufttrace -A fib@arg1 -R fib@retval fibonacci 5
```

#	DURATION	TID	FUNCTION
	0.633 us	[2851]	__monstartup();
	0.480 us	[2851]	__cxa_atexit();
		[2851]	main() {
	0.546 us	[2851]	atoi();
		[2851]	fib(5) {
		[2851]	fib(4) {
		[2851]	fib(3) {
	1.146 us	[2851]	fib(2) = 1;
	0.077 us	[2851]	fib(1) = 1;
	1.823 us	[2851]	} = 2; /* fib */
	0.062 us	[2851]	fib(2) = 1;
	2.199 us	[2851]	} = 3; /* fib */
		[2851]	fib(3) {
	0.061 us	[2851]	fib(2) = 1;
	0.067 us	[2851]	fib(1) = 1;
	0.474 us	[2851]	} = 2; /* fib */
	3.317 us	[2851]	} = 5; /* fib */
	4.343 us	[2851]	} /* main */

ufttrace




“C/C++ execution flow”

```
# Default == record + replay

$ ufttrace -A fib@arg1 -R fib@retval fibonacci 5
# DURATION      TID      FUNCTION
  0.633 us [ 2851] | __monstartup();
  0.480 us [ 2851] | __cxa_atexit();
  0.546 us [ 2851] | main() {
    0.546 us [ 2851] |     atoi();
    0.546 us [ 2851] |     fib(5) {
    0.546 us [ 2851] |         fib(4) {
    0.546 us [ 2851] |             fib(3) {
    0.546 us [ 2851] |                 fib(2) = 1;
    0.546 us [ 2851] |                 fib(1) = 1;
    0.546 us [ 2851] |             } = 2; /* fib */
    0.546 us [ 2851] |         fib(2) = 1;
    0.546 us [ 2851] |     } = 3; /* fib */
    0.546 us [ 2851] |     fib(3) {
    0.546 us [ 2851] |         fib(2) = 1;
    0.546 us [ 2851] |         fib(1) = 1;
    0.546 us [ 2851] |     } = 2; /* fib */
    0.546 us [ 2851] | } = 5; /* fib */
  4.343 us [ 2851] | } /* main */
```

ufttrace




“C/C++ execution flow”

Function Call Trace

```
# Default == record + replay

$ ufttrace -A fib@arg1 -R fib@retval fibonacci 5
# DURATION      TID      FUNCTION
0.633 us [ 2851] | __monstartup();
0.480 us [ 2851] | __cxa_atexit();
0.546 us [ 2851] | main() {
    0.546 us [ 2851] |     atoi();
    0.546 us [ 2851] |     fib(5) {
        0.546 us [ 2851] |         fib(4) {
            0.546 us [ 2851] |             fib(3) {
                1.146 us [ 2851] |                 fib(2) = 1;
                0.077 us [ 2851] |                 fib(1) = 1;
                1.823 us [ 2851] |                 } = 2; /* fib */
                0.062 us [ 2851] |                 fib(2) = 1;
                2.199 us [ 2851] |                 } = 3; /* fib */
                0.061 us [ 2851] |                 fib(3) {
                    0.067 us [ 2851] |                     fib(2) = 1;
                    0.474 us [ 2851] |                     fib(1) = 1;
                    3.317 us [ 2851] |                     } = 2; /* fib */
                    4.343 us [ 2851] |                 } = 5; /* fib */
            } /* main */
```

ufttrace



“C/C++ each function

Execution time”


```
# Default == record + replay

$ uftrace -A fib@arg1 -R fib@retval fibonacci 5
# DURATION      TID      FUNCTION
  0.633 us [ 2851] | __monstartup();
  0.480 us [ 2851] | __cxa_atexit();
  0.546 us [ 2851] | main() {
    0.546 us [ 2851] |     atoi();
    0.546 us [ 2851] |     fib(5) {
      0.546 us [ 2851] |         fib(4) {
        0.546 us [ 2851] |             fib(3) {
          1.146 us [ 2851] |                 fib(2) = 1;
          0.077 us [ 2851] |                 fib(1) = 1;
          1.823 us [ 2851] |                 } = 2; /* fib */
          0.062 us [ 2851] |                 fib(2) = 1;
          2.199 us [ 2851] |                 } = 3; /* fib */
          2.199 us [ 2851] |                 fib(3) {
            0.061 us [ 2851] |                     fib(2) = 1;
            0.067 us [ 2851] |                     fib(1) = 1;
            0.474 us [ 2851] |                     } = 2; /* fib */
            3.317 us [ 2851] |                 } = 5; /* fib */
          4.343 us [ 2851] |             } /* main */
```

uftrace




“Arguments”

based on
Function Call Trace

```
# Default == record + replay

$ uftrace -A fib@arg1 -R fib@retval fibonacci 5
# DURATION      TID      FUNCTION
  0.633 us [ 2851] | __monstartup();
  0.480 us [ 2851] | __cxa_atexit();
  0.546 us [ 2851] | main() {
    0.546 us [ 2851] |     atoi();
    0.546 us [ 2851] |     fib(5) {
      0.546 us [ 2851] |         fib(4) {
        0.546 us [ 2851] |             fib(3) {
          1.146 us [ 2851] |                 fib(2) = 1;
          0.077 us [ 2851] |                 fib(1) = 1;
          1.823 us [ 2851] |                 } = 2; /* fib */
          0.062 us [ 2851] |                 fib(2) = 1;
          2.199 us [ 2851] |                 } = 3; /* fib */
          2.199 us [ 2851] |             fib(3) {
            0.061 us [ 2851] |                 fib(2) = 1;
            0.067 us [ 2851] |                 fib(1) = 1;
            0.474 us [ 2851] |                 } = 2; /* fib */
            3.317 us [ 2851] |             } = 5; /* fib */
          4.343 us [ 2851] |         } /* main */
```

uftrace



“Return values”


based on
Function Call Trace

```
$ uftrace record fibonacci 5
```

```
$ uftrace report -s call
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
1.400 us	1.400 us	9	fib
0.398 us	0.398 us	1	__cxa_atexit
0.489 us	0.489 us	1	__monstartup
0.603 us	0.603 us	1	atoi
2.454 us	0.451 us	1	main

uftrace report



Statistics

“Duration (time)”

“Function call counts”

```
$ uftrace record fibonacci 5

$ uftrace graph
#
# function graph for 'main' (session:
58de6d06edafbe8d)
#

backtrace
=====
backtrace #0: hit 1, time    2.454 us
  [0] main (0x4006d6)

calling functions
=====
  2.454 us : (1) main
  0.603 us : +-(1) atoi
             |
  1.400 us : +-(1) fib
  1.161 us : (2) fib
  0.681 us : (4) fib
  0.139 us : (2) fib
```

uftrace **graph**



Repeated func calls

“Total **time**”

“Hit **counts**”

```
$ ufttrace dump --chrome
```

```
{ "traceEvents": [
  { "ts": 326310414447.272, "ph": "B", "pid": 11326, "name": "__monstartup" },
  { "ts": 326310414447.761, "ph": "E", "pid": 11326, "name": "__monstartup" },
  { "ts": 326310414449.567, "ph": "B", "pid": 11326, "name": "__cxa_atexit" },
  { "ts": 326310414449.965, "ph": "E", "pid": 11326, "name": "__cxa_atexit" },
  { "ts": 326310414450.491, "ph": "B", "pid": 11326, "name": "main" },
  { "ts": 326310414450.588, "ph": "B", "pid": 11326, "name": "atoi" },
  { "ts": 326310414451.191, "ph": "E", "pid": 11326, "name": "atoi" },
  { "ts": 326310414451.455, "ph": "B", "pid": 11326, "name": "fib" },
  { "ts": 326310414451.547, "ph": "B", "pid": 11326, "name": "fib" },
  { "ts": 326310414451.590, "ph": "B", "pid": 11326, "name": "fib" },
  { "ts": 326310414451.632, "ph": "B", "pid": 11326, "name": "fib" },
  { "ts": 326310414451.719, "ph": "E", "pid": 11326, "name": "fib" },
  { "ts": 326310414451.956, "ph": "B", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.008, "ph": "E", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.123, "ph": "E", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.227, "ph": "B", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.275, "ph": "E", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.355, "ph": "E", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.425, "ph": "B", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.484, "ph": "B", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.536, "ph": "E", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.673, "ph": "B", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.721, "ph": "E", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.778, "ph": "E", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.855, "ph": "E", "pid": 11326, "name": "fib" },
  { "ts": 326310414452.945, "ph": "E", "pid": 11326, "name": "main" }
], "displayTimeUnit": "ns", "metadata": {
  "command_line": "ufttrace record ./fibonacci 5 ",
  "recorded_time": "Thu Oct 19 06:03:54 2017"
} }
```

ufttrace dump



.json
(trace data)

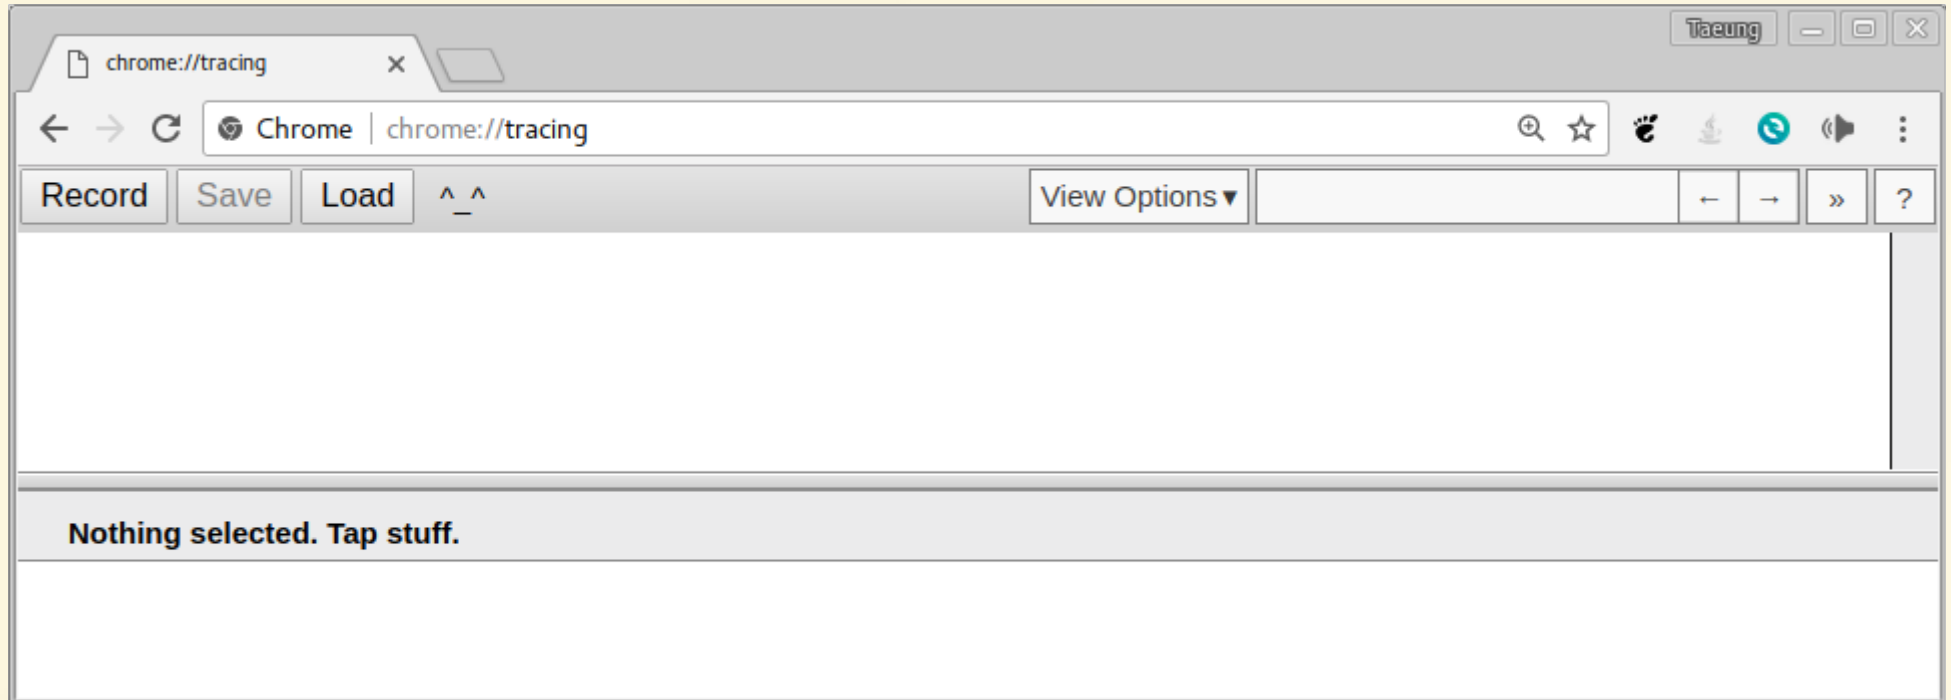
```
$ uftrace dump --chrome > fibonacci.json
```

```
$ google-chrome
```

```
$ uftrace dump --chrome > fibonacci.json
```

```
$ google-chrome
```

chrome://tracing →



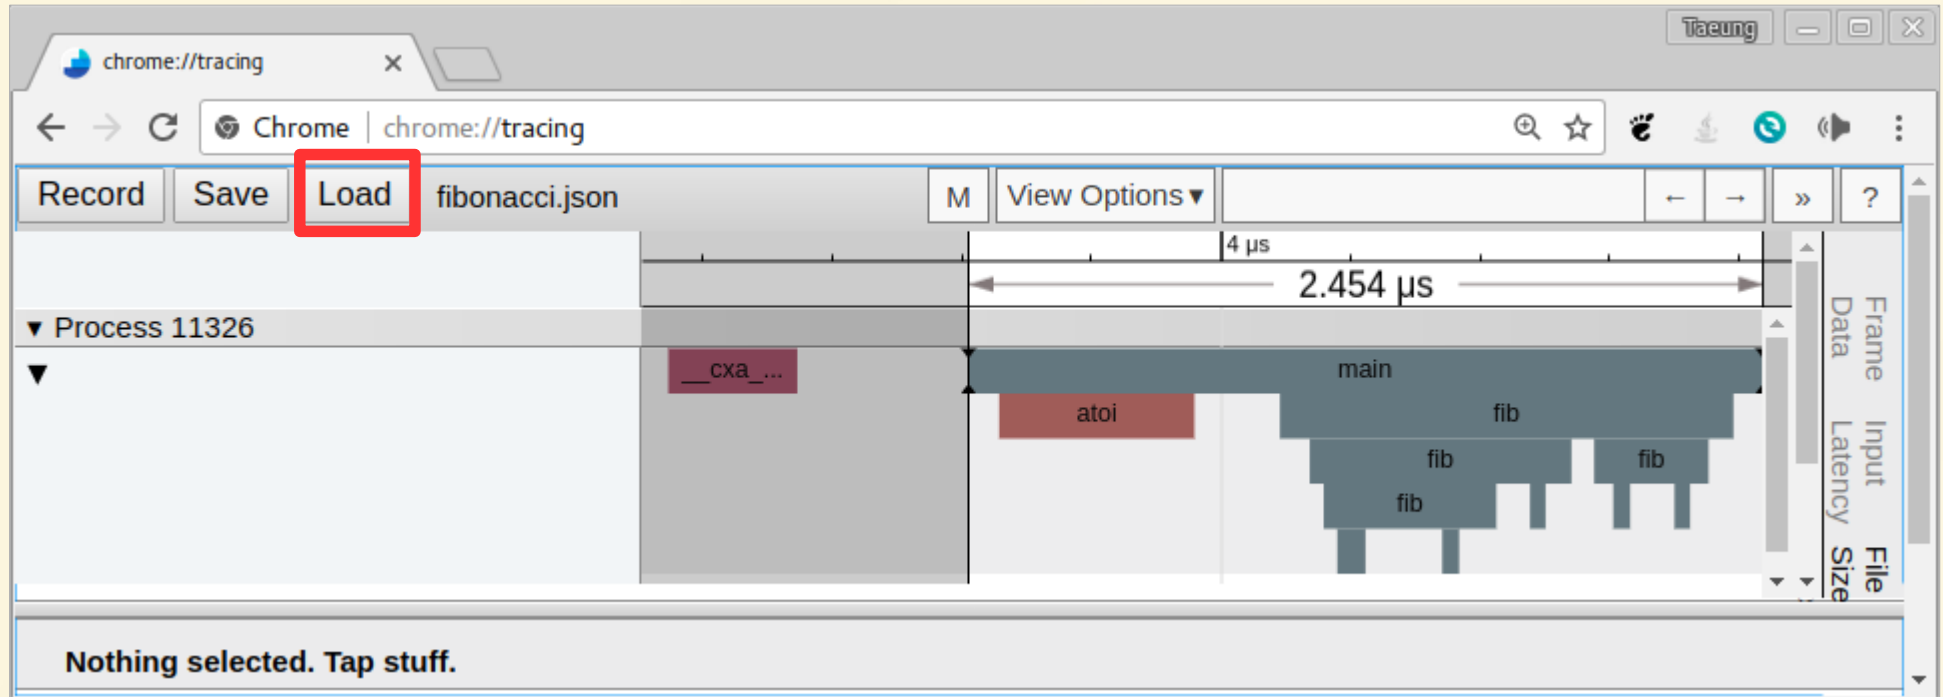
```
$ uftrace dump --chrome > fibonacci.json
```

```
$ google-chrome
```

```
$ uftrace dump --chrome > fibonacci.json
```

```
$ google-chrome
```

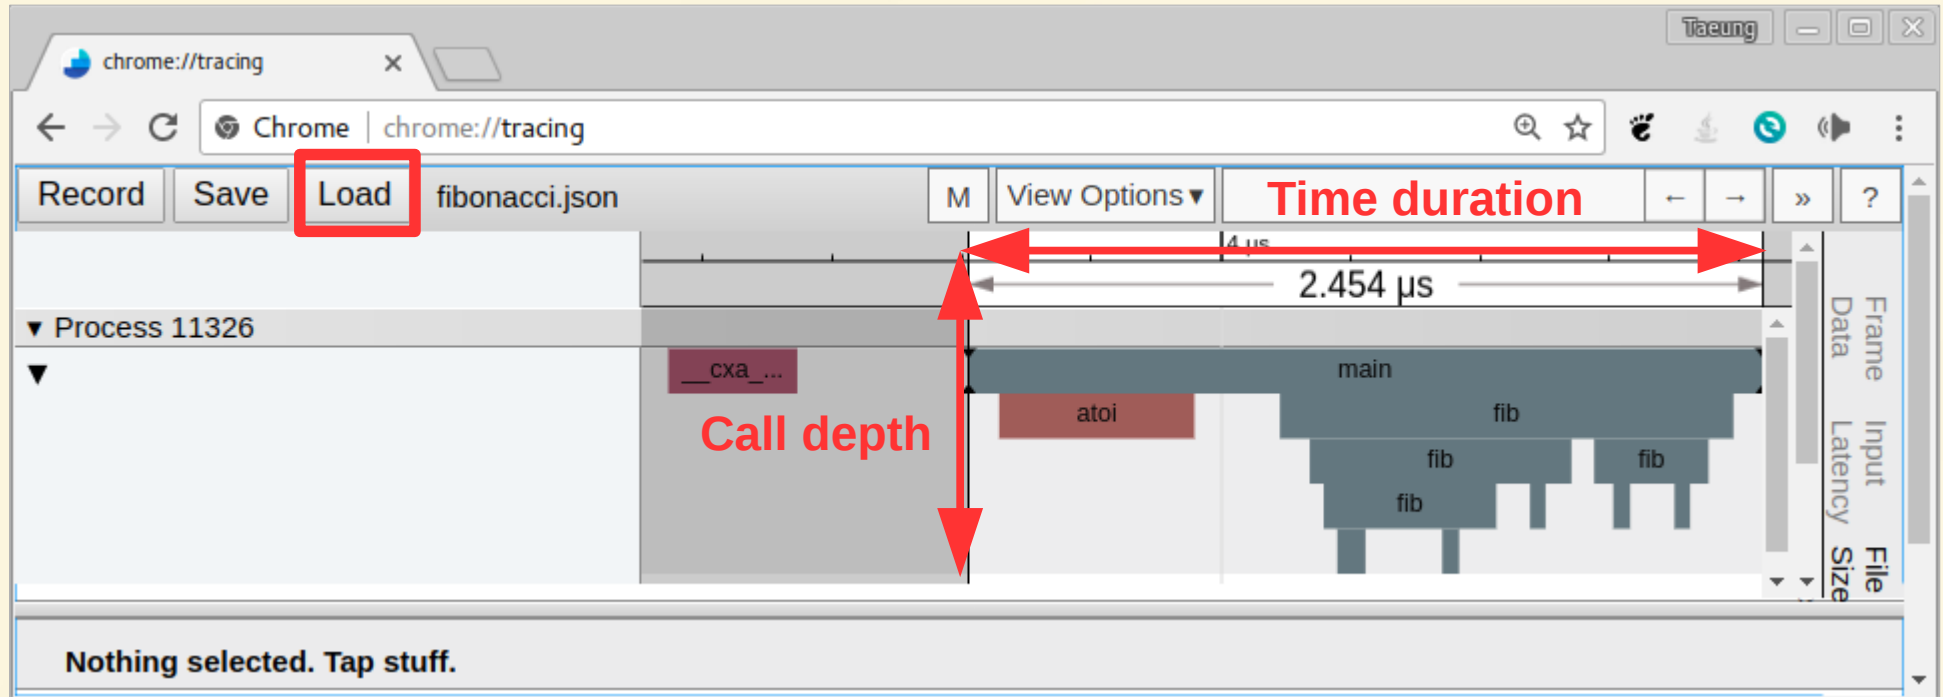
chrome://tracing → Load → fibonacci.json → Function call graph




```
$ uftrace dump --chrome > fibonacci.json
```

```
$ google-chrome
```

chrome://tracing → Load → **fibonacci.json** → Function **call graph**



```
$ uftrace dump --chrome > fibonacci.json
```

```
$ google-chrome
```

chrome://tracing → Load → **fibonacci.json** → Function call graph

The screenshot shows the Chrome DevTools Tracing interface. The top bar includes buttons for 'Record', 'Save', 'Load', and 'fibonacci.json'. A red box highlights the 'M' button, which is the 'Metadata' button. A red arrow points from this button to a dialog box titled 'Metadata for trace'. The dialog box contains a table with the following data:

name	value
metadata	{command_line: "uftrace record ./fibonacci 5 ", recorded_time: "Thu Oct 19 06:03:54 2017"}

The bottom of the interface shows the text 'Nothing selected. Tap stuff.'

uftrace can trace **User** + **Lib** + **Kernel**

showing the execution flow

uftrace = **user** function trace
+ **strace** + **ltrace** + **ftrace**

based on the execution flow

```
$ uftrace record hello
Hello OSSEU17 !!
```

```
$ uftrace replay
```

#	DURATION	TID	FUNCTION
	0.710 us	[13588]	__monstartup();
	0.713 us	[13588]	__cxa_atexit();
		[13588]	main() {
	4.107 us	[13588]	printf();
	4.046 us	[13588]	fflush();
	8.815 us	[13588]	} /* main */

user space

main()


libraries

printf()
fflush()

```
$ uftrace record hello
Hello OSSEU17 !!
```

```
$ uftrace replay
```

#	DURATION	TID	FUNCTION
	0.710 us	[13588]	__monstartup();
	0.713 us	[13588]	__cxa_atexit();
		[13588]	main() {
	4.107 us	[13588]	printf();
	4.046 us	[13588]	fflush();
	8.815 us	[13588]	} /* main */




```
$ uftrace record -k hello
Hello OSSEU17 !!
```

```
$ uftrace replay
```

#	DURATION	TID	FUNCTION
	1.060 us	[13565]	__monstartup();
	1.113 us	[13565]	__cxa_atexit();
		[13565]	main() {
		[13565]	printf() {
	3.173 us	[13565]	sys_newfstat();
	6.107 us	[13565]	__do_page_fault();
	17.713 us	[13565]	} /* printf */
		[13565]	fflush() {
	7.198 us	[13565]	sys_write();
	12.270 us	[13565]	} /* fflush */
	30.661 us	[13565]	} /* main */


Integrated tracer !




```
$ uftrace record -K 2 hello
Hello OSSEU17 !!
```

```
$ uftrace replay
```

#	DURATION	TID	FUNCTION
0.887	us	[17099]	__monstartup();
1.039	us	[17099]	__cxa_atexit();
		[17099]	main() {
		[17099]	printf() {
		[17099]	sys_newfstat() {
1.378	us	[17099]	vfs_fstat();
0.561	us	[17099]	cp_new_stat();
4.197	us	[17099]	} /* sys_newfstat */
		[17099]	__do_page_fault() {
0.233	us	[17099]	down_read_trylock();
0.239	us	[17099]	_cond_resched();
1.839	us	[17099]	find_vma();
6.288	us	[17099]	handle_mm_fault();
0.249	us	[17099]	up_read();
10.514	us	[17099]	} /* __do_page_fault */
22.183	us	[17099]	} /* printf */
		[17099]	fflush() {
		[17099]	sys_write() {
0.772	us	[17099]	__fdget_pos();
7.731	us	[17099]	vfs_write();
11.184	us	[17099]	} /* sys_write */
15.619	us	[17099]	} /* fflush */
38.504	us	[17099]	} /* main */





uftrace

Features

- C/C++ (**user** space) functions
 - compiled with **-pg** or -finstrument-functions
- **Library** functions
- Linux **kernel** functions
- Some of system events
 - schedule in-out
 - page faults

uftrace Installation

```
# Ubuntu
```

```
$ sudo apt-get install libelf-dev
```

```
# Fedora, RHEL
```

```
$ sudo dnf install elfutils-libelf-devel
```

DEPENDENCY

Need to install **libelf**

```
# Ubuntu
$ sudo apt-get install libelf-dev

# Fedora, RHEL
$ sudo dnf install elfutils-libelf-devel

$ git clone https://github.com/namhyung/uftrace.git

$ cd uftrace
```

GETTING THE SOURCE

Need to git-clone

```
# Ubuntu
$ sudo apt-get install libelf-dev

# Fedora, RHEL
$ sudo dnf install elfutils-libelf-devel

$ git clone https://github.com/namhyung/uftrace.git

$ cd uftrace

$ ./configure

$ make

$ sudo make install
```

BUILD

Need to **make**

```
$ cd uftrace
```

```
# You can check many test source files
```

```
$ cd tests/ && ls s-*
```

TEST

Various test cases

```
$ cd uftrace
```

```
# You can check many test source files
```

```
$ cd tests/ && ls s-*
```

```
# For example, you can run the 090 test case
```

```
$ ./runtest.py 090
```

```
# Or, you can run it by yourself
```

```
$ gcc -pg -o fib s-fibonacci.c
```

```
$ uftrace ./fib
```

TEST

Various test cases

Live demo

fibonacci example

Live demo

fibonacci example

Step by step, together !!

uftrace tutorials

<https://github.com/taeung/uftrace-osseu17>

<https://github.com/taeung/uftrace-osseu17>

Get uftrace **tutorial** examples

<https://github.com/taeung/uftrace-osseu17>

Get uftrace **tutorial** examples



recorded data by uftrace

Reproduce execution flow of **examples**,

Reproduce execution flow of **examples**,
with **recorded data** (e.g. uftrace.data)


```
$ git clone https://github.com/taeung/uftrace-osseu17.git
```

```
$ cd uftrace-osseu17/
```

```
$ ls  
chrome_tracing_examples  
cpython_example  
Good_bye_printf_hands_on_tutorial_uftrace<...>.pdf  
nmap_examples  
nullptr_exception_example  
optimization_level_examples  
printf_kern_examples  
process_life_cycle_example
```

git-clone

uftrace examples !!

Nmap examples

What is **Nmap** ?

The Matrix Reloaded (2003)

Trinity uses Nmap !



```
80/tcp    open      http
81/tcp    open      hosts2.nc
10.0.0.0  [mobile]
11 # nmap -v -sS -O 10.2.2.2
11
13 Starting nmap V. 2.54BETA25
13 Insufficient responses for TCP sequencing (3), OS detection
13 accurate
14 Interesting ports on 10.2.2.2:
44 (The 1539 ports scanned but not shown below are in state: closed)
51 Port      State      Service
51 22/tcp    open      ssh
58
60 No exact OS matches for host
60
24 Nmap run completed -- 1 IP address (1 host up) scanned
50 # sshnuke 10.2.2.2 -rootpw="Z10M0101"
Connecting to 10.2.2.2:ssh ... successful.
Re Attempting to exploit SSHv1 CRC32 ... successful.
IP Resetting root password to "Z10M0101".
System open: Access Level <9>
Hn # ssh 10.2.2.2 -l root
root@10.2.2.2's password: [REDACTED]
```

ACCESS GRANTED

<https://nmap.org/movies> "Trinity hacking scene"

Let's trace Nmap !

```
$ git clone https://github.com/nmap/nmap.git
```

```
$ cd nmap
```

```
$ ./configure
```

```
$ make
```

GETTING SRC & BUILD

But you don't need to do
because of example files

```
$ ./nmap nmap.org
```

```
Starting Nmap 7.60SVN ( https://nmap.org ) at 2017-10-19 ...  
Nmap scan report for nmap.org (45.33.49.119)  
Host is up (0.18s latency).  
Other addresses for nmap.org (not scanned):  
2600:3c01::f03c:91ff:fe98:ff4e  
rDNS record for 45.33.49.119: ack.nmap.org  
Not shown: 993 filtered ports  


| PORT      | STATE  | SERVICE |
|-----------|--------|---------|
| 22/tcp    | open   | ssh     |
| 25/tcp    | open   | smtp    |
| 70/tcp    | closed | gopher  |
| 80/tcp    | open   | http    |
| 113/tcp   | closed | ident   |
| 443/tcp   | open   | https   |
| 31337/tcp | closed | Elite   |

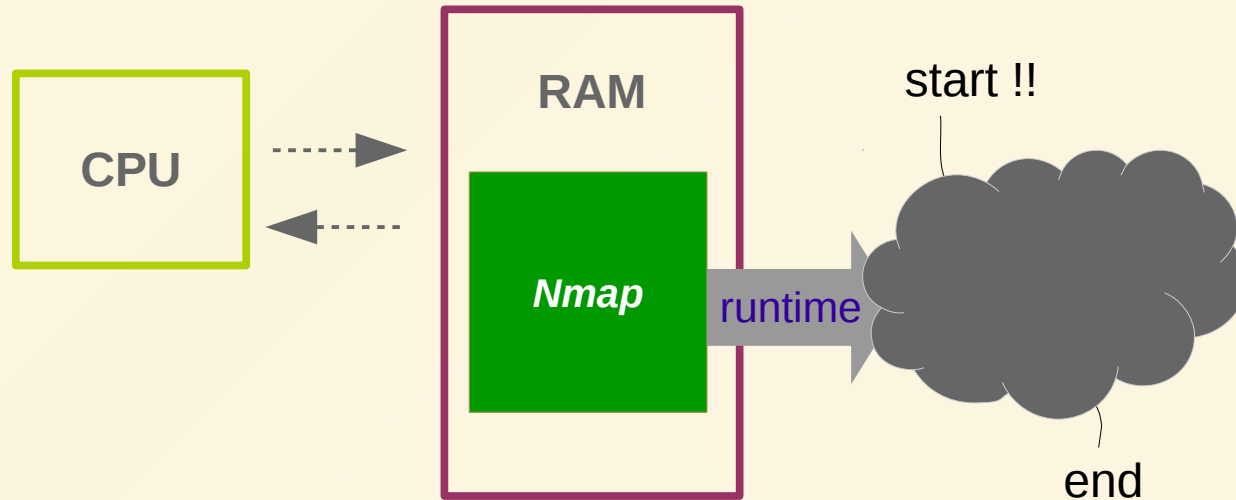
  
Nmap done: 1 IP address (1 host up) scanned in 12.13 seconds
```

SIMPLE TEST

Nmap can **port scan**

<https://nmap.org>

How does Nmap scan port numbers?

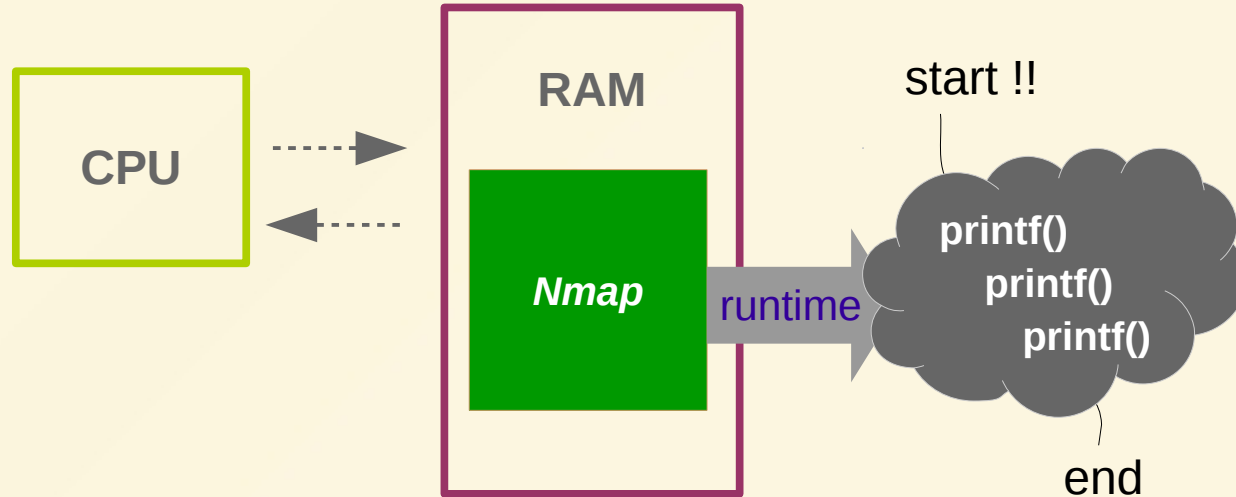


Hum..

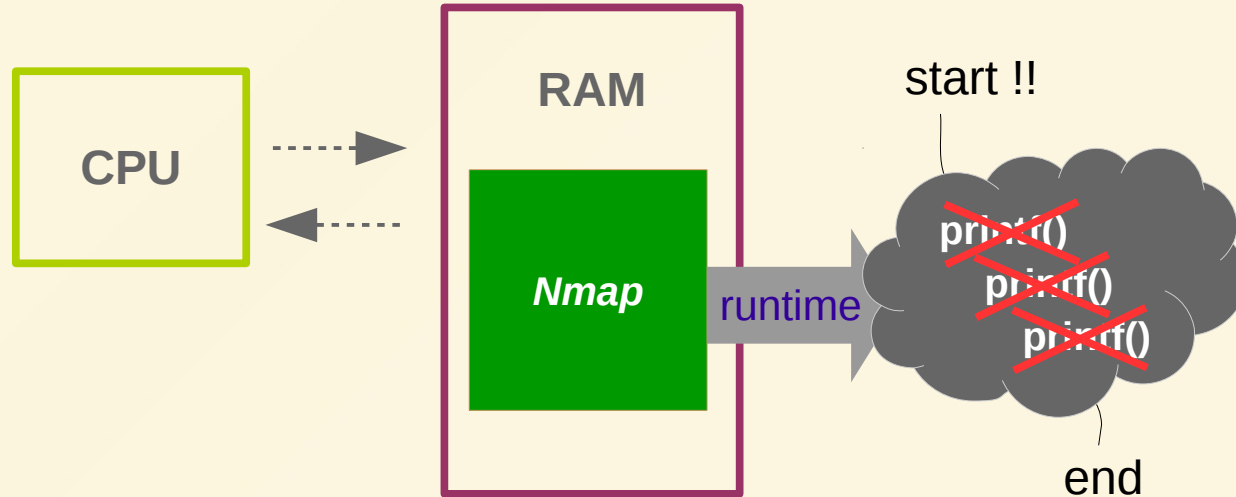
Nmap scans ports **in order** ?

0, 1, 2, ... 65535 (16 bit unsigned int)

To **investigate** it, we can read Nmap src
and dig using **printf()** !



You can use **uftrace** instead



You can **more efficiently** analyze it

You can **more efficiently** analyze it
using **uftrace** !!

Let's do it ! step by step

Step 1. Recompile Nmap

with -pg or -finstrument-functions

```
$ cd nmap/

$ emacs Makefile

$ diff Makefile.old Makefile
--- Makefile.old 2017-10-21 04:02:58.016974789 +0900
+++ Makefile      2017-10-21 04:03:10.861083599 +0900
@@ -47,9 +47,9 @@ # Level 1 only makes changes that don't
# For mtrace debugging -- see MTRACE define in main.cc for
instructions
# Should only be enabled during debugging and not in any real
release.
# DEFS += -DMTRACE=1
-CXXFLAGS = -g -O2 -Wall -fno-strict-aliasing $(DBGFLAGS) $(CCOPT)
+CXXFLAGS = -pg -g -O2 -Wall -fno-strict-aliasing $(DBGFLAGS) $(CCOPT)
CPPFLAGS = -I$(top_srcdir)/liblinear -I$(top_srcdir)/liblua -I$
(top_srcdir)/libdnet-stripped/include -I$(top_srcdir)/libssh2/include
-I$(top_srcdir)/libpcap -I$(top_srcdir)/nbase -I$
(top_srcdir)/nsock/include $(DEFS)
-CFLAGS = -g -O2 -Wall $(DBGFLAGS) $(CCOPT)
+CFLAGS = -pg -g -O2 -Wall $(DBGFLAGS) $(CCOPT)
STATIC =
LDFLAGS = -Wl,-E -lnbase -lnsock/src/ $(DBGFLAGS) $(STATIC)
LIBS = -lnsock -lnbase -lpcap $(LIBPCAPDIR)/libpcap.a $
(LIBSSH2_LIBS) $(OPENSSL_LIBS) $(ZLIB_LIBS) libnetutil/libnetutil.a $
(top_srcdir)/libdnet-stripped/src/.libs/libdnet.a $
(top_srcdir)/liblua/liblua.a $(top_srcdir)/liblinear/liblinear.a -ldl

$ make
```

RECOMPILE

with **-pg** or

-finstrument-functions

Step 2. Nmap record & replay

in 2 ~3 depth level

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```



Just use recorded data !!

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace info
```

```
# system information
```

```
# =====
```

```
# program version      : uftrace v0.8-59-g7181
```

```
# recorded on          : Fri Oct 20 13:20:03 2017
```

```
# cmdline               : uftrace record -t 5us -d uftrace.data.nmap_nmap.org ./nmap nmap.org
```

```
# cpu info              : Intel(R) Core(TM) i7-5500U CPU @ 2.40GHz
```

```
# number of cpus        : 4 / 4 (online / possible)
```

```
# memory info           : 0.1 / 7.4 GB (free / total)
```

```
# system load           : 0.12 / 0.24 / 0.34 (1 / 5 / 15 min)
```

```
# kernel version        : Linux 4.5.0-rc4+
```

```
# hostname               : taeung-ThinkPad-X1-Carbon-3rd
```

```
# distro                 : "Ubuntu 16.04.3 LTS"
```

```
#
```

```
# process information
```

```
# =====
```

```
# number of tasks       : 1
```

```
# task list             : 7673
```

```
# exe image             : /home/taeung/git/opensource/nmap/nmap
```

```
...
```

```
info - read Info documents
```

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace info
```

Check record cmdline !



```
# system information
```

```
# =====
```

```
# program version      : uftrace v0.8-59-g7181
```

```
# recorded on          : Fri Oct 20 12:20:02 2017
```

```
# cmdline              : uftrace record -t 5us -d uftrace.data.nmap_nmap.org ./nmap nmap.org
```

```
# cpu info             : Intel(R) Core(TM) i7-5500U CPU @ 2.40GHz
```

```
# number of cpus       : 4 / 4 (online / possible)
```

```
# memory info          : 0.1 / 7.4 GB (free / total)
```

```
# system load          : 0.12 / 0.24 / 0.34 (1 / 5 / 15 min)
```

```
# kernel version       : Linux 4.5.0-rc4+
```

```
# hostname             : taeung-ThinkPad-X1-Carbon-3rd
```

```
# distro               : "Ubuntu 16.04.3 LTS"
```

```
#
```

```
# process information
```

```
# =====
```

```
# number of tasks      : 1
```

```
# task list            : 7673
```

```
# exe image            : /home/taeung/git/opensource/nmap/nmap
```

```
...
```

info - read Info documents

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace replay -D 2
```

```
# DURATION      TID      FUNCTION
    7.245 us [ 7673] | _GLOBAL__sub_I__ZN3DNS7Factory13progressiveIdE() {
   21.938 us [ 7673] |   open();
   59.669 us [ 7673] |   read();
   6.125 us [ 7673] | } /* _GLOBAL__sub_I__ZN3DNS7Factory13progressiveIdE */
    5.718 us [ 7673] | _GLOBAL__sub_I__Z16set_program_namePKc();
    5.718 us [ 7673] | _GLOBAL__sub_I_o() {
    5.718 us [ 7673] |   NmapOps::NmapOps();
...
    14.193 s [ 7673] | main() {
    14.193 s [ 7673] |   nmap_main();
    14.193 s [ 7673] | } /* main */
...
```

```
-D, --depth=DEPTH
```

```
Trace functions within DEPTH
```

```
-F, --filter=FUNC
```

```
Only trace those FUNCs
```

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace replay -D 2
```

```
# DURATION      TID      FUNCTION
    7.245 us [ 7673] | _GLOBAL__sub_I__ZN3DNS7Factory13progressiveIdE() {
   21.938 us [ 7673] |   open();
   59.669 us [ 7673] |   read();
    6.125 us [ 7673] | } /* _GLOBAL__sub_I__ZN3DNS7Factory13progressiveIdE */
    5.718 us [ 7673] | _GLOBAL__sub_I_Z16set_program_namePKc();
    ...
    14.193 s [ 7673] | main() {
    14.193 s [ 7673] |   nmap_main();
    14.193 s [ 7673] | } /* main */
    ...
```

Not important !

```
-D, --depth=DEPTH
```

Trace functions within DEPTH

```
-F, --filter=FUNC
```

Only trace those FUNCS

```
$ uftrace replay -D 3 -F main
```

```
# DURATION      TID      FUNCTION
    [ 7673] | main() {
    [ 7673] |   nmap_main() {
36.916 us [ 7673] |     localtime();
15.758 us [ 7673] |     parse_options();
48.699 us [ 7673] |     tty_init();
143.925 ms [ 7673] |     apply_delayed_options();
19.436 us [ 7673] |     ctime();
106.828 us [ 7673] |     NmapOps::XSLStyleSheet();
 7.032 us [ 7673] |     xml_start_document();
 6.040 us [ 7673] |     xml_attribute();
10.399 us [ 7673] |     xml_write_escaped();
206.983 us [ 7673] |     output_xml_scaninfo_records();
...

```

```
-D, --depth=DEPTH
```

Trace functions within DEPTH

```
-F, --filter=FUNC
```

Only trace those FUNCS

Replay in **3 depth level**

with example data

But **too many** functions..

Step 3. **Statistics** and **summary**

```
$ man uftace report
```

NAME

uftace-report - Print **statistics** and **summary** for trace data

SYNOPSIS

uftace report [options]

DESCRIPTION

This command collects trace data from a given data file and prints statistics and summary information. It shows function statistics by default, but can show thread statistics with the `--threads` option and show differences between traces with the `--diff` option.

...

\$ uftrace **report**

Total time	Self time	Calls	Function
=====	=====	=====	=====
14.193 s	16.465 us	1	main
14.193 s	80.617 us	1	nmap_main
13.373 s	165.279 ms	2	ultra_scan
13.151 s	3.474 ms	1172	do_one_select_round
13.147 s	13.147 s	814	select
883.098 ms	0.501 us	1	nexthost
883.097 ms	15.743 us	1	refresh_hostbatch
381.225 ms	0.714 us	1	nmap_mass_rdns
381.225 ms	131.629 us	1	nmap_mass_rdns_core
379.970 ms	379.970 ms	1	epoll_wait
292.927 ms	2.250 us	1	TargetGroup::get_next_host
292.925 ms	13.247 us	1	NetBlockHostname::resolve
292.911 ms	1.278 us	1	resolve_all
292.910 ms	292.910 ms	1	getaddrinfo
143.925 ms	27.309 us	1	apply_delayed_options
143.661 ms	7.293 ms	1	gettoppts
135.620 ms	87.742 ms	1	nmap_services_init
46.393 ms	21.499 ms	2	std::__cxx11::list::sort
44.670 ms	18.142 ms	2002	sendConnectScanProbe
25.039 ms	25.039 ms	2002	connect

...

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
13.147 s	13.147 s	814	select
379.970 ms	379.970 ms	1	epoll_wait
292.910 ms	292.910 ms	1	getaddrinfo
13.373 s	165.279 ms	2	ultra_scan
135.620 ms	87.742 ms	1	nmap_services_init
25.039 ms	25.039 ms	2002	connect
24.875 ms	24.761 ms	719	std::__cxx11::list::merge
46.393 ms	21.499 ms	2	std::__cxx11::list::sort
44.670 ms	18.142 ms	2002	sendConnectScanProbe
143.661 ms	7.293 ms	1	gettoppts
5.533 ms	5.533 ms	865	std::_Rb_tree::_M_erase
3.727 ms	3.727 ms	318	close
13.151 s	3.474 ms	1172	do_one_select_round
3.099 ms	3.099 ms	1	std::__cxx11::list::~~list
5.880 ms	2.392 ms	607	HostScanStats::markProbeTimedout
1.851 ms	1.366 ms	197	keyWasPressed
1.346 ms	1.346 ms	196	socket
1.260 ms	1.187 ms	145	next_token
942.693 us	942.693 us	1	std::ios_base::ios_base
2.117 ms	857.682 us	1	load_payloads_from_file

```
...
```

```
-s, --sort=KEY[,KEY,...] Sort reported functions by KEYS
```

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
13.147 s	13.147 s	814	select
379.970 ms	379.970 ms	1	epoll_wait
292.910 ms	292.910 ms	1	getaddrinfo
13.373 s	165.279 ms	2	ultra_scan
135.620 ms	87.742 ms	1	nmap_services_init
25.039 ms	25.039 ms	2002	connect
24.875 ms	24.761 ms	719	std::__cxx11::list::merge
46.393 ms	21.499 ms	2	std::__cxx11::list::sort
44.670 ms	18.142 ms	2002	sendConnectScanProbe
143.661 ms	7.293 ms	1	gettoppts
5.533 ms	5.533 ms	865	std::_Rb_tree::_M_erase
3.727 ms	3.727 ms	318	close
13.151 s	3.474 ms	1172	do_one_select_round
3.099 ms	3.099 ms	1	std::__cxx11::list::~~list
5.880 ms	2.392 ms	607	HostScanStats::markProbeTimedout
1.851 ms	1.366 ms	197	keyWasPressed
1.346 ms	1.346 ms	196	socket
1.260 ms	1.187 ms	145	next_token
942.693 us	942.693 us	1	std::ios_base::ios_base
2.117 ms	857.682 us	1	load_payloads_from_file

...

-s, --sort=KEY[,KEY,...] Sort reported functions by KEYS

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
13.147 s	13.147 s	814	select
379.970 ms	379.970 ms	1	epoll_wait
292.910 ms	292.910 ms	1	getaddrinfo
13.373 s	165.279 ms	2	ultra_scan ←
135.620 ms	87.742 ms	1	nmap_services_init
25.039 ms	25.039 ms	2002	connect
24.875 ms	24.761 ms	719	std::__cxx11::list::merge
46.393 ms	21.499 ms	2	std::__cxx11::list::sort
44.670 ms	18.142 ms	2002	sendConnectScanProbe ←
143.661 ms	7.293 ms	1	gettoppts
5.533 ms	5.533 ms	865	std::_Rb_tree::_M_erase
3.727 ms	3.727 ms	318	close
13.151 s	3.474 ms	1172	do_one_select_round
3.099 ms	3.099 ms	1	std::__cxx11::list::~~list
5.880 ms	2.392 ms	607	HostScanStats::markProbeTimeout
1.851 ms	1.366 ms	197	keyWasPressed
1.346 ms	1.346 ms	196	socket
1.260 ms	1.187 ms	145	next_token
942.693 us	942.693 us	1	std::ios_base::ios_base
2.117 ms	857.682 us	1	load_payloads_from_file

```
...
```

```
-s, --sort=KEY[,KEY,...] Sort reported functions by KEYS
```

Focus on funcs related to 'scan'

ultra_scan() first !


```
$ man uftrace graph
```

NAME

uftrace-graph - Show function call graph

SYNOPSIS

uftrace graph [options] []

DESCRIPTION

This command shows a function call graph for the given function in a uftrace record datafile. If the function name is omitted, main is used by default. The function call graph contains backtrace and calling functions. Each function in the output is annotated with a hit count and the total time spent running that function.

...

```
$ uftrace graph ultra_scan
```

```
...
```

```
calling functions
```

```
=====
```

```
13.373 s : (2) ultra_scan
  2.219 ms :  +--(1) init_payloads
 92.412 us :  |  +--(1) nmap_fetchfile
 90.022 us :  |  |  (2) nmap_fetchfile_sub
 53.572 us :  |  |  +--(2) nmap_fetchfile_userdir_uid
 47.017 us :  |  |  |  (2) getpwuid
           :  |  |  |
  8.384 us :  |  |  +--(1) readlink
```

```
...
```

```
           :  |
21.104 us :  +--(2) UltraScanInfo::Init
  5.336 us :  |  (1) GroupScanStats::GroupScanStats
           :  |
44.155 ms :  +--(1995) sendConnectScanProbe
24.683 ms :  |  +--(1995) connect
           :  |  |
  1.292 ms :  |  +--(189) socket
```

```
...
```

```
$ uftrace graph ultra_scan
```

```
...
```

```
calling functions
```

```
=====
```

```
13.373 s : (2) ultra_scan
  2.219 ms : +- (1) init_payloads
 92.412 us : | +- (1) nmap_fetchfile
 90.022 us : | | (2) nmap_fetchfile_sub
 53.572 us : | | +- (2) nmap_fetchfile_userdir_uid
 47.017 us : | | | (2) getpwuid
          : | | |
  8.384 us : | | +- (1) readlink
```

```
...
```

```
          : |
21.104 us : +- (2) UltraScanInfo::Init
  5.336 us : | (1) GroupScanStats::GroupScanStats
          : |
44.155 ms : +- (1995) sendConnectScanProbe ←
24.683 ms : | +- (1995) connect
          : | |
  1.292 ms : | +- (189) socket
```

```
...
```

Oh, too many call counts

```
$ uftrace graph ultra_scan
```

```
...
```

```
calling functions
```

```
=====
```

```
13.373 s : (2) ultra_scan
  2.219 ms : +- (1) init_payloads
 92.412 us : | +- (1) nmap_fetchfile
 90.022 us : | | (2) nmap_fetchfile_sub
 53.572 us : | | +- (2) nmap_fetchfile_userdir_uid
 47.017 us : | | | (2) getpwuid
          : | | |
  8.384 us : | | +- (1) readlink
```

```
...
```

```
          : |
21.104 us : +- (2) UltraScanInfo::Init
  5.336 us : | (1) GroupScanStats::GroupScanStats
          : |
44.155 ms : +- (1995) sendConnectScanProbe
24.683 ms : | +- (1995) connect ←
          : | |
  1.292 ms : | +- (189) socket ←
```

```
...
```

Moreover nested **connect()**, **socket()**

Looking into

sendConnectScanProbe() !

```
$ cd nmap/
```

```
$ grep -n sendConnectScanProbe *.cc
```

```
scan_engine.cc:2210:    sendConnectScanProbe(USI, hss, pspec.pd.tcp.dport, 0, 0);  
scan_engine.cc:2232:    sendConnectScanProbe(USI, hss, pspec.pd.tcp.dport, pspec_tries + 1, 0);  
scan_engine.cc:2295:    sendConnectScanProbe(USI, hss, hss->target->pingprobe.pd.tcp.dport, 0,  
scan_engine.cc:2380:    newProbe = sendConnectScanProbe(USI, hss, probe->pspec()->pd.tcp.dport, probe-  
>tryno + 1, 0);  
scan_engine_connect.cc:424:UltraProbe *sendConnectScanProbe(UltraScanInfo *USI, HostScanStats *hss,
```



Ah, it is **function header**

```
$ cd nmap/
```

```
$ grep -n sendConnectScanProbe *.cc
```

```
scan_engine.cc:2210:    sendConnectScanProbe(USI, hss, pspec.pd.tcp.dport, 0, 0);  
scan_engine.cc:2232:    sendConnectScanProbe(USI, hss, pspec.pd.tcp.dport, pspec_tries + 1, 0);  
scan_engine.cc:2295:    sendConnectScanProbe(USI, hss, hss->target->pingprobe.pd.tcp.dport, 0,  
scan_engine.cc:2380:    newProbe = sendConnectScanProbe(USI, hss, probe->pspec()->pd.tcp.dport, probe-  
>tryno + 1, 0);  
scan_engine_connect.cc:424:UltraProbe *sendConnectScanProbe(UltraScanInfo *USI, HostScanStats *hss,
```

```
$ emacs scan_engine_connect.cc
```

```
...  
422 /* If this is NOT a ping probe, set pingseq to 0. Otherwise it will be the  
423    ping sequence number (they start at 1). The probe sent is returned. */  
424 UltraProbe *sendConnectScanProbe(UltraScanInfo *USI, HostScanStats *hss,  
425                                   u16 destport, u8 tryno, u8 pingseq) {  
...
```



Oh, is it **port number** ?

Step 4. Print **Arguments**


```
424 UltraProbe *sendConnectScanProbe(UltraScanInfo *USI, HostScanStats *hss,  
425                                     u16 destport, u8 tryno, u8 pingseq) {  
...
```

```
$ cd ../ && cd uftrace.data.nmap_args/
```

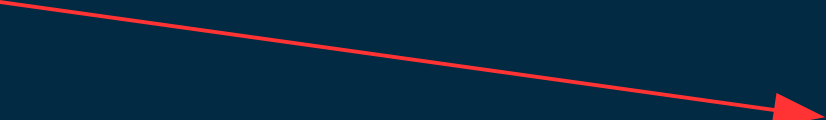
```
$ uftrace replay -F sendConnectScanProbe -D 1 -A sendConnectScanProbe@arg3
```

#	DURATION	TID	FUNCTION
	87.474 us	[31689]	sendConnectScanProbe(80);
	34.240 us	[31689]	sendConnectScanProbe(443);
	75.963 us	[31689]	sendConnectScanProbe(1723);
	35.326 us	[31689]	sendConnectScanProbe(199);
	20.016 us	[31689]	sendConnectScanProbe(25);
	17.195 us	[31689]	sendConnectScanProbe(113);
	20.688 us	[31689]	sendConnectScanProbe(143);
	17.541 us	[31689]	sendConnectScanProbe(3306);
	16.678 us	[31689]	sendConnectScanProbe(53);
	16.493 us	[31689]	sendConnectScanProbe(22);
	15.893 us	[31689]	sendConnectScanProbe(111);
	...		

```
424 UltraProbe *sendConnectScanProbe(UltraScanInfo *USI, HostScanStats *hss,  
425                                     u16 destport, u8 tryno, u8 pingseq) {  
...
```

```
$ cd ../ && cd uftrace.data.nmap_args/
```

```
$ uftrace replay -F sendConnectScanProbe -D 1 -A sendConnectScanProbe@arg3
```



#	DURATION	TID	FUNCTION
	87.474 us	[31689]	sendConnectScanProbe(80);
	34.240 us	[31689]	sendConnectScanProbe(443);
	75.963 us	[31689]	sendConnectScanProbe(1723);
	35.326 us	[31689]	sendConnectScanProbe(199);
	20.016 us	[31689]	sendConnectScanProbe(25);
	17.195 us	[31689]	sendConnectScanProbe(113);
	20.688 us	[31689]	sendConnectScanProbe(143);
	17.541 us	[31689]	sendConnectScanProbe(3306);
	16.678 us	[31689]	sendConnectScanProbe(53);
	16.493 us	[31689]	sendConnectScanProbe(22);
	15.893 us	[31689]	sendConnectScanProbe(111);
	...		

```
424 UltraProbe *sendConnectScanProbe(UltraScanInfo *USI, HostScanStats *hss,  
425                                     u16 destport, u8 tryno, u8 pingseq) {  
...
```

```
$ cd ../ && cd uftrace.data.nmap_args/
```

```
$ uftrace replay -F sendConnectScanProbe -D 1 -A sendConnectScanProbe@arg3
```

#	DURATION	TID	FUNCTION
87.474	us	[31689]	sendConnectScanProbe(80); → http
34.240	us	[31689]	sendConnectScanProbe(443);
75.963	us	[31689]	sendConnectScanProbe(1723);
35.326	us	[31689]	sendConnectScanProbe(199);
20.016	us	[31689]	sendConnectScanProbe(25);
17.195	us	[31689]	sendConnectScanProbe(113);
20.688	us	[31689]	sendConnectScanProbe(143);
17.541	us	[31689]	sendConnectScanProbe(3306);
16.678	us	[31689]	sendConnectScanProbe(53);
16.493	us	[31689]	sendConnectScanProbe(22); → ssh
15.893	us	[31689]	sendConnectScanProbe(111);
...			

```
424 UltraProbe *sendConnectScanProbe(UltraScanInfo *USI, HostScanStats *hss,  
425                                     u16 destport, u8 tryno, u8 pingseq) {  
...
```

```
$ cd ../ && cd uftrace.data.nmap_args/
```

```
$ uftrace replay -F sendConnectScanProbe -D 1 -A sendConnectScanProbe@arg3
```

#	DURATION	TID	FUNCTION
87.474	us	[31689]	sendConnectScanProbe(80);
34.240	us	[31689]	sendConnectScanProbe(443);
75.963	us	[31689]	sendConnectScanProbe(1723);
35.326	us	[31689]	sendConnectScanProbe(199);
20.016	us	[31689]	sendConnectScanProbe(25);
17.195	us	[31689]	sendConnectScanProbe(113);
20.688	us	[31689]	sendConnectScanProbe(143);
17.541	us	[31689]	sendConnectScanProbe(3306);
16.678	us	[31689]	sendConnectScanProbe(53);
16.493	us	[31689]	sendConnectScanProbe(22);
15.893	us	[31689]	sendConnectScanProbe(111);
...			

This is a port scan sequence

Nmap scans well-known port numbers first

Event Tracing (sched event)

-E linux:schedule

```
$ uftrace t-fork
```

```
# DURATION      TID      FUNCTION
    1.609 us [32272] | __monstartup();
    1.493 us [32272] | __cxa_atexit();
    189.898 us [32272] | main() {
    189.898 us [32272] |     fork();
    189.898 us [32272] |     wait() {
    189.898 us [32275] |         } /* fork */
    189.898 us [32275] |     a() {
    189.898 us [32275] |         b() {
    189.898 us [32275] |             c() {
    3.899 us [32275] |                 getpid();
    5.974 us [32275] |             } /* c */
    6.690 us [32275] |         } /* b */
    7.437 us [32275] |     } /* a */
   16.142 us [32275] | } /* main */

   956.108 us [32272] |     } /* wait */
   956.108 us [32272] |     a() {
   956.108 us [32272] |         b() {
   956.108 us [32272] |             c() {
    4.290 us [32272] |                 getpid();
    5.868 us [32272] |             } /* c */
    6.515 us [32272] |         } /* b */
    7.132 us [32272] |     } /* a */
    1.177 ms [32272] | } /* main */
```

Multiprocess program example

```
$ ufttrace -E linux:schedule t-fork
```

```
# DURATION      TID      FUNCTION
  1.609 us [32272] | __monstartup();
  1.493 us [32272] | __cxa_atexit();
189.898 us [32272] | main() {
                    [32272] |     fork();
                    [32272] |     wait() {
                        [32272] |         /* linux:sched-out */
                        [32275] |     } /* fork */
                    [32275] |     a() {
                        [32275] |         b() {
                            [32275] |             c() {
                                    getpid();
                                } /* c */
                            } /* b */
                        } /* a */
                    } /* main */
  16.142 us [32275] |
  919.404 us [32272] |     /* linux:sched-in */
  956.108 us [32272] |     } /* wait */
                    [32272] |     a() {
                        [32272] |         b() {
                            [32272] |             c() {
                                    getpid();
                                } /* c */
                            } /* b */
                        } /* a */
                    } /* main */
  4.290 us [32272] |
  5.868 us [32272] |
  6.515 us [32272] |
  7.132 us [32272] |
  1.177 ms [32272] | }
```

You can see the scheduling event

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```



```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
13.147 s	13.147 s	814	select
379.970 ms	379.970 ms	1	epoll_wait
292.910 ms	292.910 ms	1	getaddrinfo

```
...
```

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
13.147 s	13.147 s	814	select
379.970 ms	379.970 ms	1	epoll_wait
292.910 ms	292.910 ms	1	getaddrinfo

```
...
```

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
13.147 s	13.147 s	814	select
379.970 ms	379.970 ms	1	epoll_wait
292.910 ms	292.910 ms	1	getaddrinfo

```
...
```

Due to waiting time..

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
13.147 s	13.147 s	814	select
379.970 ms	379.970 ms	1	epoll_wait
292.910 ms	292.910 ms	1	getaddrinfo

```
...
```

```
# uftrace record -E linux:schedule -d uftrace.data.nmap_sched -- ./nmap nmap.org
```

```
$ cd ../ && cd uftrace.data.nmap_sched/
```

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
13.147 s	13.147 s	814	select
379.970 ms	379.970 ms	1	epoll_wait
292.910 ms	292.910 ms	1	getaddrinfo

```
...
```

```
# uftrace record -E linux:schedule -d uftrace.data.nmap_sched -- ./nmap nmap.org
```

```
$ cd ../ && cd uftrace.data.nmap_sched/
```

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
12.048 s	12.048 s	108	linux:schedule
11.884 s	150.189 ms	2	ultra_scan
140.047 ms	87.156 ms	1	nmap_services_init

```
...
```

11.673 s	8.548 ms	597	select
----------	-----------------	-----	--------

It excludes schedule-out time

```
...
```

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -d uftrace.data.nmap_nmap.org -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_nmap.org/
```

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
13.147 s	13.147 s	814	select
379.970 ms	379.970 ms	1	epoll_wait
292.910 ms	292.910 ms	1	getaddrinfo

Due to waiting time..

```
...
```

```
# uftrace record -E linux:schedule -d uftrace.data.nmap_sched -- ./nmap nmap.org
```

```
$ cd ../ && cd uftrace.data.nmap_sched/
```

```
$ uftrace report -s self
```

Total time	Self time	Calls	Function
=====	=====	=====	=====
12.048 s	12.048 s	108	linux:schedule
11.884 s	150.189 ms	2	ultra_scan
140.047 ms	87.156 ms	1	nmap_services_init

```
...
```

11.673 s	8.548 ms	597	select
----------	-----------------	-----	---------------

It excludes schedule-out time

```
...
```

User + Kernel Tracing

Nmap example

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -K 3 -d uftrace.data.nmap_kern -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_kern/
```



```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -K 3 -d uftrace.data.nmap_kern -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_kern/
```

```
$ uftrace graph sendIPScanProbe
```

```
78.812 ms : (2003) sendIPScanProbe
66.874 ms : +- (2003) send_ip_packet
64.529 ms : | +- (2003) sendto
60.876 ms : | | +- (2003) sys_sendto
 2.706 ms : | | | +- (2003) sockfd_lookup_light
 1.010 ms : | | | | +- (2003) __fdget
          : | | | | |
49.669 us : | | | | | +- (1) smp_apic_timer_interrupt
          : | | | | |
915.096 us : | | | | +- (2003) move_addr_to_kernel.part.14
          : | | | | |
51.020 ms : | | | | +- (2003) sock_sendmsg
 1.732 ms : | | | | | +- (2003) security_socket_sendmsg
22.254 us : | | | | | (1) smp_apic_timer_interrupt
          : | | | | |
41.848 ms : | | | | +- (2003) inet_sendmsg
20.826 us : | | | | | (1) smp_apic_timer_interrupt
          : | | | | |
64.012 us : | | | | +- (3) smp_apic_timer_interrupt
          : | | | | |
14.037 us : | | | | +- (2) do_IRQ
```

```
...
```

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -K 3 -d uftrace.data.nmap_kern -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_kern/
```

```
$ uftrace graph sendIPScanProbe
```

```
78.812 ms : (2003) sendIPScanProbe  
66.874 ms : +-(2003) send_ip_packet  
64.529 ms : | +-(2003) sendto  
60.876 ms : | | +-(2003) sys_sendto  
2.706 ms : | | | +-(2003) sockfd_lookup_light  
1.010 ms : | | | | +-(2003) __fdget  
:  
:  
49.669 us : | | | | +-(1) smp_apic_timer_interrupt  
:  
:  
915.096 us : | | | | +-(2003) move_addr_to_kernel.part.14  
:  
:  
51.020 ms : | | | | +-(2003) sock_sendmsg  
1.732 ms : | | | | +-(2003) security_socket_sendmsg  
22.254 us : | | | | | (1) smp_apic_timer_interrupt  
:  
:  
41.848 ms : | | | | | +-(2003) inet_sendmsg  
20.826 us : | | | | | | (1) smp_apic_timer_interrupt  
:  
:  
64.012 us : | | | | | +-(3) smp_apic_timer_interrupt  
:  
:  
14.037 us : | | | | | +-(2) do_IRQ
```

] User space

```
. . .
```

```
$ cd uftrace-osseu17/nmap_examples/
```

```
# uftrace record -K 3 -d uftrace.data.nmap_kern -- ./nmap nmap.org
```

```
$ cd uftrace.data.nmap_kern/
```

```
$ uftrace graph sendIPScanProbe
```

```
78.812 ms : (2003) sendIPScanProbe  
66.874 ms : +-(2003) send_ip_packet  
64.529 ms : | +-(2003) sendto  
60.876 ms : | | +-(2003) sys_sendto  
 2.706 ms : | | | +-(2003) sockfd_lookup_light  
 1.010 ms : | | | | +-(2003) __fdget  
          : | | | | |  
49.669 us : | | | | | +-(1) smp_apic_timer_interrupt  
          : | | | | |  
915.096 us : | | | | | +-(2003) move_addr_to_kernel.part.14  
          : | | | | |  
51.020 ms : | | | | | +-(2003) sock_sendmsg  
 1.732 ms : | | | | | +-(2003) security_socket_sendmsg  
22.254 us : | | | | | | (1) smp_apic_timer_interrupt  
          : | | | | | |  
41.848 ms : | | | | | | +-(2003) inet_sendmsg  
20.826 us : | | | | | | | (1) smp_apic_timer_interrupt  
          : | | | | | | |  
64.012 us : | | | | | | | +-(3) smp_apic_timer_interrupt  
          : | | | | | | |  
14.037 us : | | | | | | | +-(2) do_IRQ
```

User space

Kernel space

```
...
```

User + Kernel Tracing

`printf()` example
(x86_64, xen)

```
$ cd uftrace-osseu17/printf_kern_examples/
```

```
$ cd uftrace.data.printf_kern/
```

```
$ cd uftrace-osseu17/printf_kern_examples/
```

```
$ cd uftrace.data.printf_kern/  recorded data on X86_64
```

```
$ cd uftrace-osseu17/printf_kern_examples/
```

```
$ cd uftrace.data.printf_kern/  recorded data on X86_64
```

```
$ uftrace replay
```

#	DURATION	TID	FUNCTION
	1.057 us	[31071]	__monstartup();
	0.940 us	[31071]	__cxa_atexit();
		[31071]	main() {
		[31071]	printf() {
...			
		[31071]	sys_write() {
		[31071]	__fdget_pos() {
	0.470 us	[31071]	__fget_light();
	1.337 us	[31071]	} /* __fdget_pos */
		[31071]	vfs_write() {
		[31071]	rw_verify_area() {
	1.224 us	[31071]	security_file_permission();
	2.330 us	[31071]	} /* rw_verify_area */
		[31071]	__vfs_write() {
	13.830 us	[31071]	tty_write();

```
$ cd uftrace-osseu17/printf_kern_examples/
```

```
$ cd uftrace.data.printf_xen/
```



```
$ cd uftrace-osseu17/printf_kern_examples/
```

```
$ cd uftrace.data.printf_xen/  recorded data on xen
```

```
$ cd uftrace-osseu17/printf_kern_examples/
```

```
$ cd uftrace.data.printf_xen/  recorded data on xen
```

```
$ uftrace replay
```

#	DURATION	TID	FUNCTION
	3.215 us	[3148]	__monstartup();
	2.751 us	[3148]	__cxa_atexit();
		[3148]	main() {
		[3148]	printf() {
...			
		[3148]	sys_write() {
		[3148]	xen_evtchn_do_upcall() {
		[3148]	irq_enter() {
0.423 us		[3148]	rcu_irq_enter();
2.085 us		[3148]	} /* irq_enter */
...			
		[3148]	tty_write() {
1.842 us		[3148]	xen_evtchn_do_upcall();
1.050 us		[3148]	xen_maybe_preempt_hcall();
25.477 us		[3148]	} /* tty_write */
...			

User + Kernel Tracing

Process Lifecycle

`fork()` → `exec()` → `main()` → `exit()` → `schedule()`

```
$ cd uftrace-osseu17/process_life_cycle_example/
```

```
# uftrace record -K 5 ./hello
```

```
$ cd uftrace.data
```

```
$ cd uftrace-osseu17/process_life_cycle_example/
```

```
# uftrace record -K 5 ./hello
```

```
$ cd uftrace.data
```

```
$ uftrace replay --kernel-full
```

```
# DURATION      TID      FUNCTION
  3.817 us [ 1037] | finish_task_switch();
  0.434 us [ 1037] | _raw_spin_lock_irq();
  0.375 us [ 1037] | __fsnotify_parent();
  0.335 us [ 1037] | fsnotify();
           [ 1037] | sys_execve() {
...
           [ 1037] | main() {
           [ 1037] |     printf() {
...
           [ 1037] |         sys_write() {
...
20.088 us [ 1037] |         } /* sys_write */
39.679 us [ 1037] |     } /* printf */
  1.166 us [ 1037] |     fflush();
41.453 us [ 1037] | } /* main */
...
           [ 1037] | sys_exit_group() {
           [ 1037] |     do_group_exit() {
...
           [ 1037] |         schedule() {
0.386 us [ 1037] |             rcu_note_context_switch();
```

```
--kernel-full Show kernel functions outside of user
```

```
$ cd uftrace-osseu17/process_life_cycle_example/
```

```
# uftrace record -K 5 ./hello
```

```
$ cd uftrace.data
```

```
$ uftrace replay --kernel-full
```

```
# DURATION    TID     FUNCTION
  3.817 us [ 1037] | finish_task_switch();
  0.434 us [ 1037] | _raw_spin_lock_irq();
  0.375 us [ 1037] | __fsnotify_parent();
  0.335 us [ 1037] | fsnotify();
           [ 1037] | sys_execve() {
...
           [ 1037] | main() {
           [ 1037] |   printf() {
...
           [ 1037] |     sys_write() {
...
20.088 us [ 1037] |     } /* sys_write */
39.679 us [ 1037] |   } /* printf */
  1.166 us [ 1037] |   fflush();
41.453 us [ 1037] | } /* main */
...
           [ 1037] | sys_exit_group() {
           [ 1037] |   do_group_exit() {
...
           [ 1037] |     schedule() {
0.386 us [ 1037] |       rcu_note_context_switch();
```

```
--kernel-full Show kernel functions outside of user
```

```
$ cd ufttrace-osseu17/process_life_cycle_example/
```

```
# ufttrace record -K 5 ./hello
```

```
$ cd ufttrace.data
```

```
$ ufttrace replay --kernel-full
```

```
# DURATION    TID     FUNCTION
 3.817 us [ 1037] | finish_task_switch();
 0.434 us [ 1037] | _raw_spin_lock_irq();
 0.375 us [ 1037] | __fsnotify_parent();
 0.335 us [ 1037] | fsnotify();
...          [ 1037] | sys_execve() {
...          [ 1037] |   main() {
...          [ 1037] |     printf() {
...          [ 1037] |       sys_write() {
...          [ 1037] |         } /* sys_write */
20.088 us [ 1037] |       } /* printf */
39.679 us [ 1037] |     fflush();
41.453 us [ 1037] |   } /* main */
...          [ 1037] | sys_exit_group() {
...          [ 1037] |   do_group_exit() {
...          [ 1037] |     schedule() {
0.386 us [ 1037] |       rcu_note_context_switch();
```

fork()




exec()



main()



exit()



schedule()

It shows process life cycle

Production quality programs examples

with **--chrome** option

(V8, clang and perf)


```
$ cd uftrace-osseu17/chrome_tracing_examples
```

```
$ ls  
clang_internal.json      v8_CPU_intensive.json  
perf_record_internal.json v8_mem_intensive.json
```

```
$ google-chrome
```

```
--chrome    Dump recorded data in chrome trace format
```

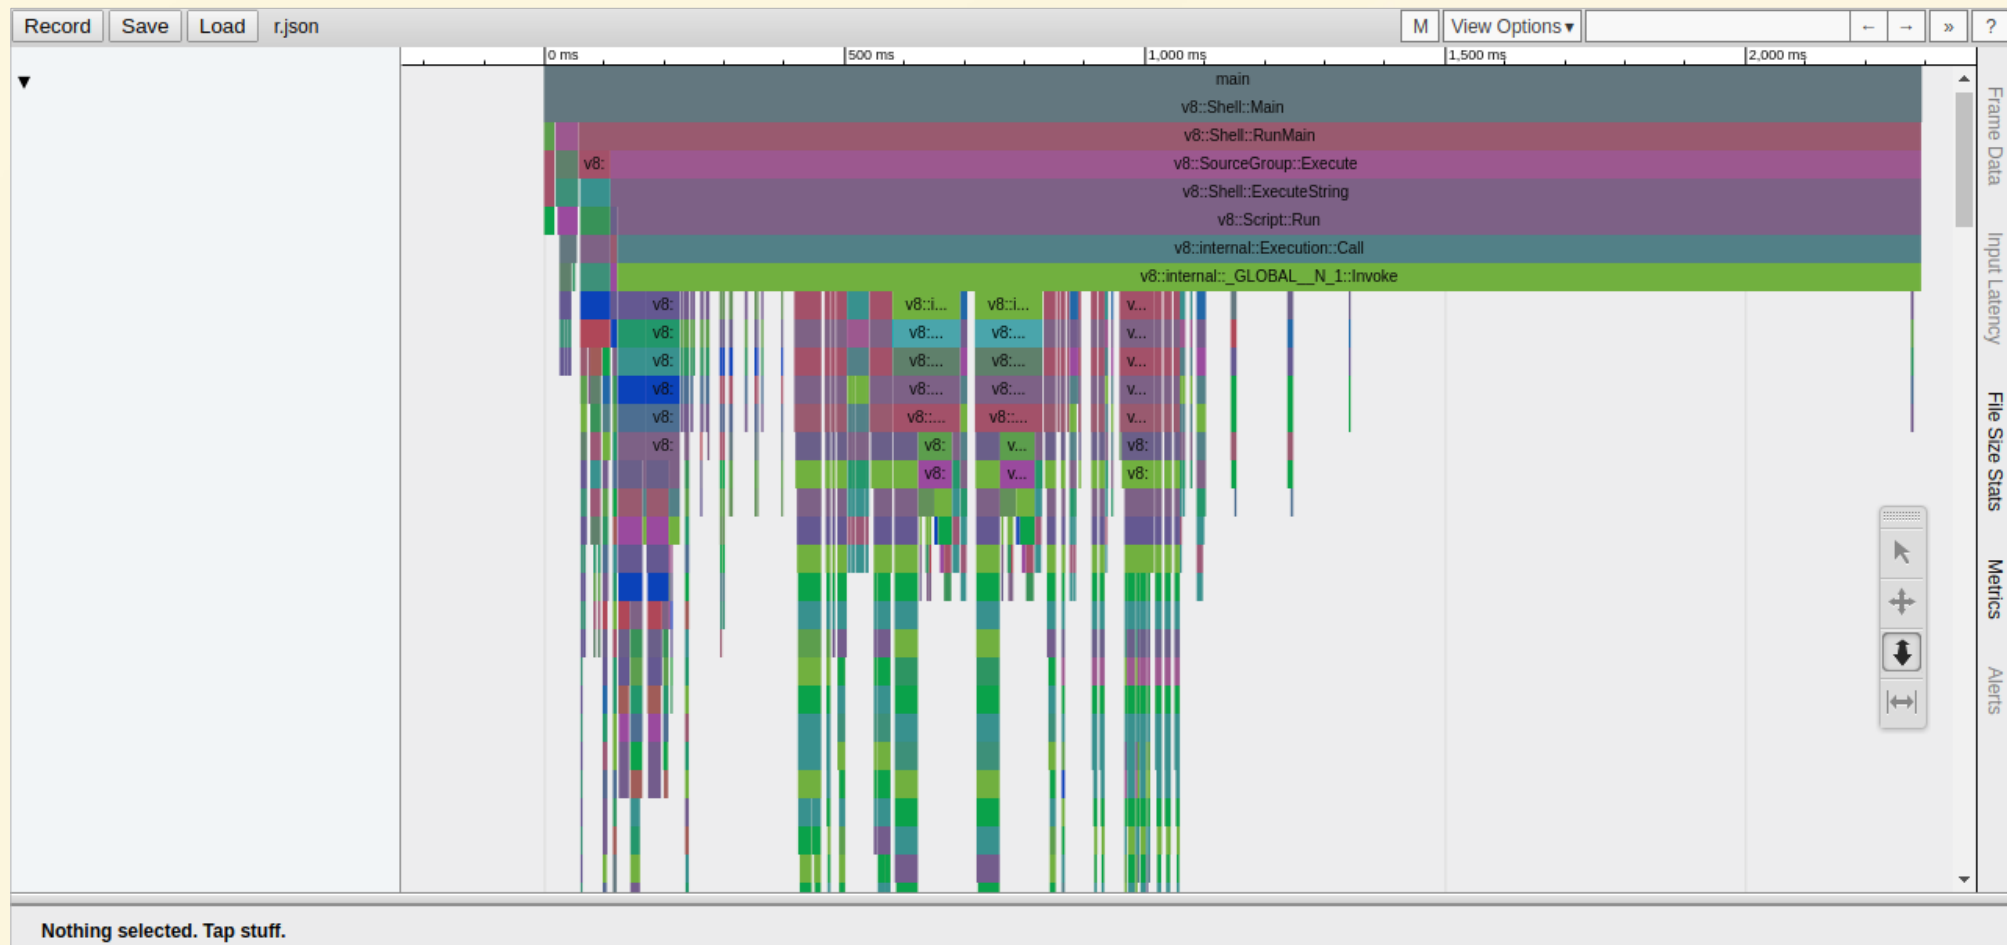
Run chrome browser

and

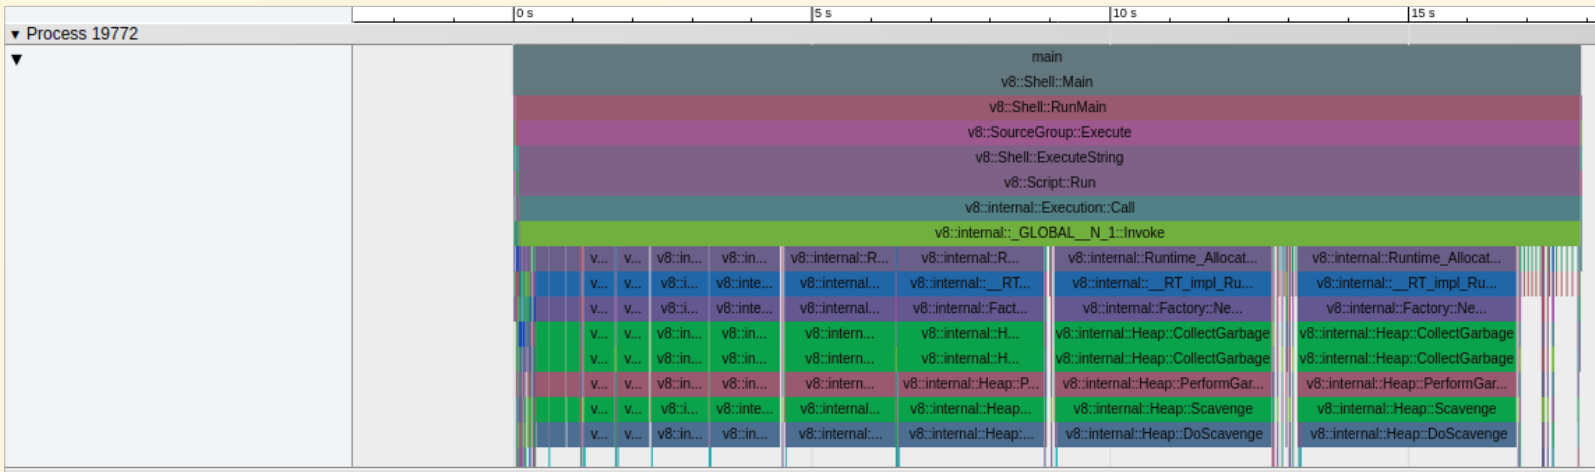

Load .json file

on chrome://tracing

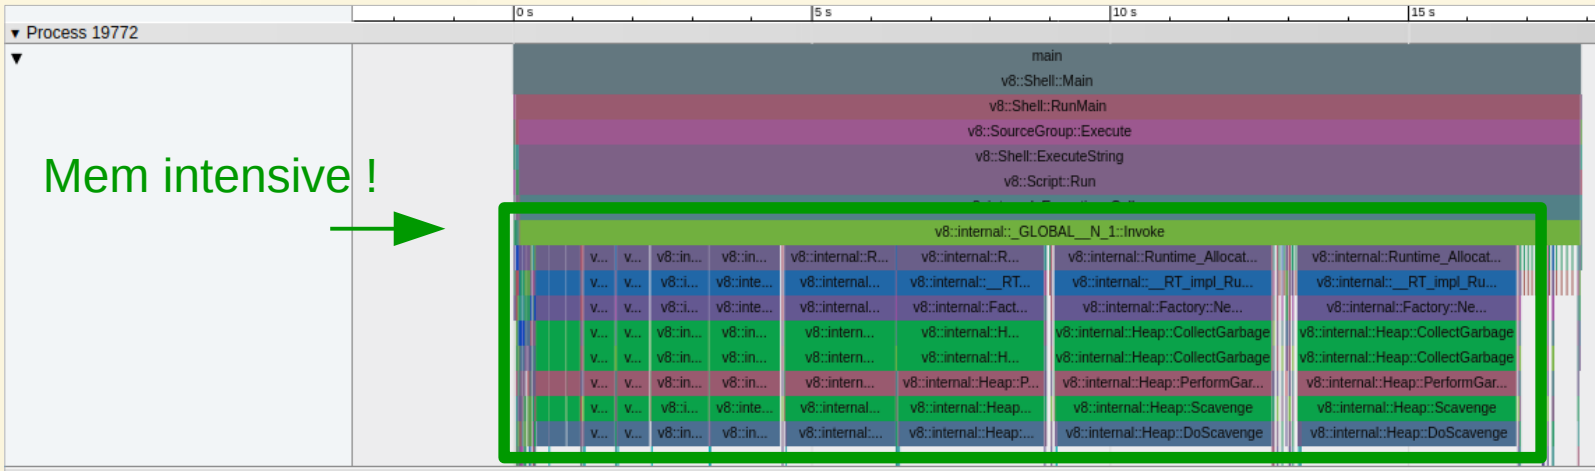

You can **zoom in-out** on **chrome://tracing**



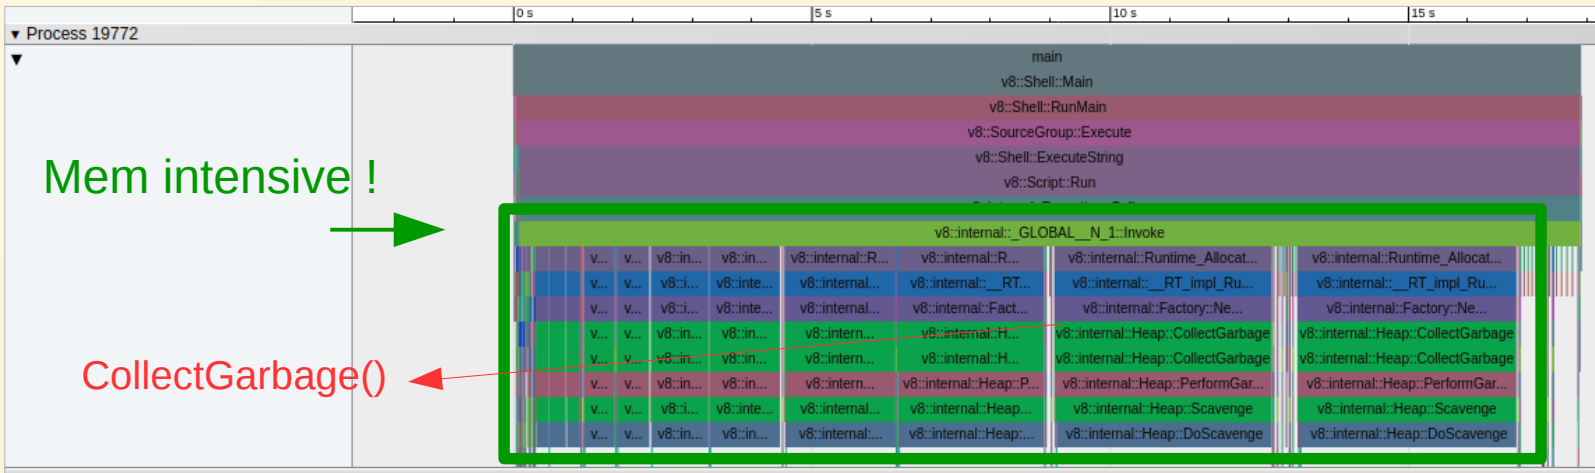

What is **different** between v8_CPU_intensive.json and v8_mem_intensive.json ?




What is **different** between v8_CPU_intensive.json and v8_mem_intensive.json ?



What is **different** between v8_CPU_intensive.json and v8_mem_intensive.json ?

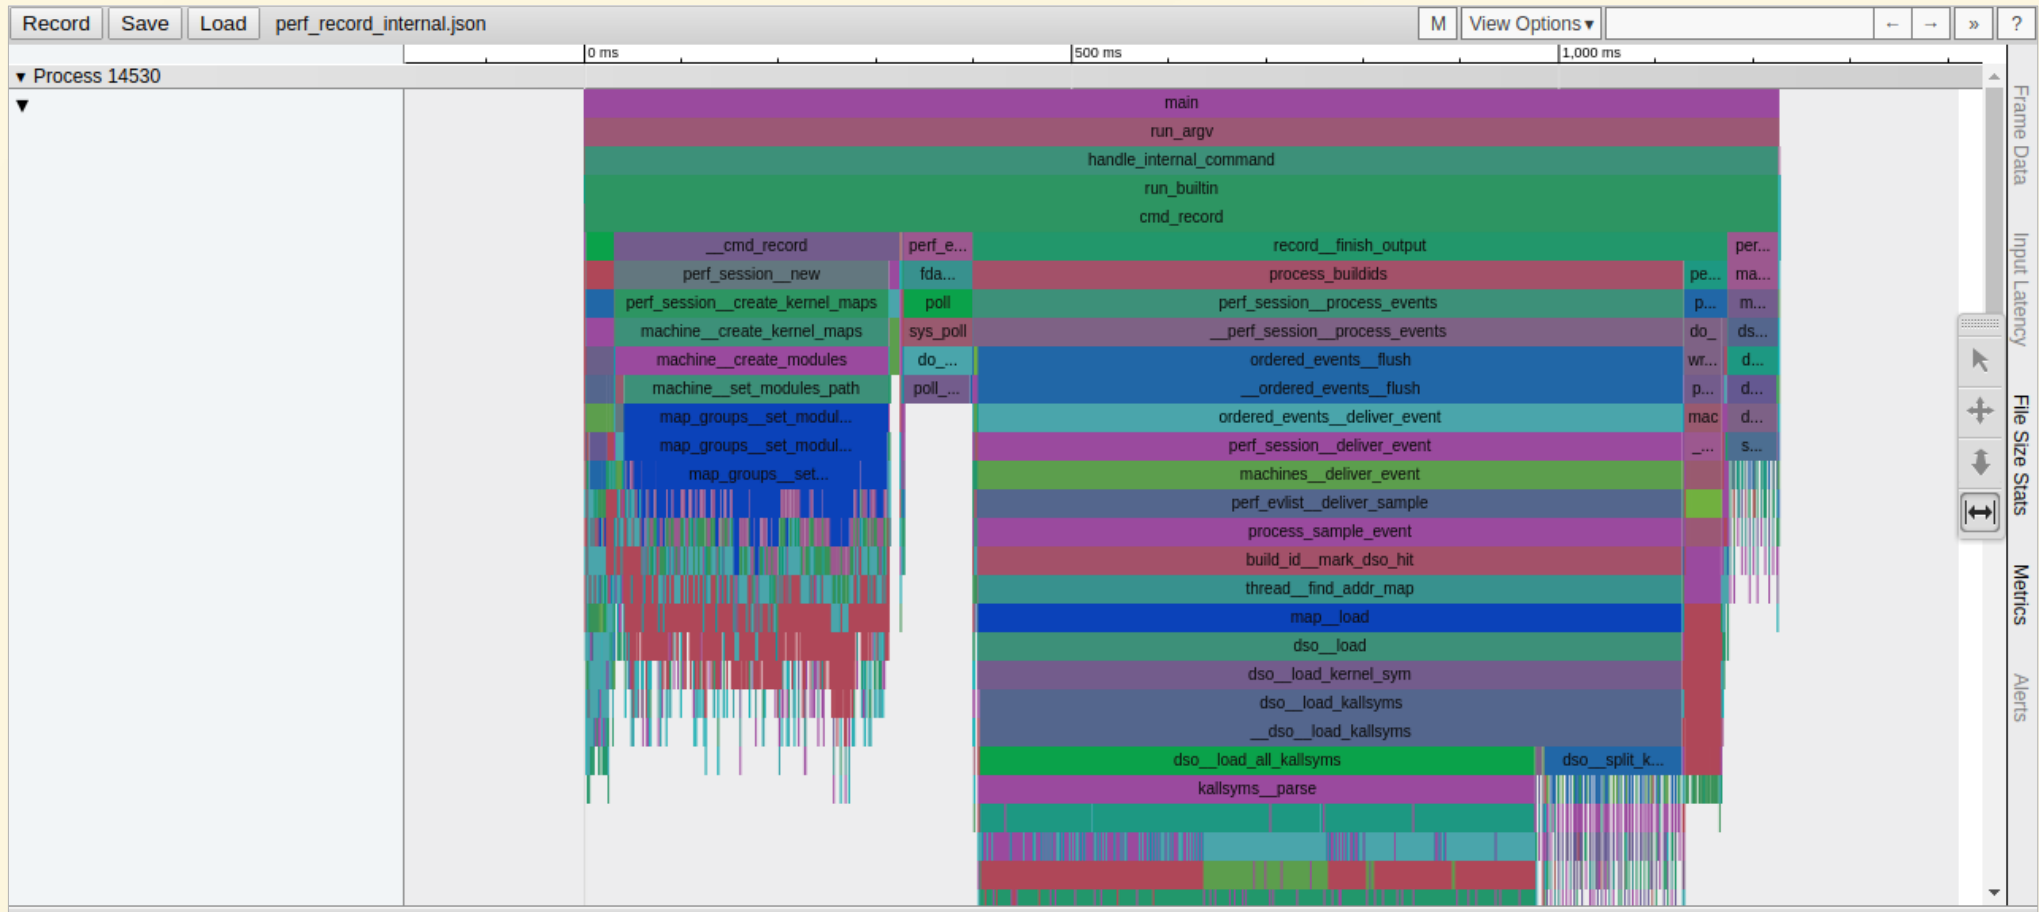


(A trace of clang (LLVM) compiling uftrace src)



perf_internal.json → chrome://tracing

(A trace of perf profiling ufttrace)



Python interpreter tracing

.py VS. **.pyc**


```
$ cd uftrace-osseu17/cpython_example
```

```
$ cd uftrace.data.hello_py && uftrace graph
```

```
calling functions
```

```
=====
```

```
101.647 ms : (1) main
101.563 ms : (1) Py_Main
85.513 ms : +- (1) Py_Initialize
85.513 ms : | (1) _Py_InitializeEx_Private.part.8
6.298 ms : | +- (1) _Py_ReadyTypes
: | |
1.368 ms : | +- (1) _PyExc_Init
: | |
1.425 ms : | +- (1) _PySys_Init
: | |
16.101 ms : | +- (1) import_init
2.894 ms : | | +- (1) PyImport_ImportFrozenModule
2.892 ms : | | | (1) PyImport_ImportFrozenModuleObject
1.273 ms : | | | +- (1) PyMarshal_ReadObjectFromString
1.271 ms : | | | | (1) r_object
1.238 ms : | | | | (1) r_object
: | | |
1.585 ms : | | | +- (1) PyEval_EvalCode
1.584 ms : | | | (1) _PyEval_EvalCodeWithName
1.581 ms : | | | (1) _PyEval_EvalFrameDefault
: | |
12.220 ms : | | +- (1) PyObject_CallMethod
12.215 ms : | | (1) PyObject_Call
12.215 ms : | | (1) function_call
12.215 ms : | | (1) PyEval_EvalCodeEx
```

```
. . .
```

recorded data
'python hello.py'

A trace of python
running **.py**

```
$ cd .. & cd uftrace.data.hello_pyc
```

```
$ uftrace graph
```

```
calling functions
```

```
=====
```

```
95.862 ms : (1) main
95.778 ms : (1) Py_Main
81.528 ms : +-(1) Py_Initialize
81.528 ms : | (1) _Py_InitializeEx_Private.part.8
6.309 ms : | +-(1) _Py_ReadyTypes
: | |
1.329 ms : | +-(1) _PyExc_Init
: | |
1.461 ms : | +-(1) _PySys_Init
: | |
15.573 ms : | +-(1) import_init
2.889 ms : | | +-(1) PyImport_ImportFrozenModule
2.888 ms : | | | (1) PyImport_ImportFrozenModuleObject
1.268 ms : | | | +-(1) PyMarshal_ReadObjectFromString
1.266 ms : | | | | (1) r_object
1.233 ms : | | | | (1) r_object
: | | | |
1.585 ms : | | | +-(1) PyEval_EvalCode
1.584 ms : | | | (1) _PyEval_EvalCodeWithName
1.581 ms : | | | (1) _PyEval_EvalFrameDefault
: | | |
12.034 ms : | | +-(1) PyObject_CallMethod
12.029 ms : | | (1) PyObject_Call
12.029 ms : | | (1) function_call
12.029 ms : | | (1) PyEval_EvalCodeEx
```

```
. . .
```

recorded data
'python hello.pyc'

Check **call graph**
of cpython
running **.pyc**

```
$ ufttrace report --diff-policy=percent -d ufttrace.data.hello_py --diff ufttrace.data.hello_pyc
```

```
#
```

```
# ufttrace diff
```

```
# [0] base: ufttrace.data.hello_py (from ufttrace record -t 1ms -d ufttrace.data.hello_py ./python hello.py )
```

```
# [1] diff: ufttrace.data.hello_pyc (from ufttrace record -t 1ms -d ufttrace.data.hello_pyc ./python __pycache__/hello.cpython-36.pyc )
```

```
#
```

Total time (diff)			Self time (diff)			Calls (diff)			Function
=====	=====	=====	=====	=====	=====	=====	=====	=====	=====
12.005 ms	9.328 ms	-22.30%	29.956 us	19.007 us	-36.55%	1	1	0	_PyCodecRegistry_Init
11.975 ms	9.309 ms	-22.26%	1.551 us	1.044 us	-32.69%	1	1	0	PyImport_ImportModuleNoBlock
13.531 ms	11.039 ms	-18.41%	7.693 us	6.273 us	-18.46%	1	1	0	initfsencoding.isra.5
13.523 ms	11.033 ms	-18.41%	4.264 us	4.986 us	+16.93%	1	1	0	_PyCodec_Lookup
6.556 ms	5.668 ms	-13.55%	14.558 us	12.911 us	-11.31%	4	4	0	marshal_loads
1.513 ms	1.700 ms	+12.32%	0.139 us	0.130 us	-6.47%	1	1	0	PyEval_CallObjectWithKeywords
10.436 ms	9.508 ms	-8.89%	1.203 ms	1.066 ms	-11.36%	1	1	0	PyImport_Cleanup
9.232 ms	8.441 ms	-8.56%	1.575 us	1.007 us	-36.06%	2	2	0	_PyGC_CollectNoFail
9.520 ms	8.735 ms	-8.24%	9.520 ms	8.735 ms	-8.24%	14	14	0	r_object
15.142 ms	13.948 ms	-7.88%	312.846 us	237.433 us	-24.11%	1	1	0	Py_FinalizeEx.part.3
15.142 ms	13.948 ms	-7.88%	0.322 us	0.307 us	-4.66%	1	1	0	Py_FinalizeEx
2.541 ms	2.344 ms	-7.76%	2.541 ms	2.344 ms	-7.76%	1	1	0	PyInit_posix
13.623 ms	12.642 ms	-7.20%	13.623 ms	12.642 ms	-7.20%	3	3	0	collect
56.833 ms	53.235 ms	-6.33%	45.423 us	38.517 us	-15.20%	4	4	0	PyImport_Import
56.787 ms	53.196 ms	-6.32%	6.530 us	6.232 us	-4.56%	4	4	0	PyObject_CallFunction
1.710 ms	1.817 ms	+6.30%	1.020 us	1.441 us	+41.27%	1	1	0	_imp_get_frozen_object
58.193 ms	54.778 ms	-5.87%	20.930 us	18.956 us	-9.43%	7	7	0	builtin__import__
63.565 ms	59.863 ms	-5.82%	10.850 us	8.682 us	-19.85%	16	17	+1	PyCFunction_Call
101.563 ms	95.778 ms	-5.70%	907.452 us	301.311 us	-66.80%	1	1	0	Py_Main
101.647 ms	95.862 ms	-5.69%	82.999 us	84.121 us	+0.15%	1	1	0	main
11.886 ms	11.209 ms	-5.69%	9.094 us	8.175 us	-10.11%	7	7	0	_PyCFunction_FastCallKeywords
70.315 ms	66.710 ms	-5.13%	357.353 us	297.186 us	-16.84%	90	89	-1	_PyFunction_FastCall
69.878 ms	66.297 ms	-5.12%	38.914 us	27.980 us	-28.10%	100	100	0	call_function
70.510 ms	66.920 ms	-5.09%	2.002 us	1.945 us	-2.85%	6	6	0	PyObject_Call
72.052 ms	68.449 ms	-5.00%	57.987 ms	55.468 ms	-4.34%	123	124	+1	_PyEval_EvalFrameDefault
62.573 ms	59.486 ms	-4.93%	25.997 us	23.394 us	-10.01%	18	18	0	PyObject_CallMethodIdObjArgs
62.563 ms	59.476 ms	-4.93%	2.505 us	2.462 us	-1.72%	18	18	0	_PyObject_FastCallDict
62.562 ms	59.475 ms	-4.93%	2.921 us	3.039 us	+4.04%	18	18	0	_PyFunction_FastCallDict

```
...
```

```
--diff=DATA
```

```
--diff-policy=POLICY
```

```
Report differences
```

```
Control diff report policy
```

Optimization Level Comparison

-O1 / -O2 / -O3

```
$ cd uftrace-osseu17/optimization_level_examples
```

```
$ cd uftrace.data.knap_01/ && uftrace replay
```

#	DURATION	TID	FUNCTION
	2.180 us	[5863]	__monstartup();
	2.651 us	[5863]	__cxa_atexit();
		[5863]	main() {
19.250 us		[5863]	fgets();
		[5863]	get_values_from() {
4.090 us		[5863]	__strdup();
3.082 us		[5863]	strchr();
3.268 us		[5863]	strtol();
0.767 us		[5863]	strtol();
14.764 us		[5863]	} /* get_values_from */
7.246 us		[5863]	malloc();
0.947 us		[5863]	fgets();
		[5863]	get_values_from() {
1.007 us		[5863]	__strdup();
0.617 us		[5863]	strtol();
4.586 us		[5863]	} /* get_values_from */
0.546 us		[5863]	fgets();
		[5863]	get_values_from() {
0.661 us		[5863]	__strdup();
0.501 us		[5863]	strtol();
3.644 us		[5863]	} /* get_values_from */
		[5863]	get_values_from() {
0.697 us		[5863]	__strdup();
0.531 us		[5863]	strtol();
0.517 us		[5863]	strtol();
4.030 us		[5863]	} /* get_values_from */
0.567 us		[5863]	fgets();

```
...
```

Too many
Repeated func calls

```
$ uftrace graph
calling functions
=====
306.501 ms : (1) main
 63.685 us : +-(46) fgets
           : |
  1.747 ms : +-(941) get_values_from
 78.848 us : | +-(72) __strdup
           : |
   3.082 us : | +-(1) strchr
           : |
 27.050 us : | +-(46) strtol
           : |
   9.912 us : +-(2) malloc
           : |
303.826 ms : +-(940) pack_knapsack
  1.445 ms : | (182) get_cond_maxprice
           : |
 14.071 us : +-(1) printf
           : |
   1.353 us : +-(1) free

...
```

So, **change** view



uftrace **graph**

```
$ uftrace graph
```

```
calling functions
```

```
=====
```

```
306.501 ms : (1) main
 63.685 us : +-(46) fgets
           : |
  1.747 ms : +-(941) get_values_from
 78.848 us : | +-(72) __strdup
           : |
  3.082 us : | +-(1) strchr
           : |
 27.050 us : | +-(46) strtol
           : |
  9.912 us : +-(2) malloc
           : |
303.826 ms : +-(940) pack_knapsack
  1.445 ms : | (182) get_cond_maxprice
           : |
 14.071 us : +-(1) printf
           : |
  1.353 us : +-(1) free
```

```
...
```

-O1



-O2

```
$ cd ../ && ufttrace.data.knap_02
```

```
$ ufttrace graph
```

```
calling functions
```

```
=====
```

```
15.019 ms : (1) main
13.344 us : +-(3) fgets
           : |
           : |
1.166 ms : +-(941) get_values_from
26.613 us : | +-(12) __strdup
           : | |
           : | |
1.493 us : | +-(1) strchr
           : | |
           : | |
1.570 us : | +-(1) strtol
           : | |
           : | |
6.408 us : +-(2) malloc
           : | |
           : | |
13.242 ms : +-(940) pack_knapsack
           : | |
           : | |
16.331 us : +-(1) printf
           : | |
           : | |
1.641 us : +-(1) free
```

`get_cond_maxprice()` **inlined**

```
...
```

-O1



-O2


```
$ uftrace report -d uftrace.data.knap_01 --diff uftrace.data.knap_02
```

```
#
```

```
# uftrace diff
```

```
# [0] base: uftrace.data.knap_01 (from uftrace record -t 0.5us -d uftrace.data.knap_01 knapsack_01 )
```

```
# [1] diff: uftrace.data.knap_02 (from uftrace record -t 0.5us -d uftrace.data.knap_02 knapsack_02 )
```

```
#
```

Total time (diff)			Self time (diff)			Calls (diff)			Function
=====			=====			=====			=====
306.501 ms	15.019 ms	-291.481 ms	838.433 us	573.861 us	-264.572 us	1	1	0	main
303.826 ms	13.242 ms	-290.584 ms	302.381 ms	13.242 ms	-289.139 ms	940	940	0	pack_knapsack
1.445 ms	-	-1.445 ms	1.445 ms	-	-1.445 ms	182	0	-182	get_cond_maxprice
1.747 ms	1.166 ms	-581.066 us	1.638 ms	1.136 ms	-501.762 us	941	941	0	get_values_from
78.848 us	26.613 us	-52.235 us	78.848 us	26.613 us	-52.235 us	72	12	-60	__strdup
63.685 us	13.344 us	-50.341 us	63.685 us	13.344 us	-50.341 us	46	3	-43	fgets
27.050 us	1.570 us	-25.480 us	27.050 us	1.570 us	-25.480 us	46	1	-45	strtol
9.912 us	6.408 us	-3.504 us	9.912 us	6.408 us	-3.504 us	2	2	0	malloc
14.071 us	16.331 us	+2.260 us	14.071 us	16.331 us	+2.260 us	1	1	0	printf
3.082 us	1.493 us	-1.589 us	3.082 us	1.493 us	-1.589 us	1	1	0	strchr
2.651 us	1.189 us	-1.462 us	2.651 us	1.189 us	-1.462 us	1	1	0	__cxa_atexit
2.180 us	1.115 us	-1.065 us	2.180 us	1.115 us	-1.065 us	1	1	0	__monstartup
1.353 us	1.641 us	+0.288 us	1.353 us	1.641 us	+0.288 us	1	1	0	free

```
...
```

```
--diff=DATA
```

```
Report differences
```

```
--diff-policy=POLICY
```

```
Control diff report policy
```

```
$ cd ../ && ufttrace.data.knap_02
```

```
$ ufttrace graph
```

```
calling functions
```

```
=====
15.019 ms : (1) main
13.344 us : +-(3) fgets
           : |
1.166 ms : +-(941) get_values_from
26.613 us : | +-(12) __strdup
           : | |
1.493 us : | +-(1) strchr
           : | |
1.570 us : | +-(1) strtol
           : |
6.408 us : +-(2) malloc
           : |
13.242 ms : +-(940) pack_knapsack
           : |
16.331 us : +-(1) printf
           : |
1.641 us : +-(1) free
```

```
...
```

O2

```
$ cd ../ && ufttrace.data.knap_02
```

```
$ ufttrace graph
```

```
calling functions
```

```
=====
```

```
15.019 ms : (1) main
13.344 us : +-(3) fgets
          : |
1.166 ms : + (941) get_values_from
26.613 us : | +-(12) __strdup
          : | |
1.493 us : | +-(1) strchr
          : | |
1.570 us : | +-(1) strtol
          : | |
6.408 us : +-(2) malloc
          : | |
13.242 ms : + (940) pack_knapsack
          : | |
16.331 us : +-(1) printf
          : | |
1.641 us : +-(1) free
```

```
...
```

O2



O3

```
$ cd ../ && ufttrace.data.knap_03
```

```
$ ufttrace graph  
calling functions
```

```
=====
```

		Reduced call counts
7.143 ms	: (1) main	
8.452 us	: +-(3) fgets	941 → 1
	:	
3.917 us	: +-(1) get_values_from	
1.060 us	: +-(1) __strdup	
	:	
0.760 us	: +-(1) strchr	
	:	
0.896 us	: +-(1) strtol	
	:	
3.249 us	: +-(2) malloc	
	:	
9.401 us	: +-(7) __strdup	pack_knapsack() inlined into main()
	:	
11.063 us	: +-(1) memset	
	:	
11.866 us	: +-(1) printf	
	:	
1.370 us	: +-(1) free	

```
...
```

O2



O3

```
$ uftrace report -d uftrace.data.knap_02 --diff uftrace.data.knap_03
```

```
#
```

```
# uftrace diff
```

```
# [0] base: uftrace.data.knap_02 (from uftrace record -t 0.5us -d uftrace.data.knap_02 knapsack_02 )
```

```
# [1] diff: uftrace.data.knap_03 (from uftrace record -t 0.5us -d uftrace.data.knap_03 knapsack_03 )
```

```
#
```

Total time (diff)			Self time (diff)			Calls (diff)			Function
=====			=====			=====			=====
13.242 ms	-	-13.242 ms	13.242 ms	-	-13.242 ms	940	0	-940	pack_knapsack
15.019 ms	7.143 ms	-7.875 ms	573.861 us	7.094 ms	+6.520 ms	1	1	0	main
1.166 ms	3.917 us	-1.162 ms	1.136 ms	1.201 us	-1.135 ms	941	1	-940	get_values_from
26.613 us	10.461 us	-16.152 us	26.613 us	10.461 us	-16.152 us	12	8	-4	__strdup
-	11.063 us	+11.063 us	-	11.063 us	+11.063 us	0	1	+1	memset
13.344 us	8.452 us	-4.892 us	13.344 us	8.452 us	-4.892 us	3	3	0	fgets
16.331 us	11.866 us	-4.465 us	16.331 us	11.866 us	-4.465 us	1	1	0	printf
6.408 us	3.249 us	-3.159 us	6.408 us	3.249 us	-3.159 us	2	2	0	malloc
1.493 us	0.760 us	-0.733 us	1.493 us	0.760 us	-0.733 us	1	1	0	strchr
1.570 us	0.896 us	-0.674 us	1.570 us	0.896 us	-0.674 us	1	1	0	strtol
1.189 us	0.709 us	-0.480 us	1.189 us	0.709 us	-0.480 us	1	1	0	__cxa_atexit
1.115 us	0.663 us	-0.452 us	1.115 us	0.663 us	-0.452 us	1	1	0	__monstartup
1.641 us	1.370 us	-0.271 us	1.641 us	1.370 us	-0.271 us	1	1	0	free

```
...
```

```
--diff=DATA
```

```
Report differences
```

```
--diff-policy=POLICY
```

```
Control diff report policy
```

NULL Pointer Exception case

GDB core VS. uftrace.data

```
$ cd uftrace-osseu17/nullptr_exception_example

$ ./benchmark_STL_containers_nullptr
*** Error in `./benchmark_STL_containers_nullptr': double free or
corruption (fasttop): 0x0000000002389710 ***
===== Backtrace: =====
/lib/x86_64-linux-gnu/libc.so.6(+0x777e5)[0x7f29349ae7e5]
/lib/x86_64-linux-gnu/libc.so.6(+0x8037a)[0x7f29349b737a]
/lib/x86_64-linux-gnu/libc.so.6(cfree+0x4c)[0x7f29349bb53c]
./benchmark_STL_containers_nullptr[0x400c8a]
/lib/x86_64-linux-gnu/libc.so.6(__libc_start_main+0xf0)[0x7f2934957830]
./benchmark_STL_containers_nullptr[0x400ea9]
===== Memory map: =====
...
Aborted (core dumped)
```

Crash case

If some program crash,

core file contains the last stacktrace


```
$ gdb ./benchmark_STL_containers_nullptr core
```

```
GNU gdb (Ubuntu 7.11.1-0ubuntu1~16.5) 7.11.1
```

```
...
```

```
Reading symbols from ./benchmark_STL_containers_nullptr...(no debugging symbols found)...done.
```

```
warning: core file may not match specified executable file.
```

```
[New LWP 15820]
```

```
Core was generated by `./benchmark_STL_containers_nullptr'.
```

```
Program terminated with signal SIGABRT, Aborted.
```

```
#0 0x00007f607e166428 in __GI_raise (sig=sig@entry=6) at ../sysdeps/unix/sysv/linux/raise.c:54
```

```
54  ../sysdeps/unix/sysv/linux/raise.c: No such file or directory.
```

```
(gdb) bt
```

```
#0 0x00007f607e166428 in __GI_raise (sig=sig@entry=6) at ../sysdeps/unix/sysv/linux/raise.c:54
```

```
#1 0x00007f607e16802a in __GI_abort () at abort.c:89
```

```
#2 0x00007f607e1a87ea in __libc_message (do_abort=do_abort@entry=2, fmt=fmt@entr=0x7f607e2c1e98 ...
```

```
#3 0x00007f607e1b137a in malloc_printerr (ar_ptr=<optimized out>, ptr=<optimized out>, ..
```

```
#4 _int_free (av=<optimized out>, p=<optimized out>, have_lock=0) at malloc.c:3867
```

```
#5 0x00007f607e1b553c in __GI___libc_free (mem=<optimized out>) at malloc.c:2968
```

```
#6 0x000000000000400c8a in main ()
```

```
(gdb)
```

uftrace contains full **trace**
until the **crash point**

```
# cd uftrace.data.nullptr/ && uftrace replay  
$ uftrace ./benchmark_STL_containers_nullptr
```

```
# cd uftrace.data.nullptr/ && uftrace replay
$ uftrace ./benchmark_STL_containers_nullptr
*** Error in `./benchmark_STL_containers_nullptr': double free or corruption (fasttop): ...
```

```
...
```

```
process crashed by signal 6: Aborted (si_code: -6)
```

```
child terminated by signal: 6: Aborted
```

```
# DURATION      TID      FUNCTION
    [30869] | main() {
    [30869] |   bench_vector_push_back() {
0.842 us [30869] |     std::vector::_M_insert_aux();
3.036 us [30869] |     std::vector::_M_insert_aux();
1.246 us [30869] |     std::vector::_M_insert_aux();
1.648 us [30869] |     std::vector::_M_insert_aux();
10.665 us [30869] |   } /* bench_vector_push_back */
2.153 us [30869] |   bench_deque_push_back();
    [30869] |   bench_list_push_back() {
...
11.736 us [30869] |   } /* bench_list_push_back */
0.139 us [30869] |   operator new();
0.168 us [30869] |   operator delete();
0.110 us [30869] |   operator delete();
    [30869] |   operator delete() {
```

```
uftrace stopped tracing with remaining functions
```

```
=====
```

```
task: 3651
```

```
[1] operator delete
```

```
[0] main
```

```
# cd uftrace.data.nullptr/ && uftrace replay
$ uftrace ./benchmark_STL_containers_nullptr
*** Error in `./benchmark_STL_containers_nullptr': double free or corruption (fasttop): ...
```

```
...
process crashed by signal 6: Aborted (si_code: -6)
child terminated by signal: 6: Aborted
```

#	DURATION	TID	FUNCTION
		[30869]	main() {
		[30869]	bench_vector_push_back() {
0.842	us	[30869]	std::vector::_M_insert_aux();
3.036	us	[30869]	std::vector::_M_insert_aux();
1.246	us	[30869]	std::vector::_M_insert_aux();
1.648	us	[30869]	std::vector::_M_insert_aux();
10.665	us	[30869]	} /* bench_vector_push_back */
2.153	us	[30869]	bench_deque_push_back();
		[30869]	bench_list_push_back() {
...			
11.736	us	[30869]	} /* bench_list_push_back */
0.139	us	[30869]	operator new();
0.168	us	[30869]	operator delete();
0.110	us	[30869]	operator delete();
		[30869]	operator delete() {

Crash point !

uftrace stopped tracing with remaining functions

```
task: 3651
[1] operator delete
[0] main
```

```
# cd uftrace.data.nullptr/ && uftrace replay
$ uftrace ./benchmark_STL_containers_nullptr
*** Error in `./benchmark_STL_containers_nullptr': double free or corruption (fasttop): ...
```

```
...
process crashed by signal 6: Aborted (si_code: -6)
child terminated by signal: 6: Aborted
```

#	DURATION	TID	FUNCTION
		[30869]	main() {
		[30869]	bench_vector_push_back() {
0.842	us	[30869]	std::vector::_M_insert_aux();
3.036	us	[30869]	std::vector::_M_insert_aux();
1.246	us	[30869]	std::vector::_M_insert_aux();
1.648	us	[30869]	std::vector::_M_insert_aux();
10.665	us	[30869]	} /* bench_vector_push_back */
2.153	us	[30869]	bench_deque_push_back();
		[30869]	bench_list_push_back() {
...			
11.736	us	[30869]	} /* bench_list_push_back */
0.139	us	[30869]	operator new();
0.168	us	[30869]	operator delete();
0.110	us	[30869]	operator delete();
		[30869]	operator delete() {

Full trace
until the Crash point

uftrace stopped tracing with remaining functions

```
task: 3651
[1] operator delete
[0] main
```

Crash point !

Dynamic Tracing

-pg built binary calls **mcount()**
at the entry of each function

uftrace **preloads** its **mcount()**
instead of **mcount()** in libc.so

uftrace **preloads** its **mcount()**
instead of **mcount()** in libc.so



LD_PRELOAD=libmcount.so

```
$ gcc -pg -o fibonacci tests/s-fibonacci.c
```

```
$ readelf -s fibonacci | grep mcount
 4: 0000000000000000      0 FUNC      GLOBAL DEFAULT  UND mcount@GLIBC_2.2.5 (2)
63: 0000000000000000      0 FUNC      GLOBAL DEFAULT  UND mcount@@GLIBC_2.2.5
```

You can check
mcount() !

```
$ gcc -pg -o fibonacci tests/s-fibonacci.c
```

```
$ readelf -s fibonacci | grep mcount
```

```
 4: 0000000000000000      0 FUNC      GLOBAL DEFAULT  UND mcount@GLIBC_2.2.5 (2)
63: 0000000000000000      0 FUNC      GLOBAL DEFAULT  UND mcount@@GLIBC_2.2.5
```

```
$ objdump -d fibonacci
```

```
<fib>:
```

```
    push    %rbp
    mov     %rsp,%rbp
    call    <mcount@plt>
```

```
<main>:
```

```
    push    %rbp
    mov     %rsp,%rbp
    call    <mcount@plt>
```

You can check
mcount() !

Performance issue ?

Dynamic Tracing

- compile time: **mcount()** → **NOP**
- runtime : **NOP** → **trace function** (e.g. mcount, fentry)

```
$ gcc -pg -mfentry -o fibonacci \  
> tests/s-fibonacci.c
```

mcount()

→ **__fentry__()**

```
$ gcc -pg -mfentry -o fibonacci \  
> tests/s-fibonacci.c
```

```
$ objdump -d fibonacci
```

```
<fib>:  
    call    <__fentry__@plt>  
    push    %rbp  
    mov     %rsp,%rbp
```

```
<main>:  
    call    <__fentry__@plt>  
    push    %rbp  
    mov     %rsp,%rbp
```

mcount()

→ __fentry__()

at the very beginning


```
$ gcc -pg -mfentry -mnop-mcount -o fibonacci \  
> tests/s-fibonacci.c
```

```
$ objdump -d fibonacci
```

```
<fib>:
```

```
    nop
```

```
    push    %rbp
```

```
    mov     %rsp,%rbp
```

```
<main>:
```

```
    nop
```

```
    push    %rbp
```

```
    mov     %rsp,%rbp
```

__fentry__()

→ **nop**

at the very beginning

```
$ man uftrace record
```

```
-P FUNC, --patch=FUNC
```

```
Patch FUNC dynamically. This is only applicable  
binaries built with -pg -mfentry -mnop-mcount on  
x86_64. This option can be used more than once.  
See DYNAMIC TRACING.
```

Dynamic Tracing

```
$ uftrace fibonacci 5
```

```
ERROR: Can't find 'mcount' symbol in the 'fibonacci'.
```

```
<fib>:
```

```
    nop
```

```
    push    %rbp
```

```
    mov     %rsp,%rbp
```

```
<main>:
```

```
    nop
```

```
    push    %rbp
```

```
    mov     %rsp,%rbp
```

Normal Tracing

Error !

Because of **No** mcount()

```
$ uftrace -P fib fibonacci 5
```

```
<fib>:  
    call    <__fentry__@plt>  
    push    %rbp  
    mov     %rsp,%rbp
```

```
<main>:  
    nop  
    push    %rbp  
    mov     %rsp,%rbp
```

-P, --patch=FUNC Apply dynamic patching for FUNCs

Dynamic Tracing

nop → **__fentry__()**

```

$ uftrace -P fib fibonacci 5
# DURATION      TID      FUNCTION
  0.639 us [18039] | __monstartup();
  0.481 us [18039] | __cxa_atexit();
                [18039] | fib() {
                [18039] |     fib() {
                [18039] |         fib() {
  0.076 us [18039] |             fib();
  0.059 us [18039] |             fib();
  0.546 us [18039] |         } /* fib */
  0.042 us [18039] |         fib();
  0.830 us [18039] |     } /* fib */
                [18039] |     fib() {
  0.059 us [18039] |         fib();
  0.045 us [18039] |         fib();
  0.367 us [18039] |     } /* fib */
  1.447 us [18039] | } /* fib */

```

-P, --patch=FUNC Apply dynamic patching for FUNCS

Dynamic Tracing

nop → **__fentry__()**

```

$ ufttrace -P . fibonacci 5
# DURATION      TID      FUNCTION
  0.610 us [17960] | __monstartup();
  0.487 us [17960] | __cxa_atexit();
                [17960] | main() {
                [17960] |     fib() {
                [17960] |         fib() {
                [17960] |             fib() {
  0.074 us [17960] |                 fib();
  0.055 us [17960] |                 fib();
  0.492 us [17960] |             } /* fib */
  0.059 us [17960] |         fib();
  0.765 us [17960] |     } /* fib */
                [17960] |     fib() {
  0.045 us [17960] |         fib();
  0.042 us [17960] |         fib();
  0.332 us [17960] |     } /* fib */
  1.293 us [17960] | } /* fib */
  2.384 us [17960] | } /* main */

```

-P, --patch=FUNC Apply dynamic patching for FUNCS

You can **dynamic trace**

all functions with '**-P .**'

Use case

Qt QML engine **bugfix** using ufttrace

Blog: <https://www.kdab.com/fixing-bugs-via-lateral-thinking/>

Youtube: <https://www.youtube.com/watch?v=lbCaBqhncKQ>



Fixing bugs via lateral thinking

21.12.2016

Giuseppe D'Angelo

No comments



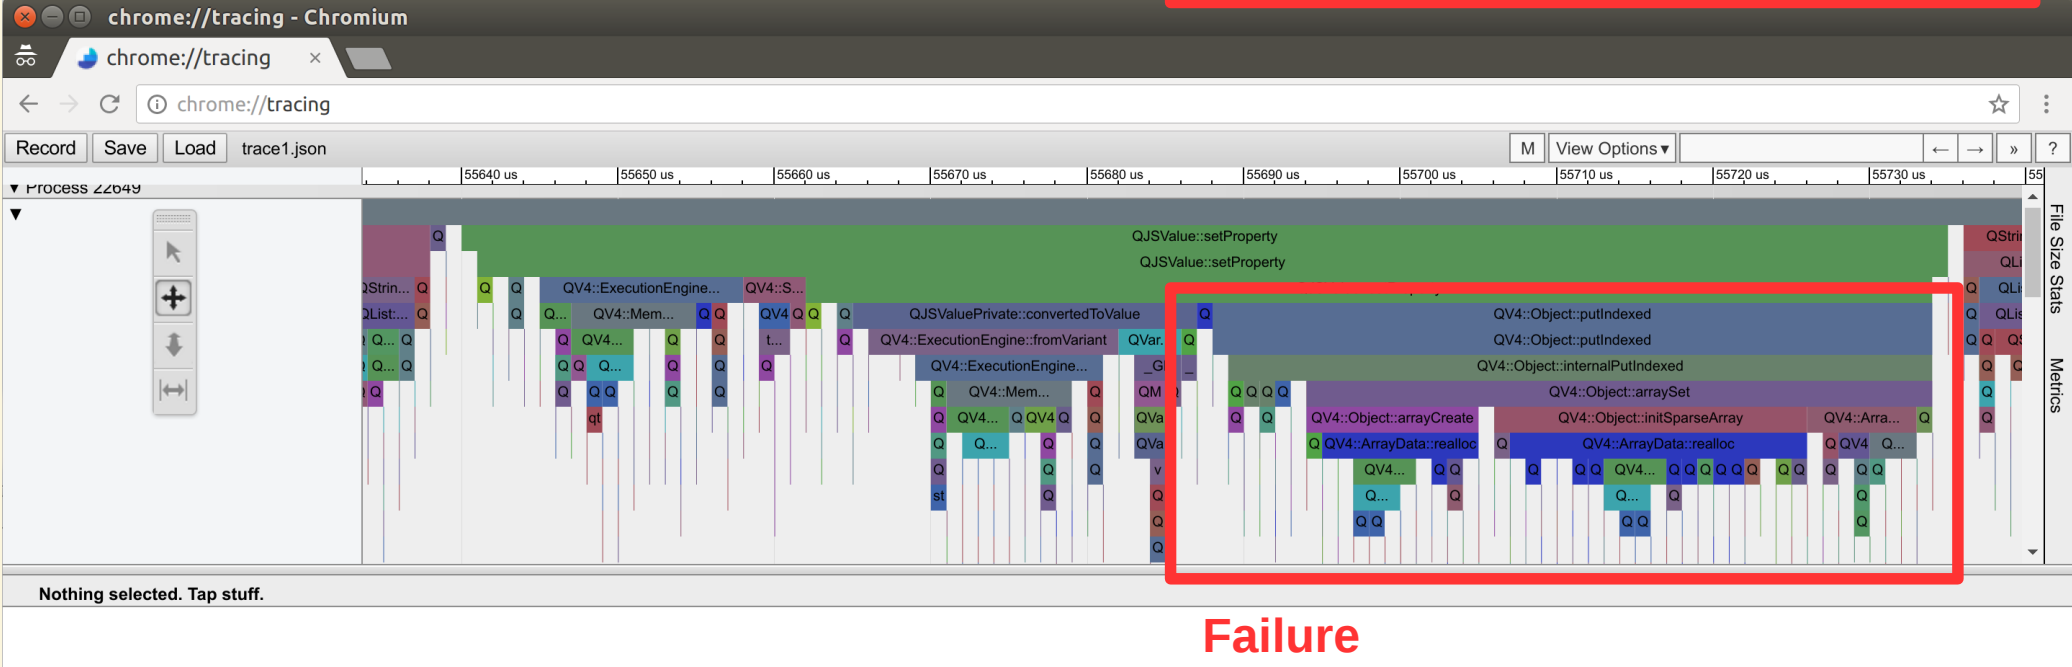
For today's blog I would like to share with you the little adventure I had when fixing a very strange bug in Qt.

Pop quiz

Don't think about this too much, just guess: what does this QML snippet print?

```
1 import QtQuick 2.0
2
3 QtObject {
4     Component.onCompleted: {
5         console.log("2400000000000" == "3776798720");
6     }
7 }
```

There are no JavaScript semantic tricks involved; and using either `==` or `===` does not change the result in any way.



bacardi opensource project:

Node.js C++ binding code **improvement** using uftrace

Slide: <https://www.slideshare.net/devview/131chromium-binding-nodejs/108>

PR: <https://github.com/lunchclass/bacardi/pull/108>

uftime를 사용한 call-graph 추적

DEVIEW
2017

```
zino — romandev@romandev-high-cpu: ~/bacardi — ssh romandev@165.227.3.183 — 80x24
0.113 us [ 2079] |      Napi::Value::IsNumber() {
0.299 us [ 2079] |      Napi::Value::Type();
[ 2079] |      } /* Napi::Value::IsNumber */
[ 2079] |      Napi::Value::ToNumber() {
[ 2079] |      Napi::Number::Number() {
0.060 us [ 2079] |      Napi::Value::Value();
0.217 us [ 2079] |      } /* Napi::Number::Number */
0.437 us [ 2079] |      } /* Napi::Value::ToNumber */
1.126 us [ 2079] |      Napi::Number::DoubleValue();
2.579 us [ 2079] |      } /* NativeTypeTraits::NativeValue */
0.064 us [ 2079] |      Napi::CallbackInfo::operator []() {
0.216 us [ 2079] |      Napi::Value::Value();
0.065 us [ 2079] |      } /* Napi::CallbackInfo::operator [] */
0.205 us [ 2079] |      Napi::CallbackInfo::Env() {
[ 2079] |      Napi::Env::Env();
[ 2079] |      } /* Napi::CallbackInfo::Env */
[ 2079] |      NativeTypeTraits::NativeValue() {
[ 2079] |      Napi::Value::IsNumber() {
0.079 us [ 2079] |      Napi::Value::Type();
0.245 us [ 2079] |      } /* Napi::Value::IsNumber */
[ 2079] |      Napi::Value::ToNumber() {
[ 2079] |      Napi::Number::Number() {
0.063 us [ 2079] |      Napi::Value::Value();
:
```

Native Type
Converting을 위해
Type Checking을 반
복적으로 하고 있음을
알 수 있음.

Thank you



`treeze.taeung@gmail.com`

`taeung@kosslab.kr`

<https://github.com/namhyung/uftrace>

Appendix

uftrace **Internals**

LD_PRELOAD

fork / exec

a.out

libmcount.so

uftrace

pipe

Code

```
...  
<foo>:  
  call mcount@plt  
...  
<main>:  
  ...  
  call foo
```

Stack

...

PLT

GOT

plt_hooker()

...

```
__attribute__((constructor))  
mcount_init()  
(running before main(),  
 to prepare tracing)  
- uftrace options  
- plthook setup  
- dynamic tracing setup  
...
```

shmem

rstack

...
...
...

Event

- read trigger `/proc/statm` - **SDT** user event
- perf scheduler in/out - Tracepoint
syscall: `SYS_perf_event_open`

Kernel

Ftrace: `/sys/kernel/debug/tracing/`
- enable / disable
- collect data

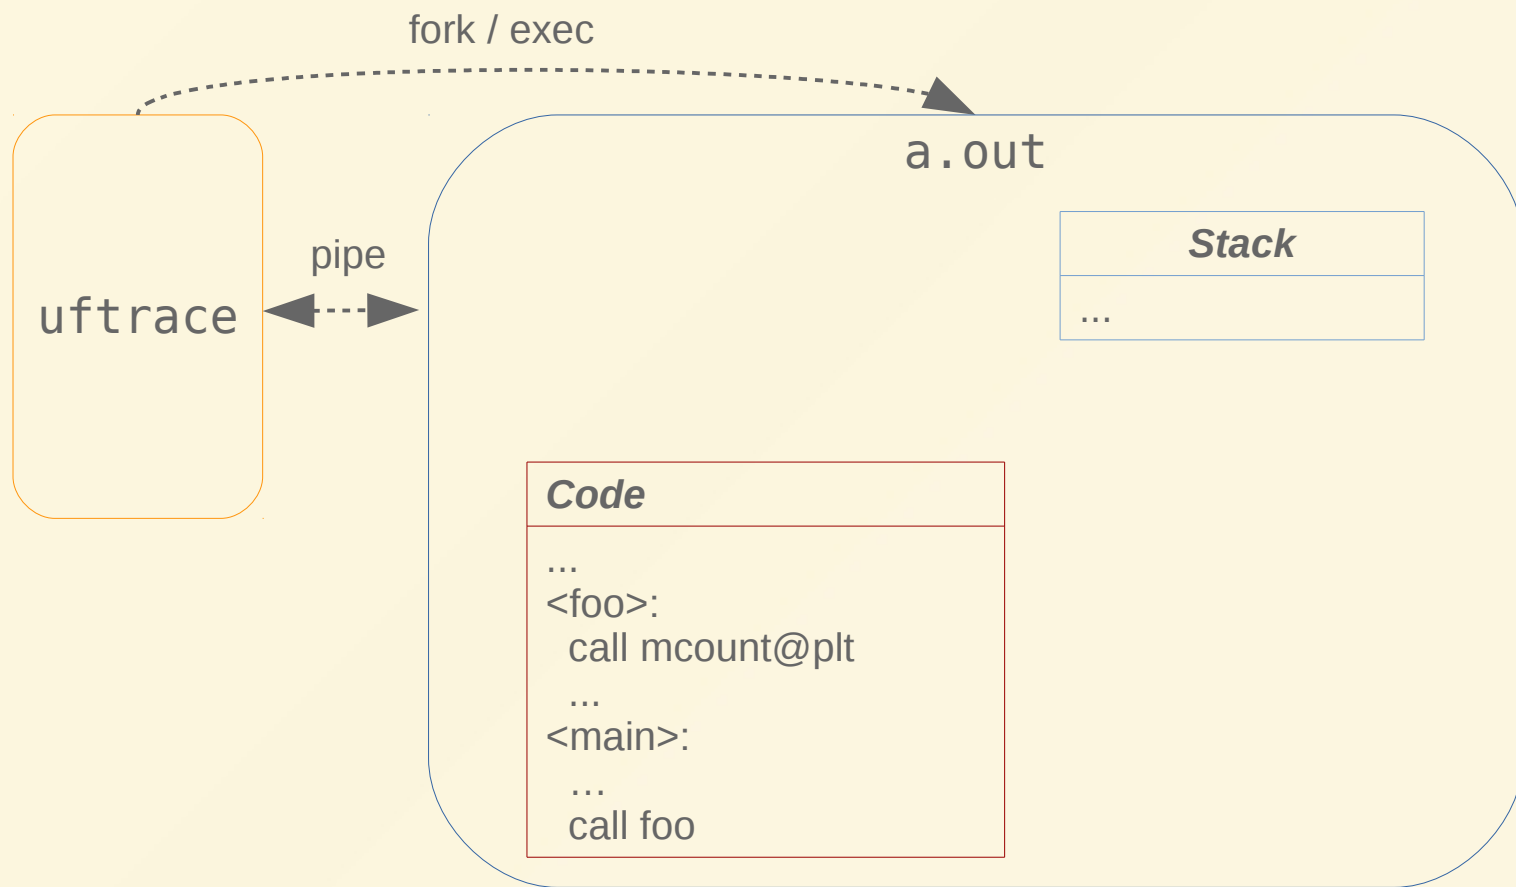
mcount hooking (user space function tracing)

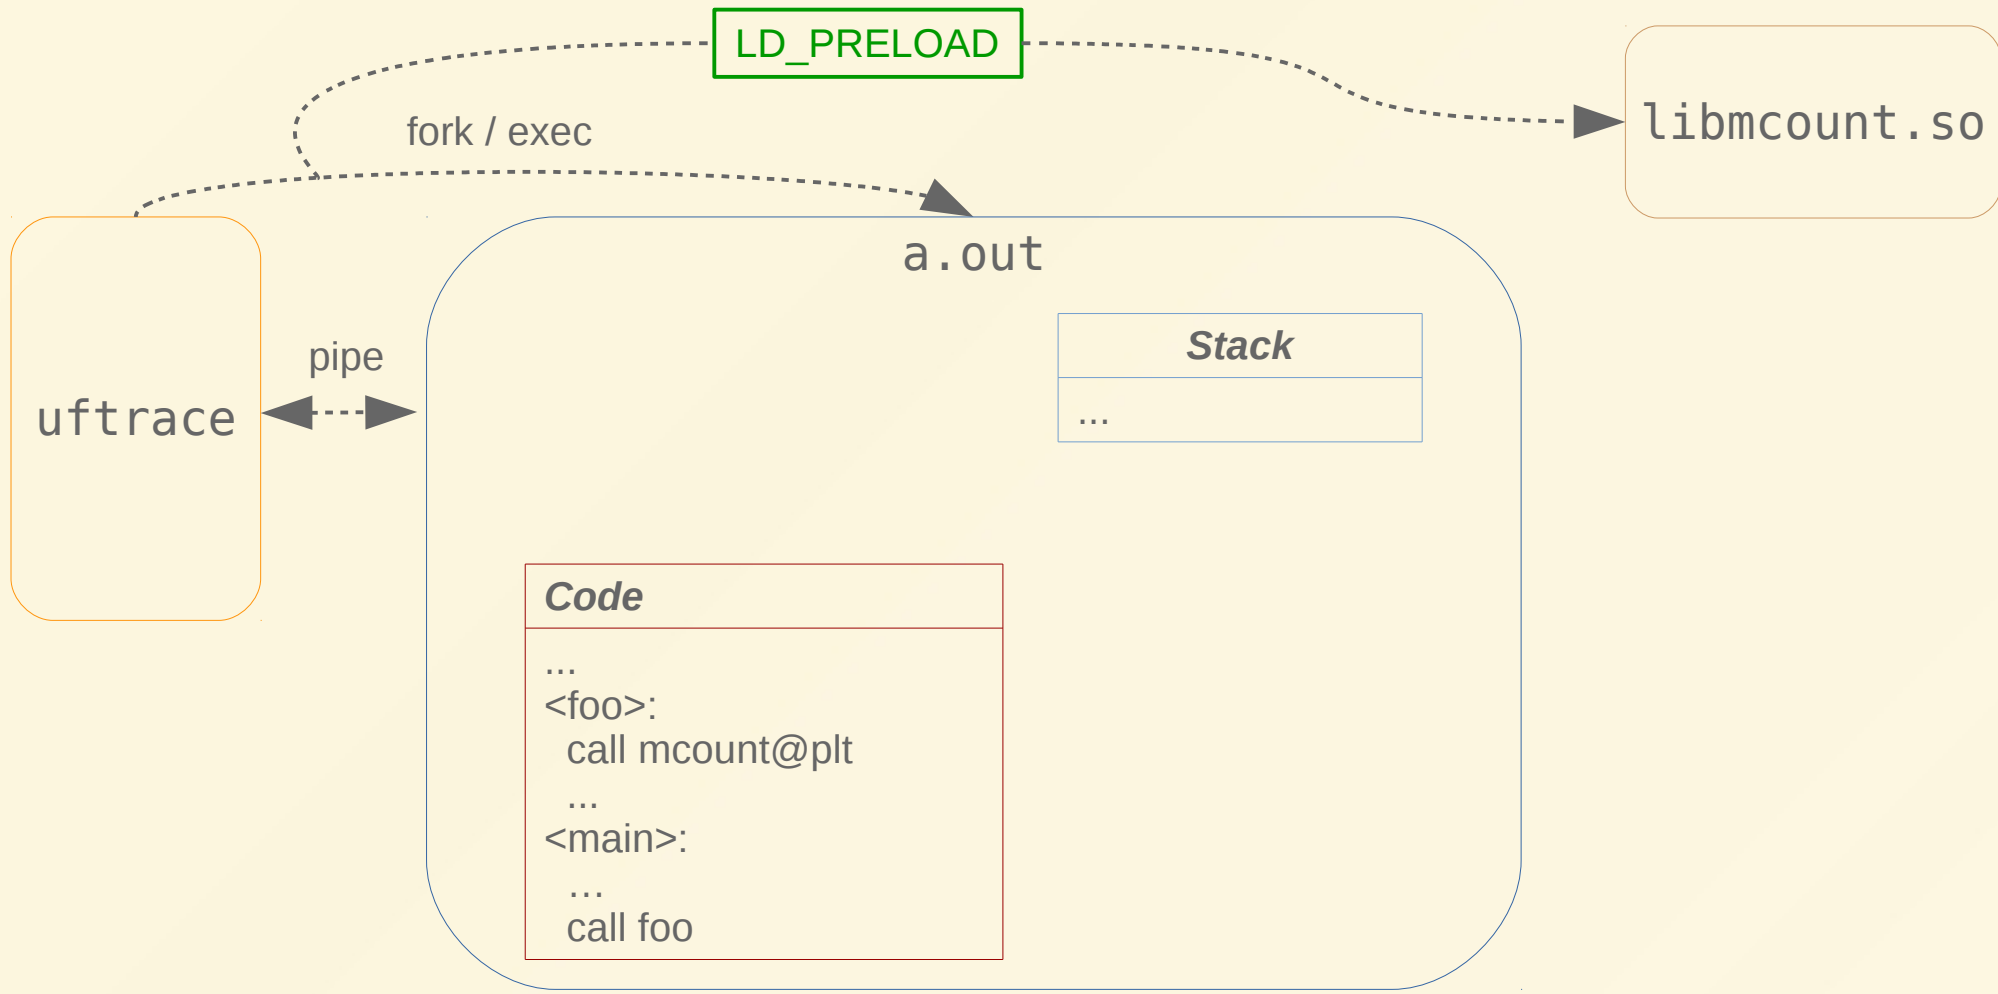
\$ gcc -pg

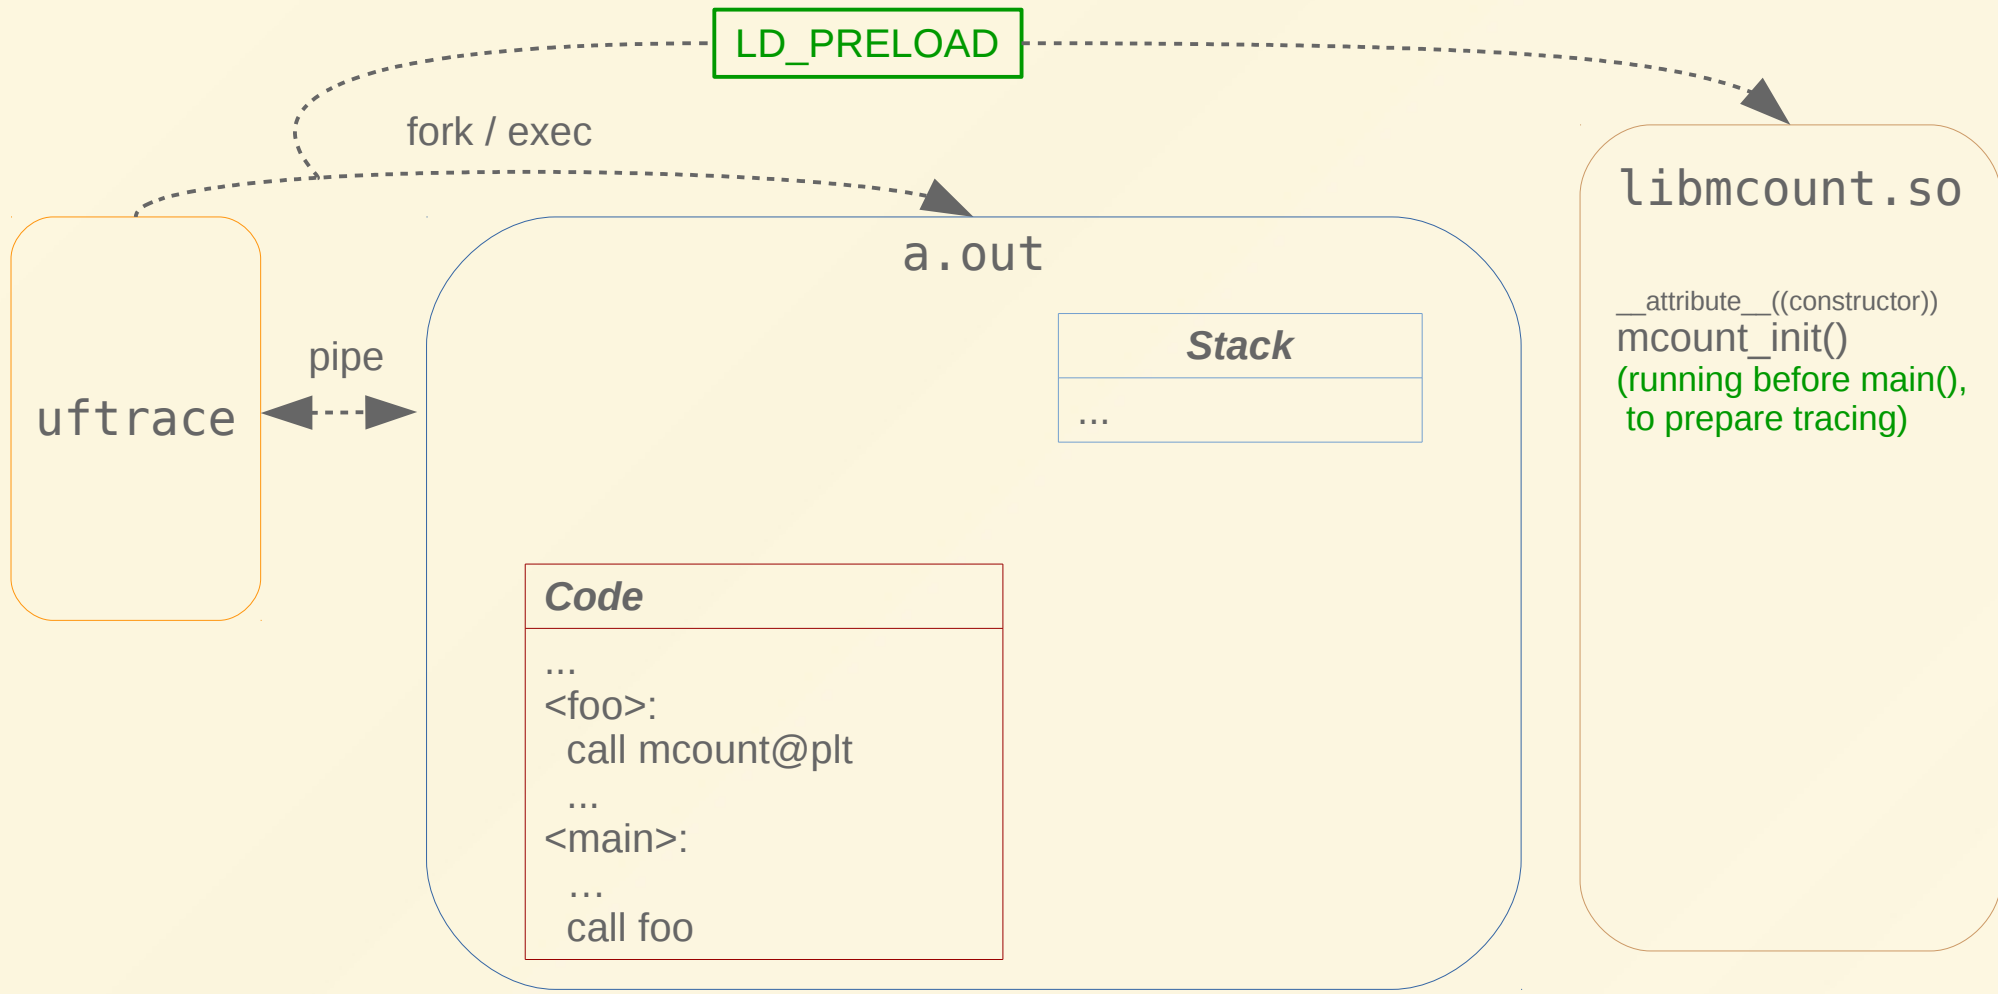
a.out

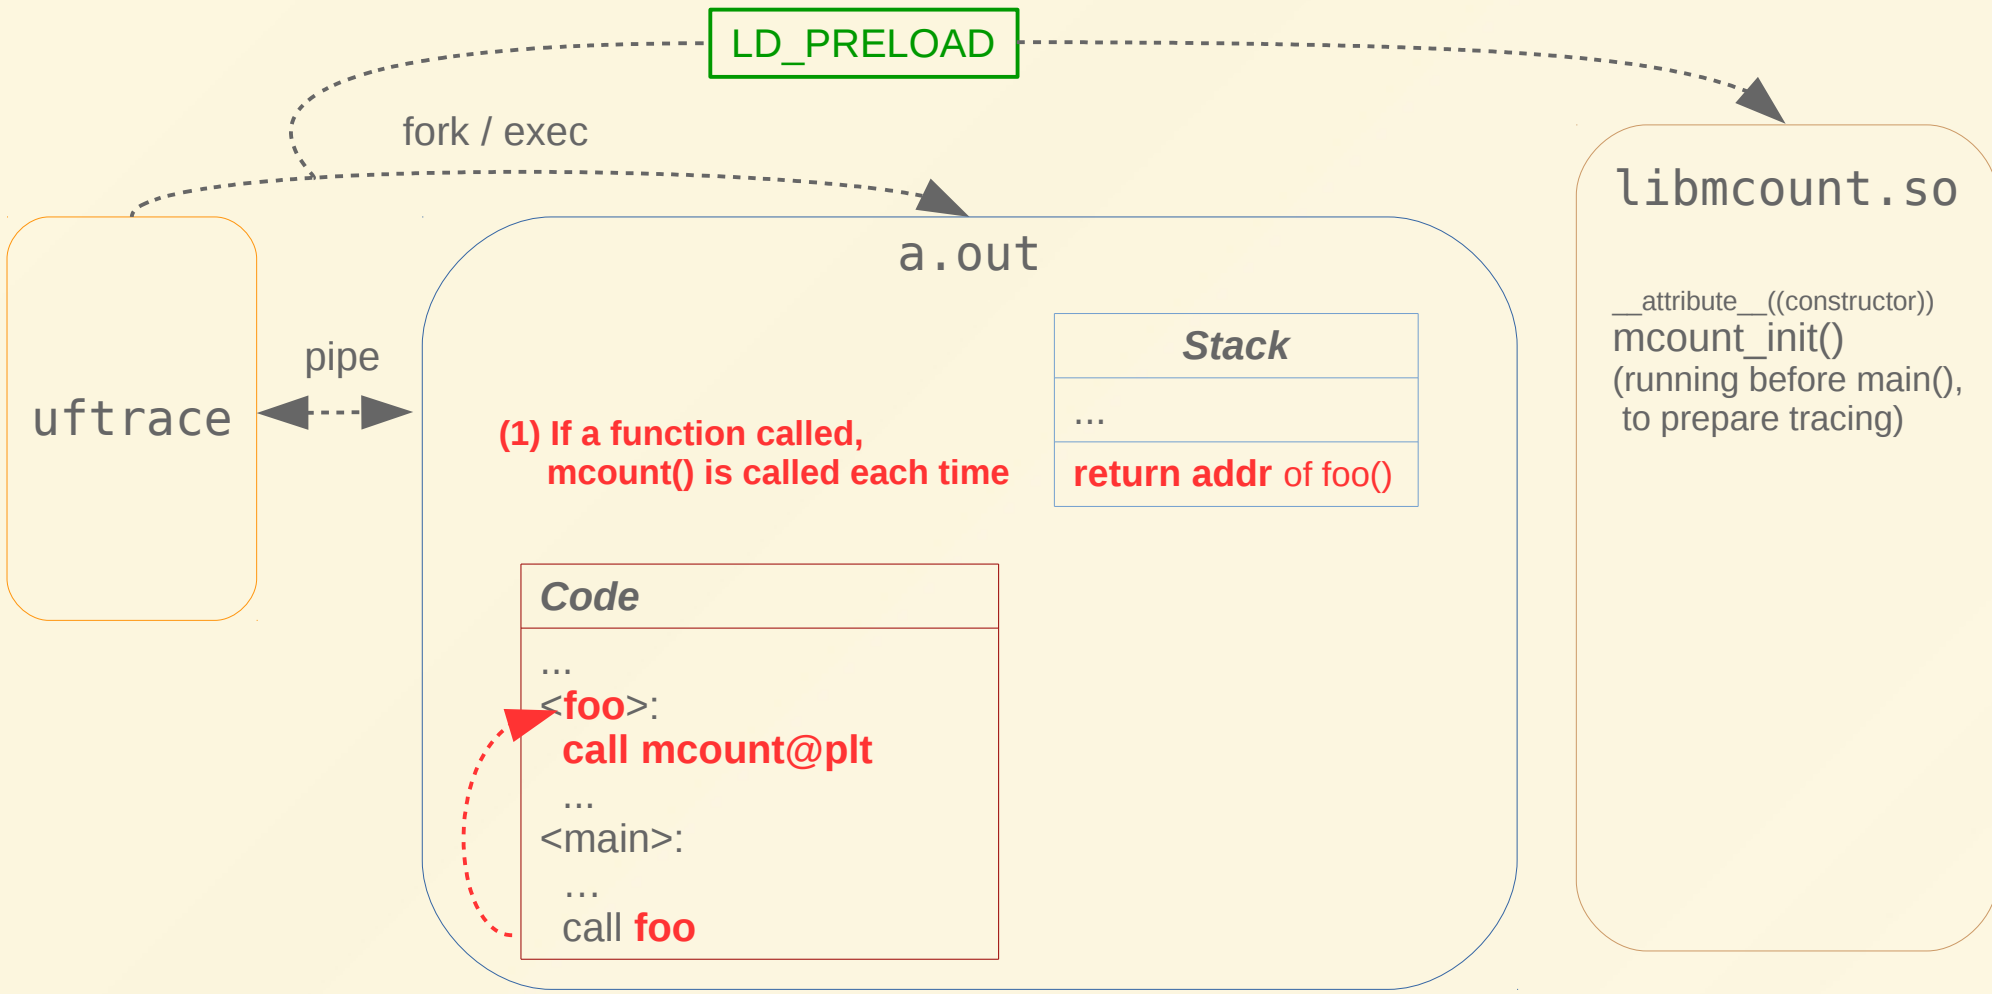
Code

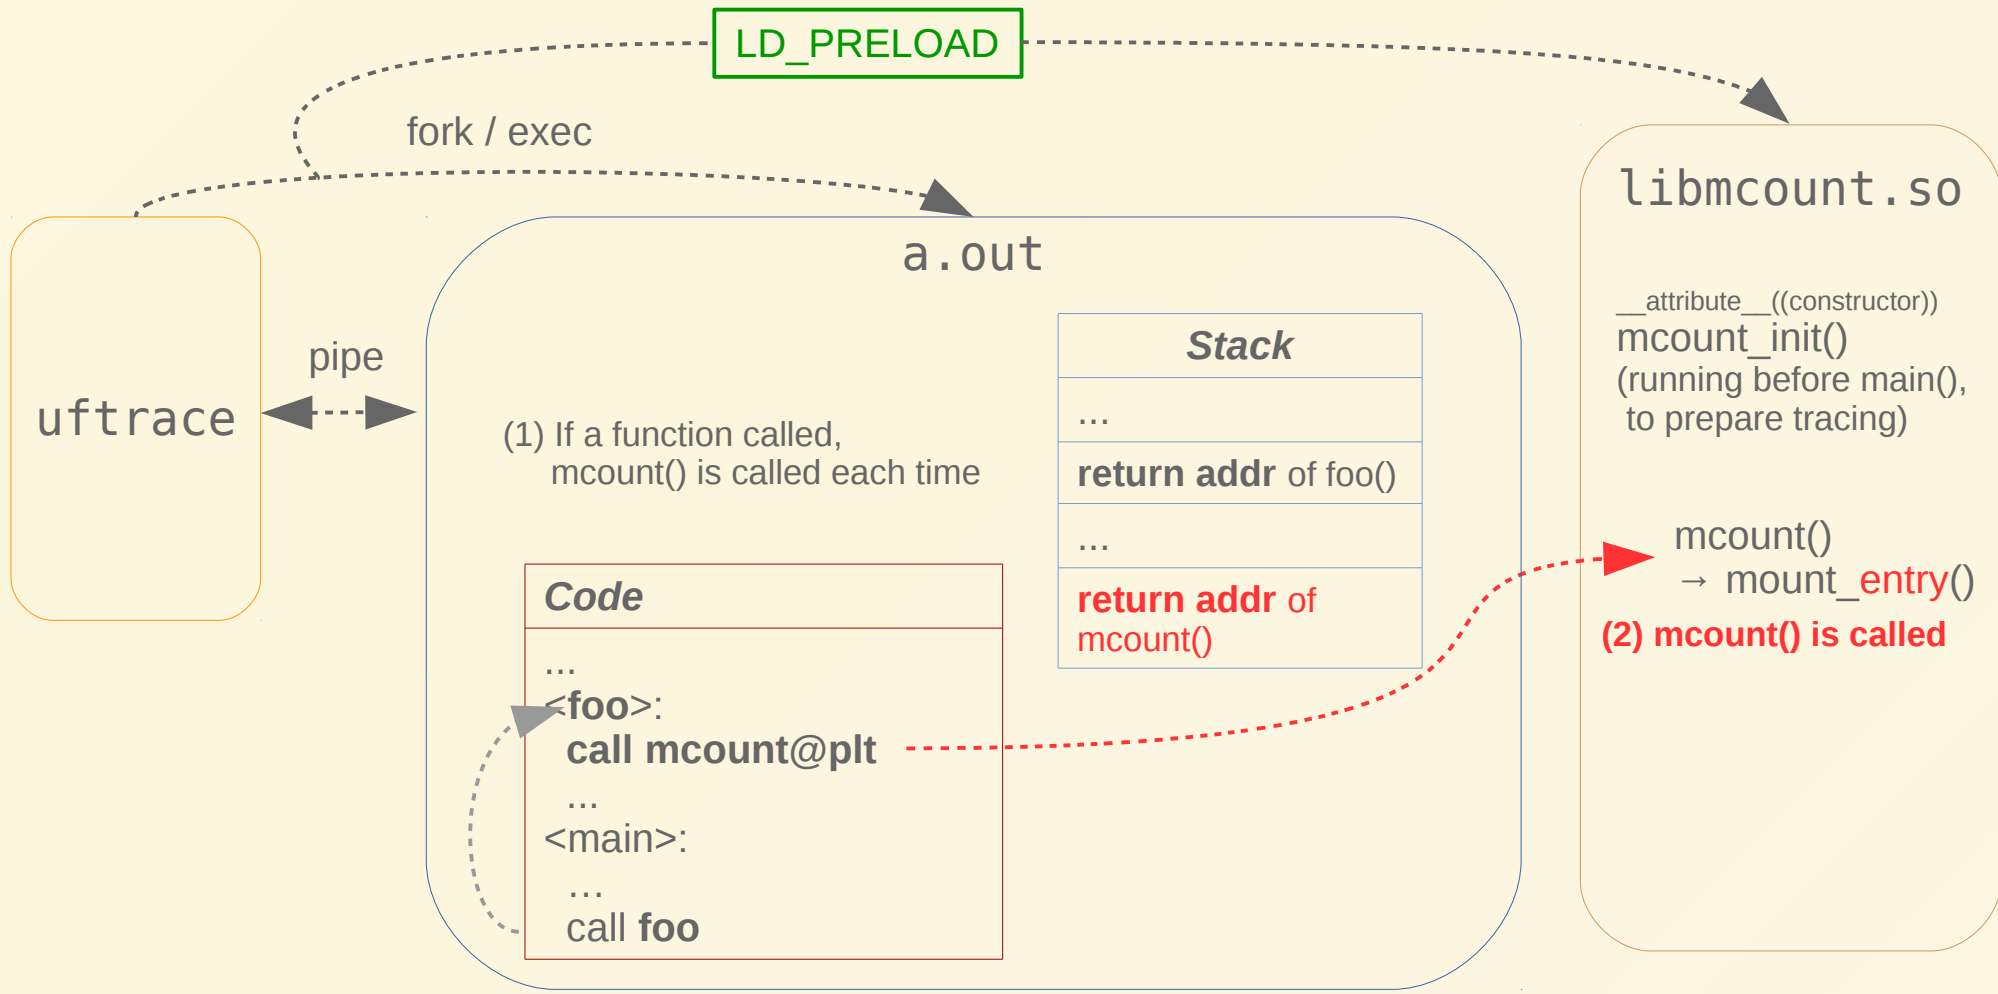
```
...  
<foo>:  
  call <mcount@plt>  
...  
<main>:  
  call <mcount@plt>  
...
```

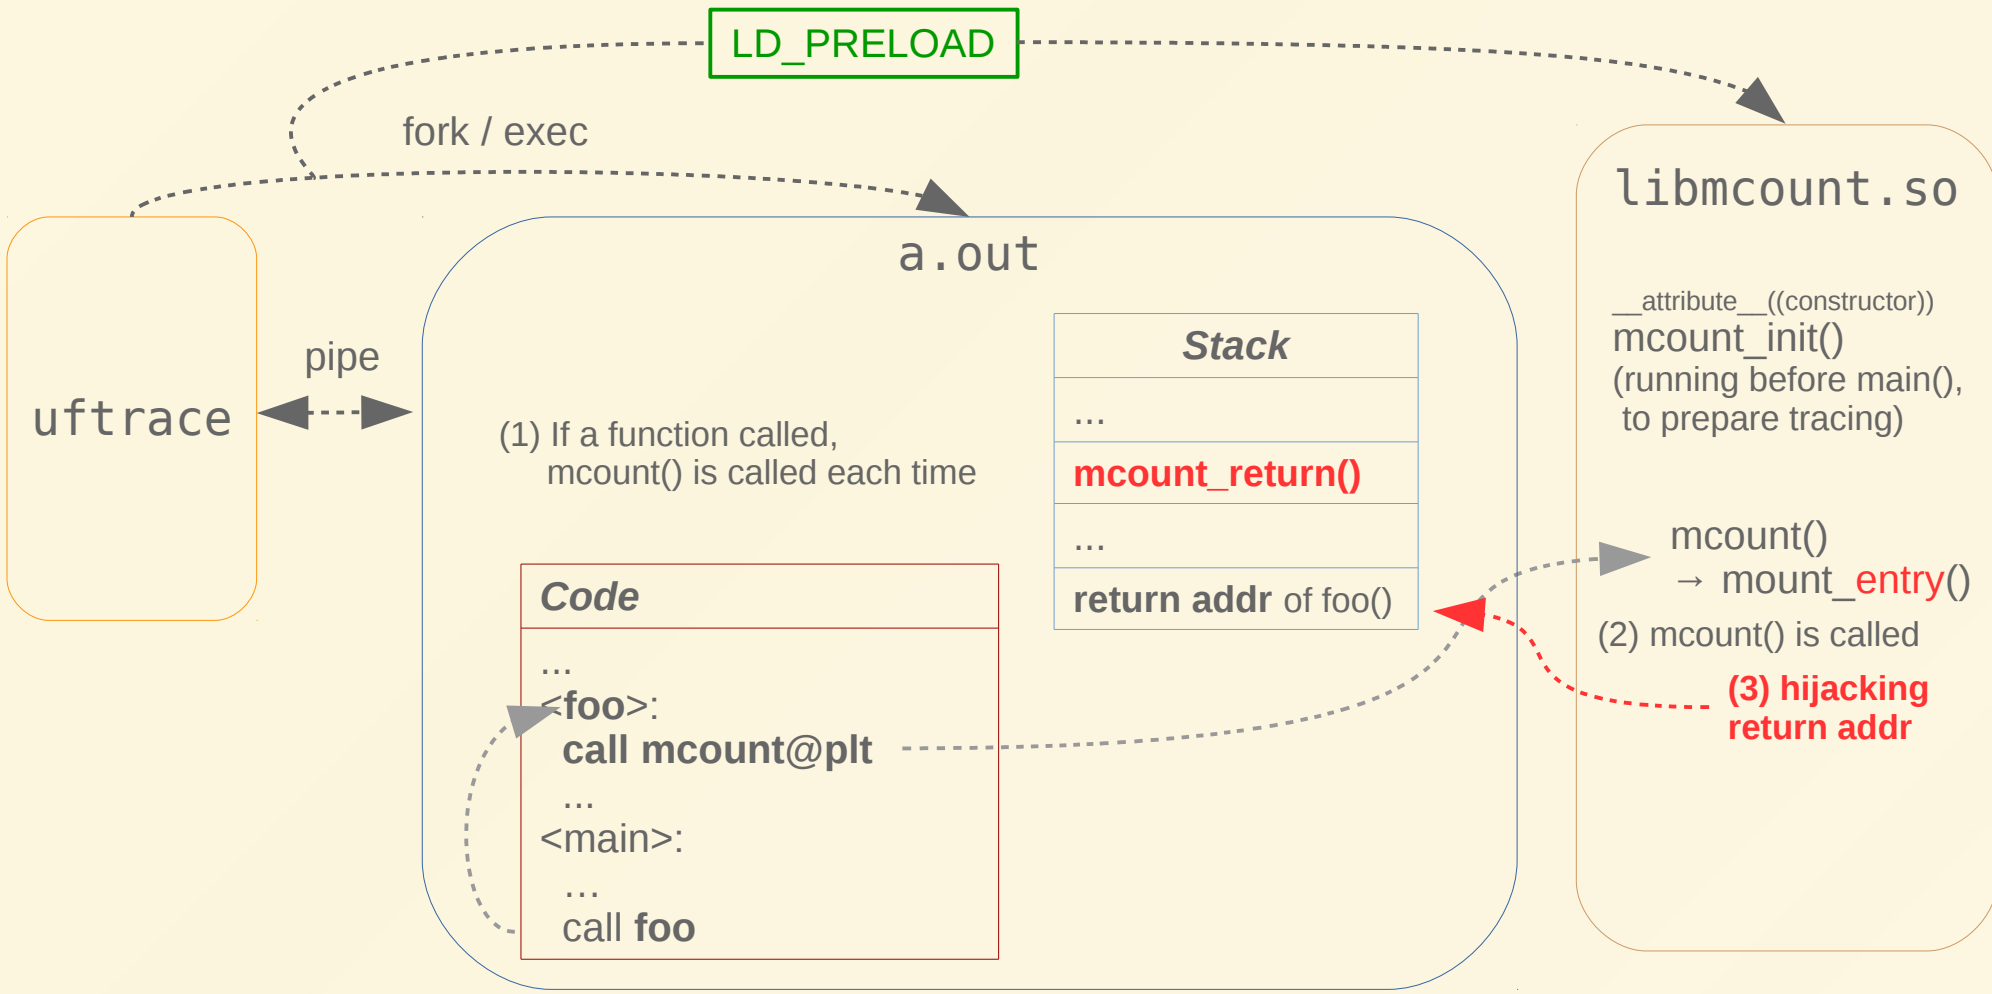


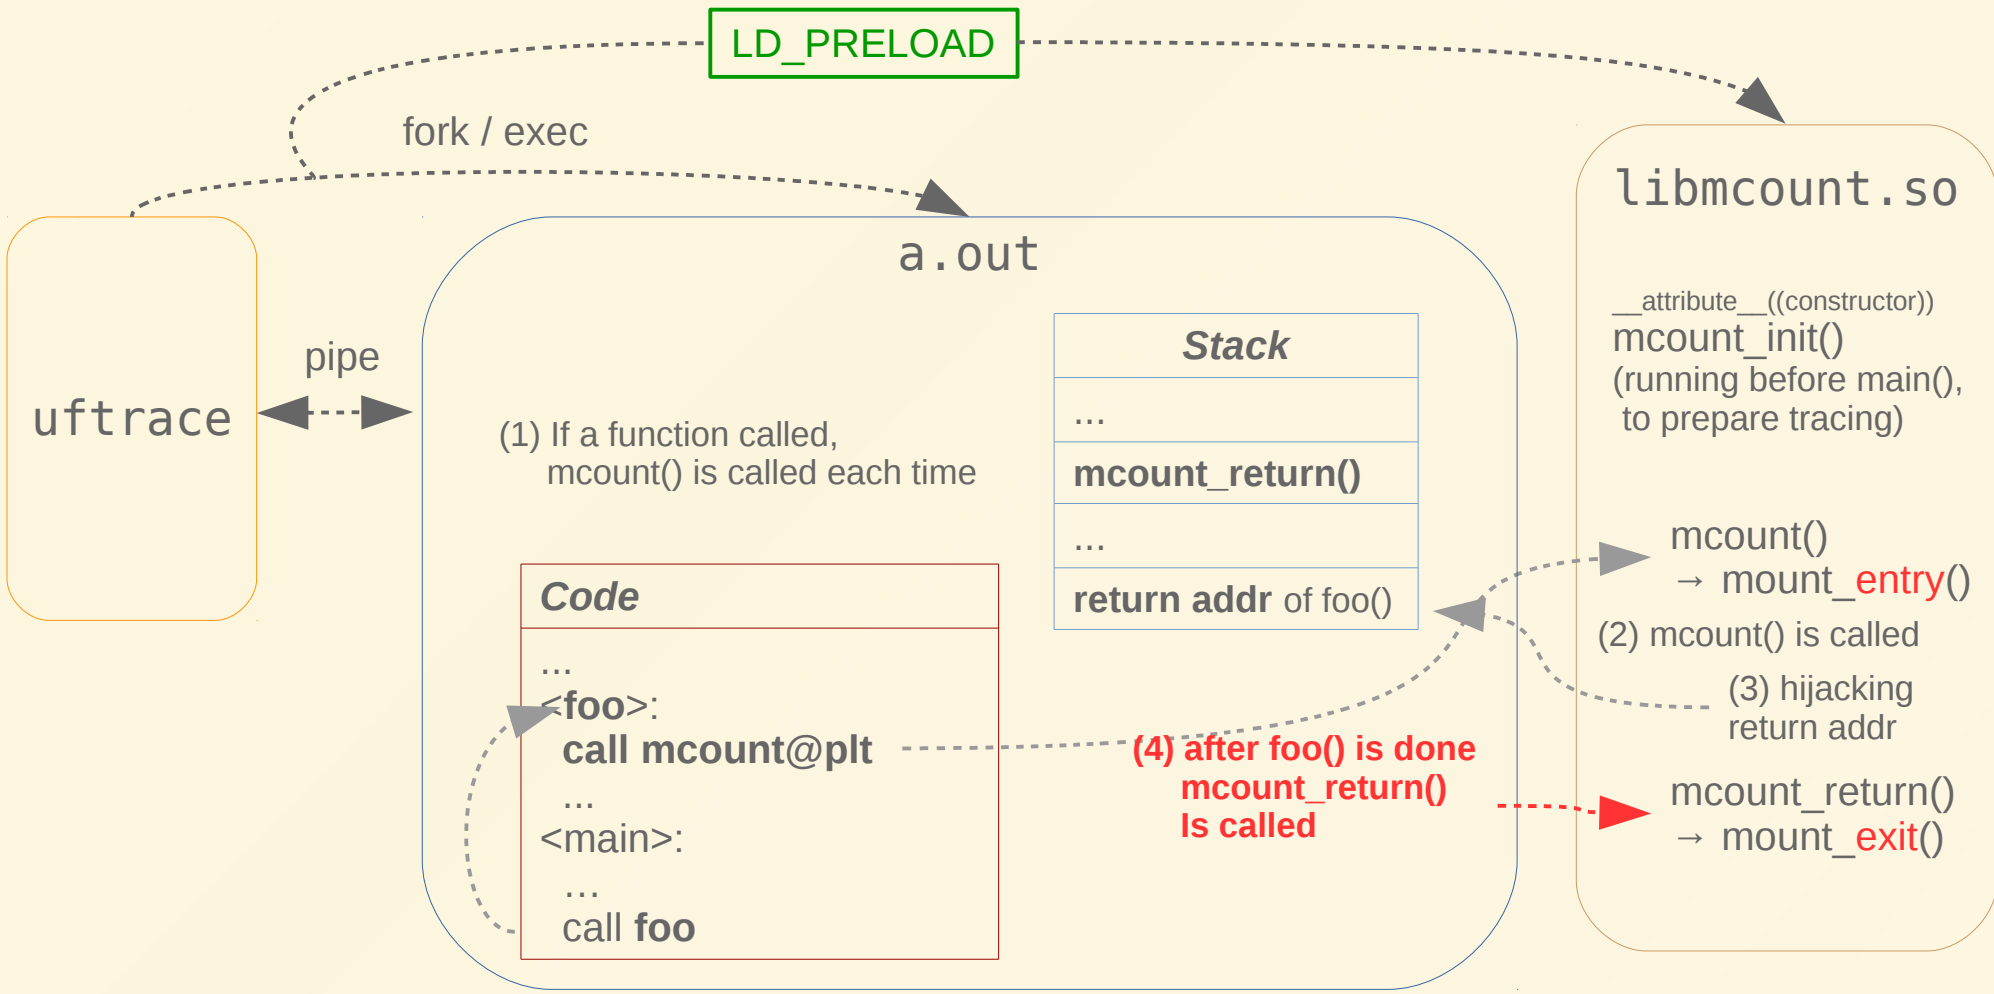












LD_PRELOAD

fork / exec

a.out

uftrace

pipe

(1) If a function called,
mcount() is called each time

Code

```
...  
<foo>:  
  call mcount@plt  
...  
<main>:  
  ...  
  call foo
```

Stack

```
...  
mcount_return()  
...  
return addr of foo()
```

libmcount.so

```
__attribute__((constructor))  
mcount_init()  
(running before main(),  
to prepare tracing)
```

mcount()
→ mount_entry()

(2) mcount() is called
(3) hijacking
return addr

mcount_return()
→ mount_exit()

record_trace_data()

(4) after foo() is done
mcount_return()
is called

shmem

rstack

...
...
...

(5) save trace data

- Func duration(exit time – enrty time)
- Func arguments, retun values
- Function call graph ...

LD_PRELOAD

fork / exec

a.out

libmcount.so

uftrace

uftrace.data/<tid>.dat

pipe

(1) If a function called,
mcount() is called each time

Code

```
...  
<foo>:  
  call mcount@plt  
...  
<main>:  
  ...  
  call foo
```

Stack

```
...  
mcount_return()  
...  
return addr of foo()
```

__attribute__((constructor))
mcount_init()
(running before main(),
to prepare tracing)

mcount()
→ mount_entry()

(2) mcount() is called
(3) hijacking
return addr

mcount_return()
→ mount_exit()

record_trace_data()

(4) after foo() is done
mcount_return()
is called

shmem

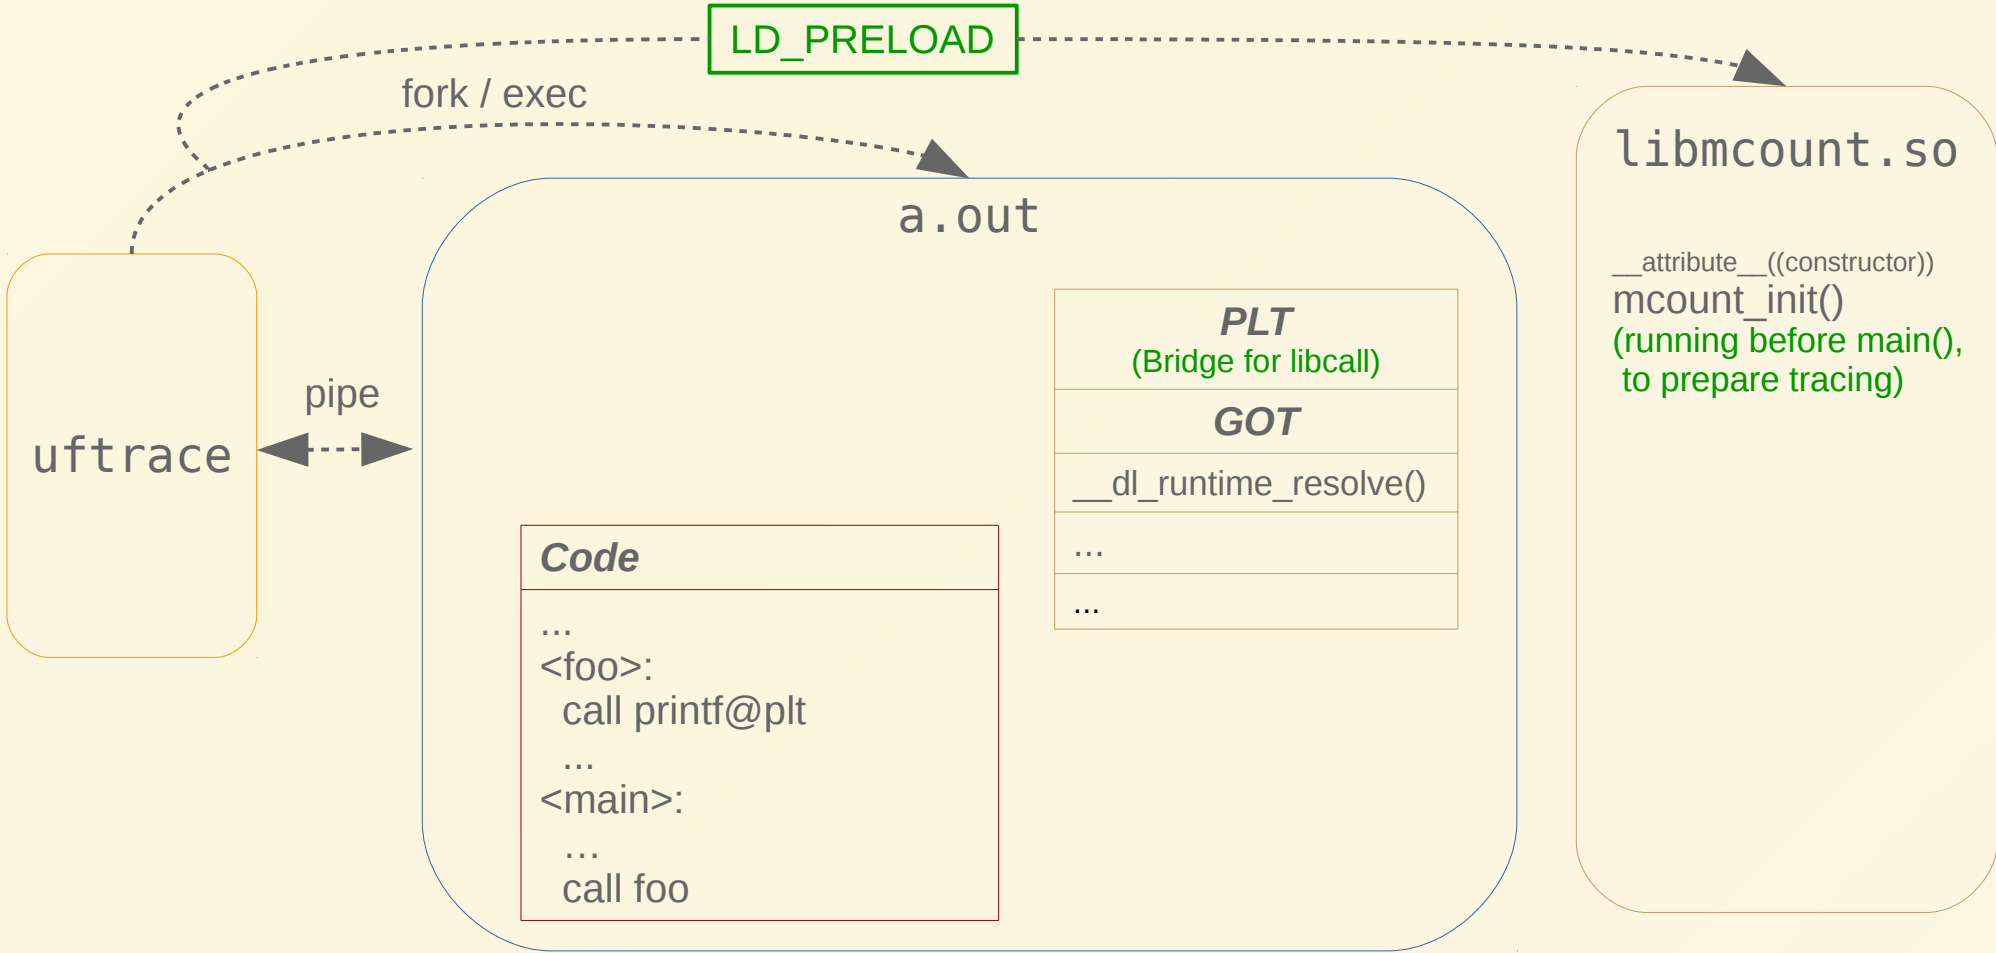
rstack

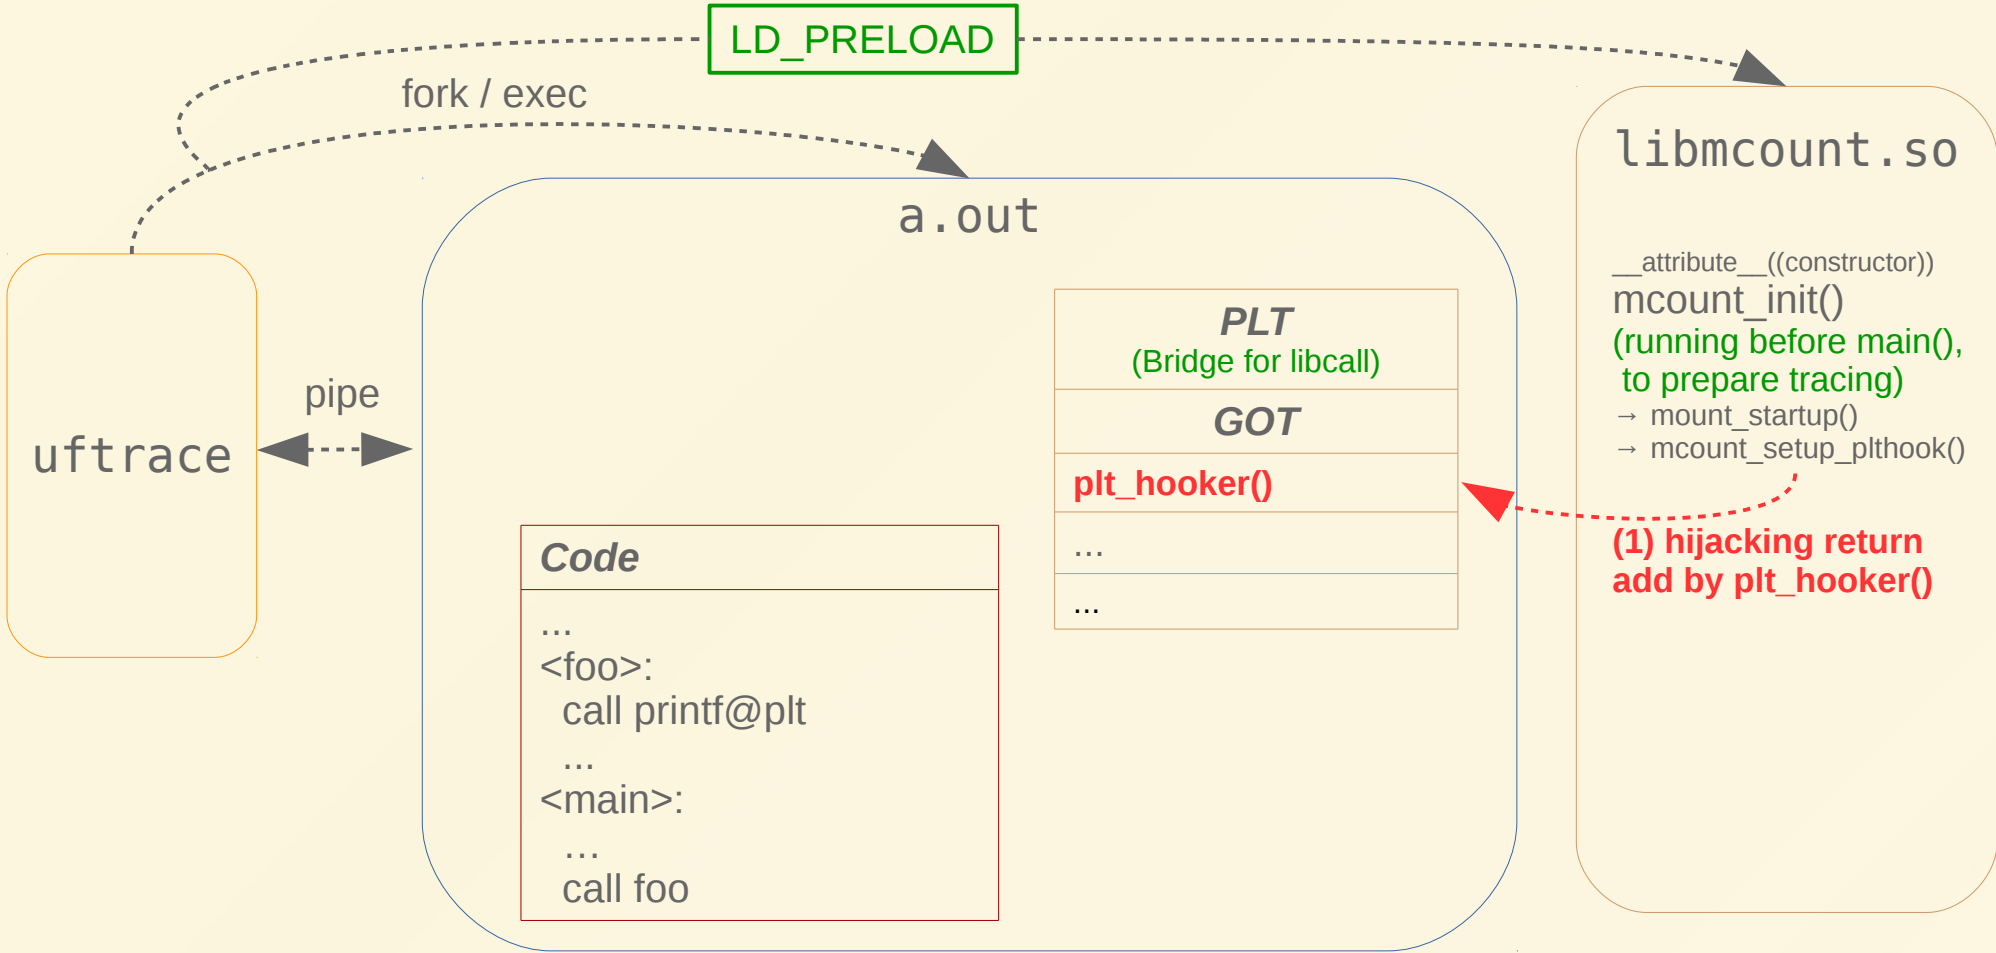
...
...
...

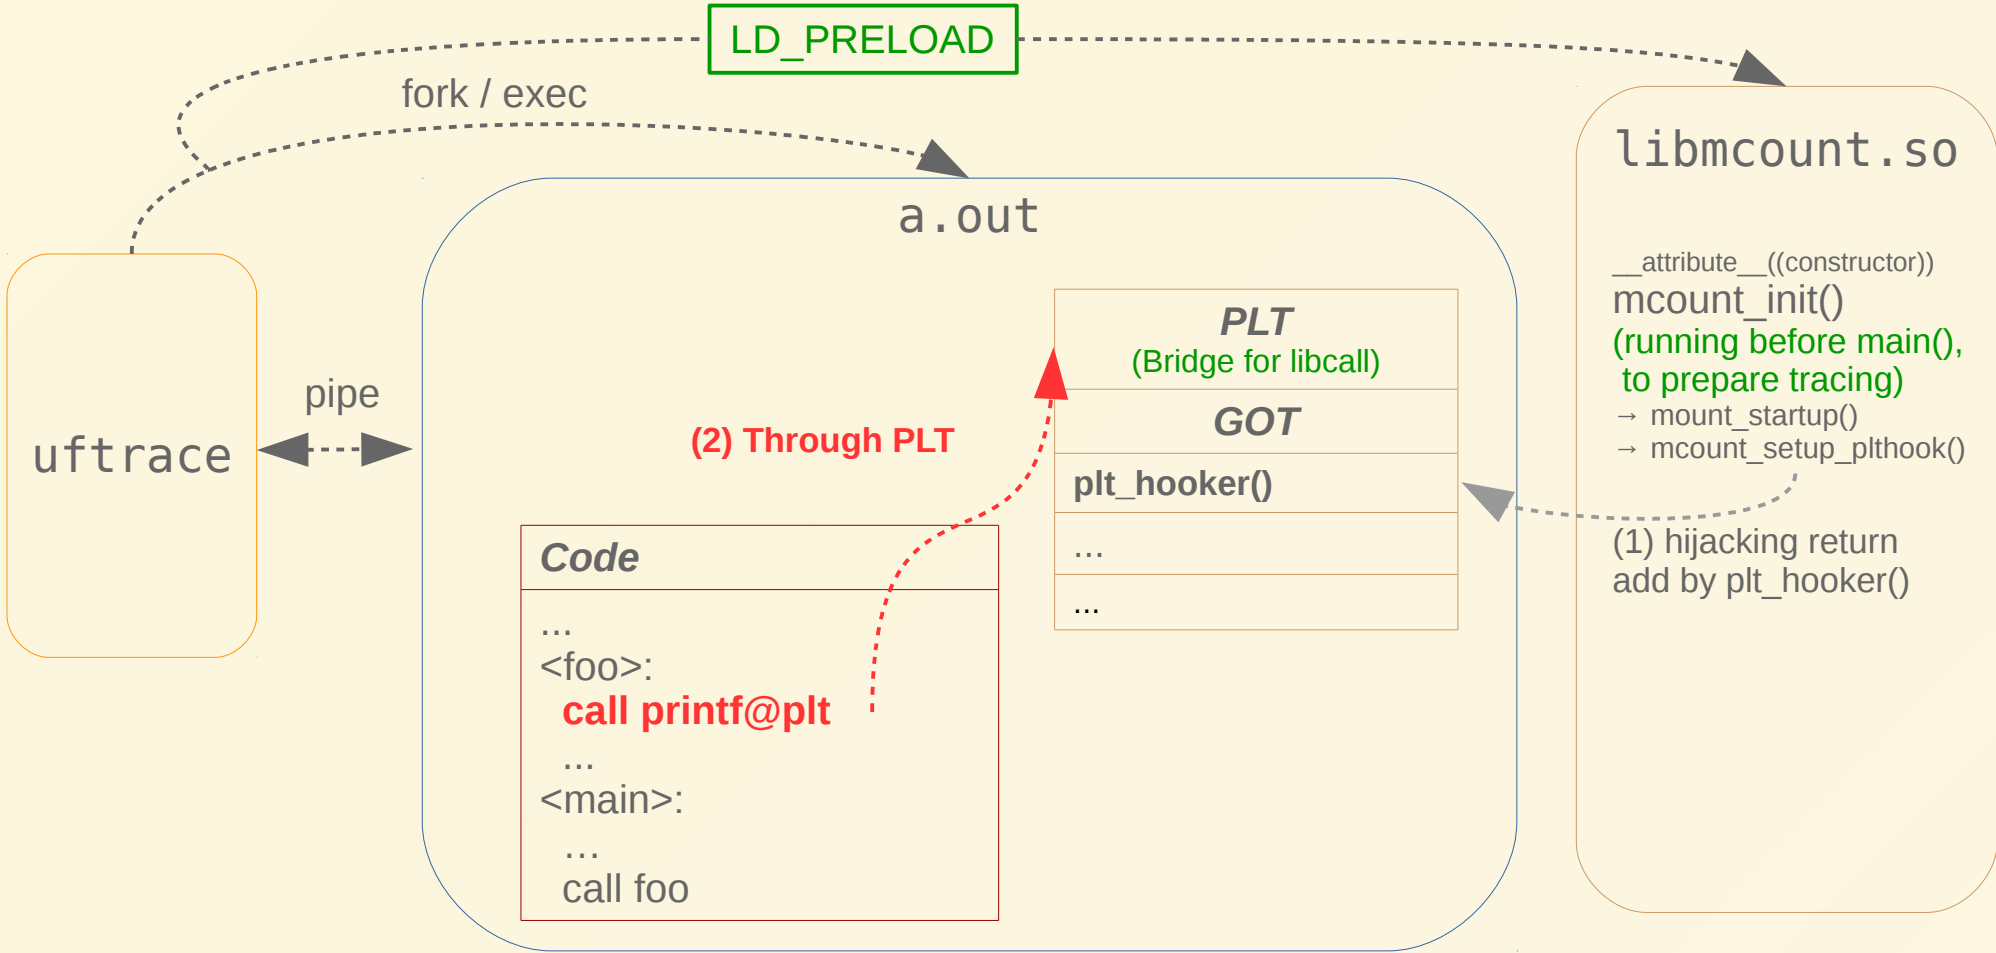
(5) save trace data

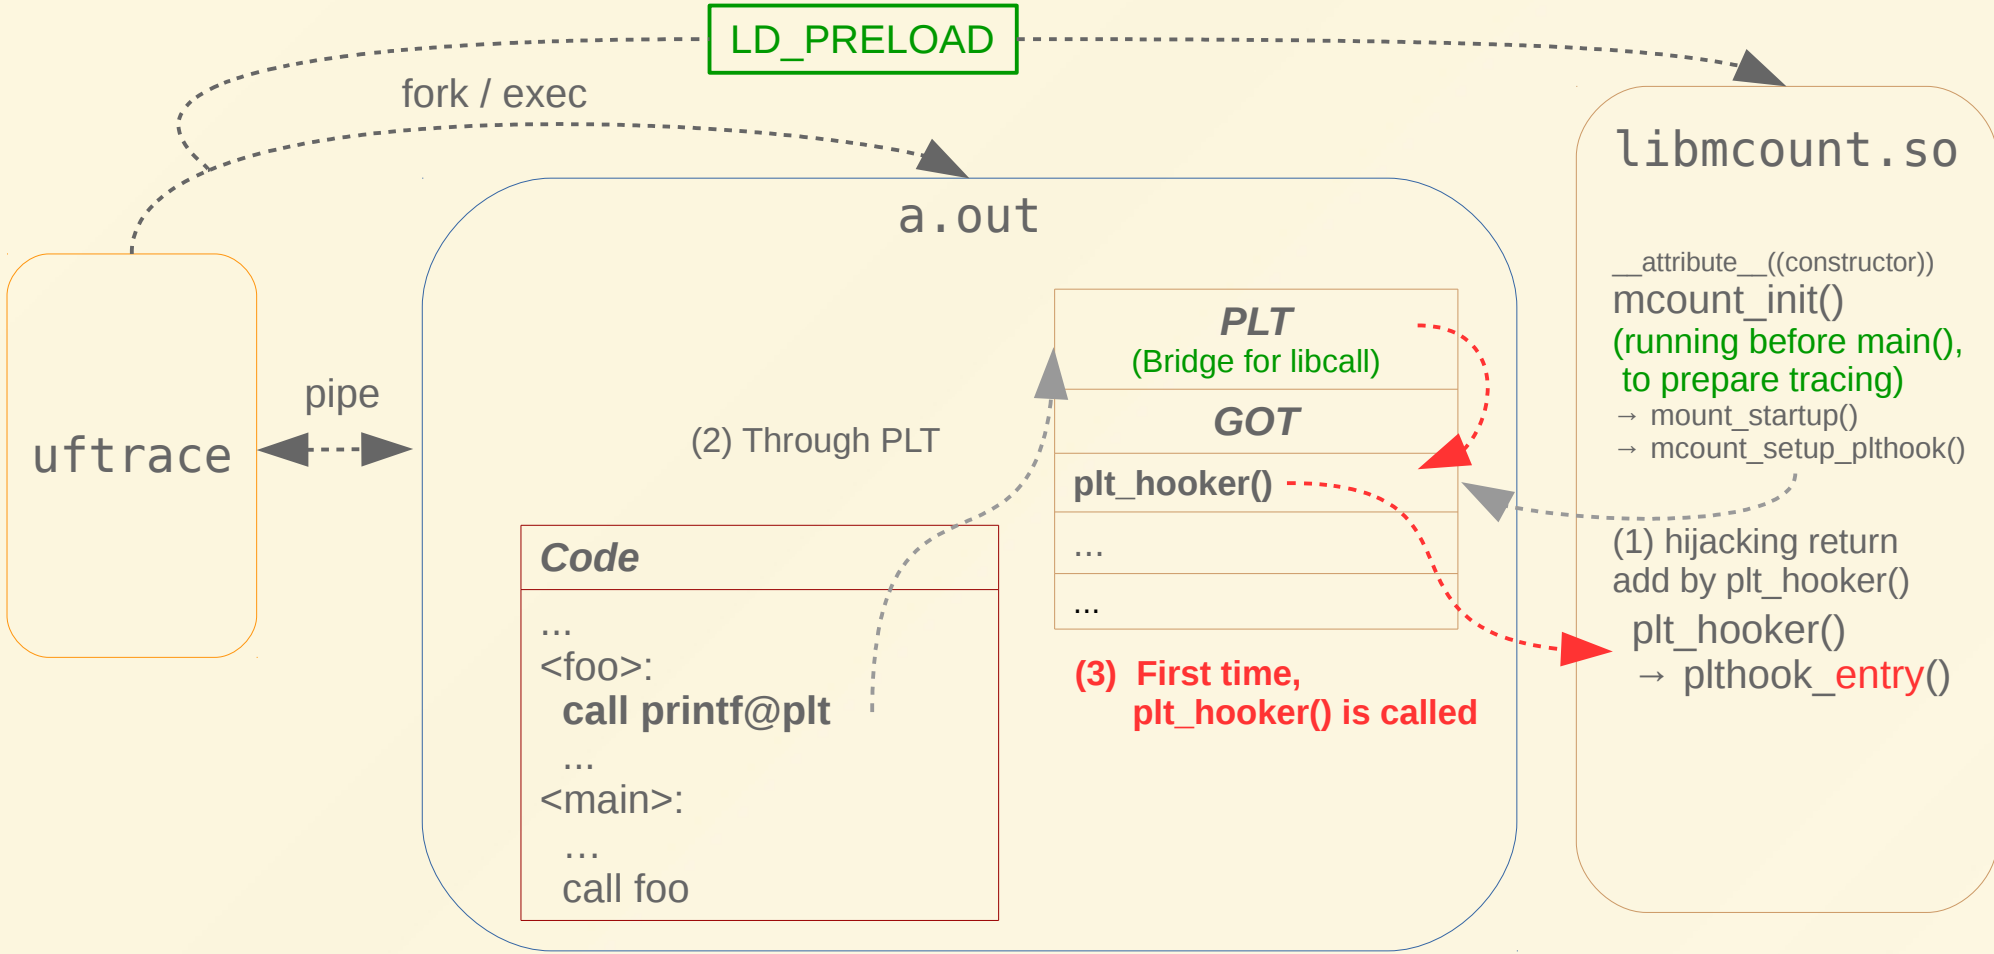
- Func duration(exit time – enrty time)
- Func arguments, retun values
- Function call graph ...

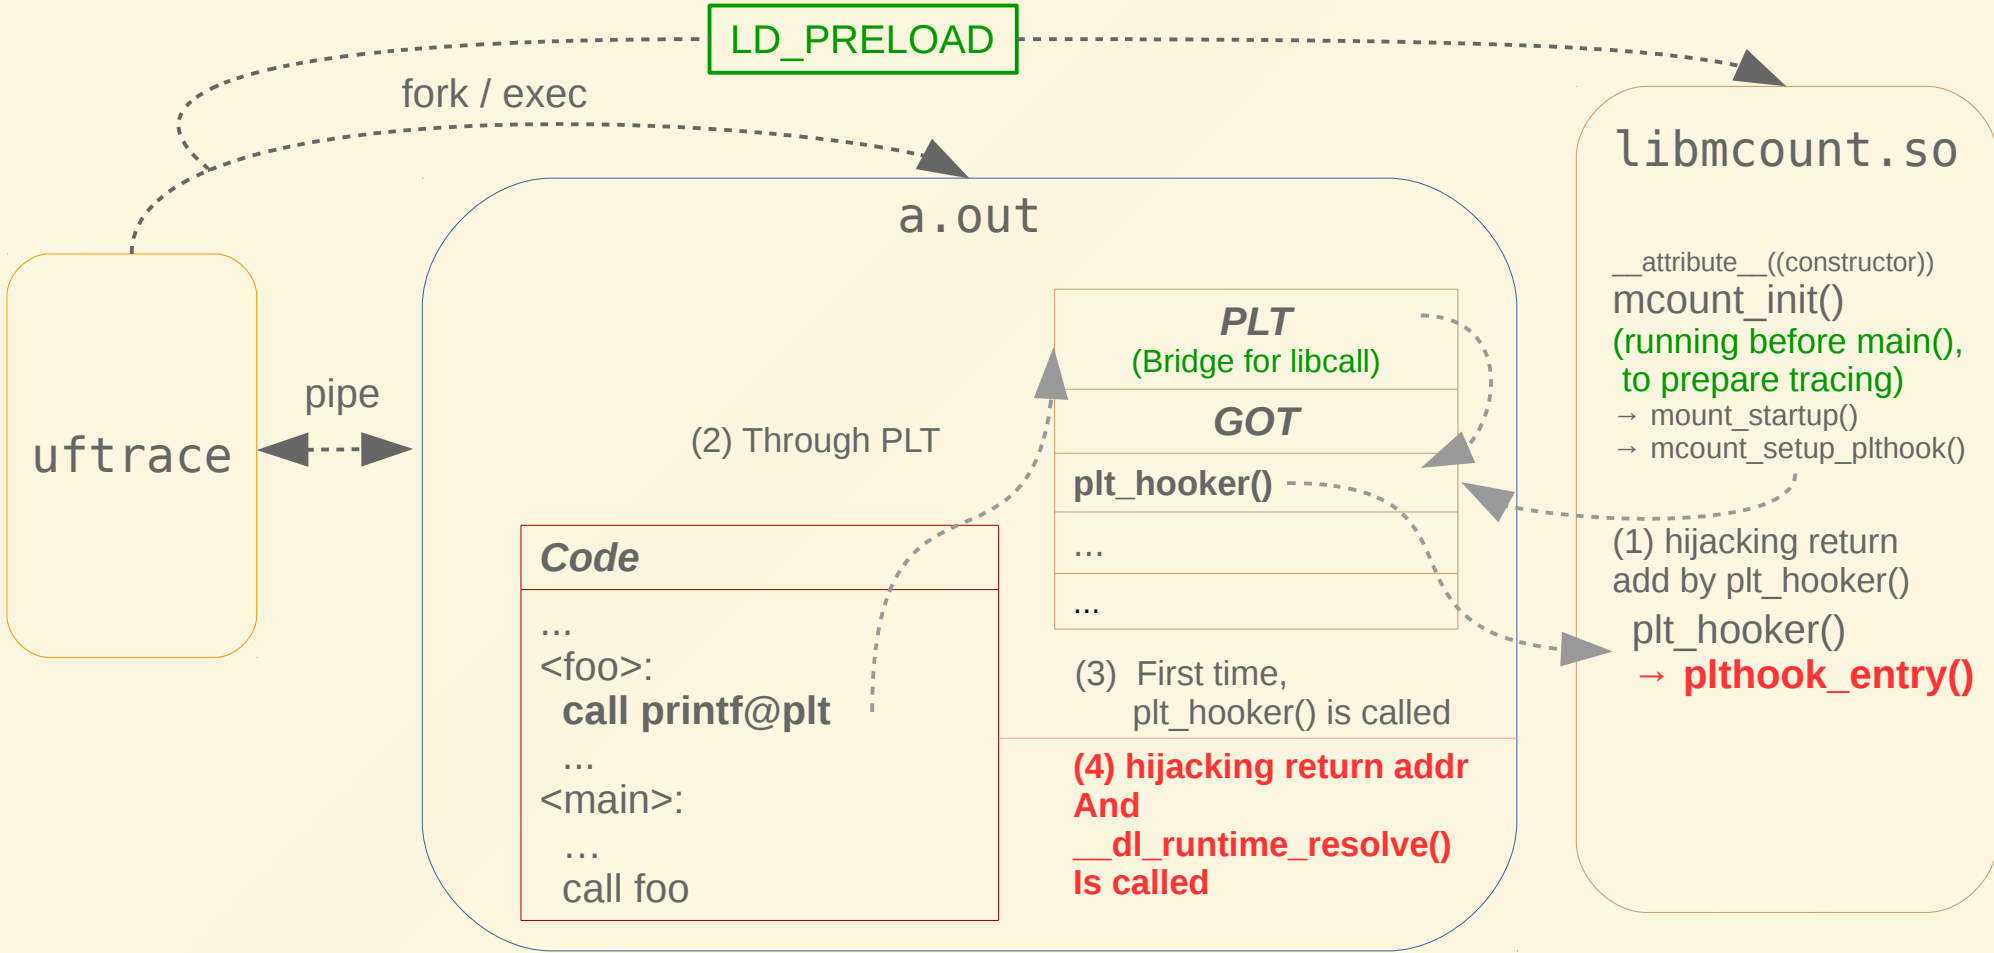
PLT hooking (library call tracing)











LD_PRELOAD

fork / exec

a.out

libmcount.so

uftrace

pipe

(2) Through PLT

Code

```
...  
<foo>:  
  call printf@plt  
...  
<main>:  
  ...  
  call foo
```

PLT

(Bridge for libcall)

GOT

plt_hooker()

...

...

(3) First time,
plt_hooker() is called

(4) hijacking return addr
And
__dl_runtime_resolve()
Is called

```
__attribute__((constructor))  
mcount_init()  
(running before main(),  
to prepare tracing)  
→ mount_startup()  
→ mcount_setup_plthook()
```

(1) hijacking return
addr by plt_hooker()

plt_hooker()
→ plthook_entry()

plt_return()
→ plthook_exit()

(5) separately save resolved
addr (e.g. addr of printf)

And restore contents of GOT

LD_PRELOAD

fork / exec

a.out

libmcount.so

uftrace

pipe

(2) Through PLT

Code

```
...  
<foo>:  
  call printf@plt  
...  
<main>:  
  ...  
  call foo
```

PLT

(Bridge for libcall)

GOT

plt_hooker()

...
...

(3) First time,
plt_hooker() is called

(4) hijacking return addr
And
__dl_runtime_resolve()
Is called

```
__attribute__((constructor))  
mcount_init()  
(running before main(),  
to prepare tracing)  
→ mount_startup()  
→ mcount_setup_plthook()
```

(1) hijacking return
addr by plt_hooker()

```
plt_hooker()  
→ plthook_entry()  
record_trace_data()  
plt_return()  
→ plthook_exit()
```

(5) separately save resolved
addr (e.g. addr of printf)

And restore contents of GOT

shmem

rstack

...
...
...

(6) trace data

- Func duration(exit time – enrtly time)
- Func arguments, return values
- Function call graph ...